
Reinforcement Learning For Automated Trading

Pierpaolo G. Necchi

Mathematical Engineering

Politecnico di Milano

Milano, IT 20123

pierpaolo.necchi@gmail.com

Abstract

The impact of *Automated Trading Systems* (ATS) on financial markets is growing every year and the trades generated by an algorithm now account for the majority of orders that arrive at stock exchanges. In this paper we explore how to find a trading strategy via *Reinforcement Learning* (RL), a branch of *Machine Learning* (ML) that allows to find an optimal strategy for a sequential decision problem by directly interacting with the environment. We show that the long-short strategy learned for a synthetic asset, whose price follows a stochastic process with some exploitable patterns, consistently outperforms the market. RL thus shows the potential to deal with many financial problems, that can be often formulated as sequential decision problems.

1 Introduction

The trading pit of a stock exchange is often imagined by outsiders as a frenzy place, with telephones constantly ringing and traders shouting orders across the room at a frenetic rhythm. This was probably the reality thirty years ago, when open outcry was still the main communication system between pit traders. Since then the floors have become more and more quiet as the majority of the orders moved to electronic trading systems. Notwithstanding, investment decisions were still made by humans who could now execute their orders without passing through the pit traders. In the last decade, the markets have witnessed the widespread adoption of *Automated Trading Systems* (ATS), that can make investment decisions in a fully automatized way at speeds with orders of magnitude greater than any human equivalent. In 2014, more than 75% of the stock shares traded on United States exchanges were originated from ATS orders and this amount kept growing since then. Quantitative hedge funds, such as Renaissance Technologies, D.E. Shaw, Citadel and many others, are employing mathematicians, physicists and other scientists to develop algorithms able to extract trading signals from large amount of data and automatically trade. These algorithms are typically based on advanced statistics, signal processing, machine learning and other fields of mathematics. However, few of these hedge funds publish their profit-generating “secret sauce” and not much can be found in the literature. In this project we develop an automated trading algorithm based on *Reinforcement Learning* (RL), a branch of *Machine Learning* (ML) which has recently been in the spotlight for being at the core of the system who beat the Go world champion in a 5-match series [1].

This document is organized as follows. In Section 2 we introduce the basic concepts of RL and present two learning algorithms that allow to determine an approximation for the optimal policy of a sequential decision problem. In section 3 we discuss the asset allocation problem from a mathematical point of view and show how these learning algorithms can be applied in this setting. In Section 4 we start discussing the implementation of the model in Python, which has been used during the prototyping phase. In Section 5 we discuss a more efficient C++ implementation. In Section 6 we describe the execution pipeline used to run the learning experiment. In Section 7 we present the numerical results for a synthetic asset, whose price follows a particular stochastic process. In Section 8 we conclude with some final remarks and we discuss some future research directions.




Figure 1: Agent-environment interaction in sequential decision problems.

2 Basics of Reinforcement Learning

Reinforcement Learning (RL) is a general class of algorithms in the field of *Machine Learning* (ML) that allows an agent to learn how to behave in a stochastic and possibly unknown environment, where the only feedback consists of a scalar reward signal [2]. The goal of the agent is to learn by trial-and-error which actions maximize his long-run rewards. However, since the environment evolves stochastically and may be influenced by the actions chosen, the agent must balance his desire to obtain a large immediate reward by acting greedily and the opportunities that will be available in the future. Thus, RL algorithms can be seen as computational methods to solve sequential decision problems by directly interacting with the environment.

2.1 Markov Decision Processes

Sequential decision problems are typically formalized using *Markov Decision Processes* (MDP). An MDP is a stochastic dynamical system specified by the tuple $\langle \mathbb{S}, \mathbb{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$, where $(\mathbb{S}, \mathcal{S})$ is a measurable state space, $(\mathbb{A}, \mathcal{A})$ is a measurable action space, $\mathcal{P} : \mathbb{S} \times \mathbb{A} \times \mathcal{S} \rightarrow \mathbb{R}$ is a Markov transition kernel, $\mathcal{R} : \mathbb{S} \times \mathbb{A} \rightarrow \mathbb{R}$ is a reward function and $0 < \gamma < 1$ is the discount factor. Suppose that at time t the system is in state $S_t = s$ and that the agent takes action $A_t = a$, then, regardless of the previous history of the system, the probability to find the system in a state belonging to $B \in \mathcal{S}$ at time $t + 1$ is given by

$$\mathcal{P}(s, a, B) = \mathbb{P}(S_{t+1} \in B | S_t = s, A_t = a) \quad (1)$$

Following this random transition, the agent receives a stochastic reward R_{t+1} . The reward function $\mathcal{R}(s, a)$ gives the expected reward obtained when action a is taken in state s , i.e.

$$\mathcal{R}(s, a) = \mathbb{E}[R_{t+1} | S_t = s, A_t = a] \quad (2)$$

This feedback mechanism between the environment and the agent is illustrated in Figure 1. At any time step, the agent selects his actions according to a certain policy $\pi : \mathbb{S} \times \mathcal{A} \rightarrow \mathbb{R}$ such that for every $s \in \mathbb{S}$, $C \mapsto \pi(s, C)$ is a probability distribution over $(\mathbb{A}, \mathcal{A})$. Hence, a policy π and an initial state $s_0 \in \mathbb{S}$ determine a random state-action-reward sequence $\{(S_t, A_t, R_{t+1})\}_{t \geq 0}$ with values on $\mathbb{S} \times \mathbb{A} \times \mathbb{R}$. In an infinite horizon task, the agent's performance is typically measured as the total discounted reward obtained following a specific policy

$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \quad (3)$$

Since this gain is stochastic, the agent considers its expected value, which is typically called *state-value function*

$$V_{\pi}(s) = \mathbb{E}_{\pi}[G_t | S_t = s] \quad (4)$$

where the subscript in \mathbb{E}_{π} indicates that all the actions are selected according to policy π . The state-value function measures how good it is for the agent to be in a given state and follow a certain policy. Similarly, we introduce the *action-value function*

$$Q_{\pi}(s, a) = \mathbb{E}_{\pi}[G_t | S_t = s, A_t = a] \quad (5)$$

We have the following relationship between V_π and Q_π

$$V_\pi(s) = \int_{\mathbb{A}} \pi(s, a) Q_\pi(s, a) da \quad (6)$$

Almost all reinforcement learning algorithms are designed to estimate these value functions and are typically based on the Bellman equations.

$$V_\pi(s) = \mathcal{R}(s) + \gamma T_\pi V_\pi(s) \quad (7)$$

$$Q_\pi(s, a) = \mathcal{R}(s, a) + \gamma T_a V_\pi(s) \quad (8)$$

where we denoted by T_a (resp. T_π) the transition operator for action a (resp. for policy π)

$$T_a F(s) = \mathbb{E}[F(S_{t+1}) | S_t = s, A_t = a] = \int_{\mathbb{S}} \mathcal{P}(s, a, s') F(s') ds' \quad (9)$$

$$T_\pi F(s) = \mathbb{E}_\pi[F(S_{t+1}) | S_t = s] = \int_{\mathbb{A}} \pi(s, a) \int_{\mathbb{S}} \mathcal{P}(s, a, s') F(s') ds' da \quad (10)$$

These equations can be rewritten as fixed-point equations which, under some formal assumptions on the reward functions, admit a unique solution by the contraction mapping theorem. The agent's goal is to select a policy π_* that maximizes his expected return in all possible states. Such a policy is called *optimal* and the corresponding value functions are called *Optimal State-Value Function*

$$V_*(s) = \sup_{\pi} V_\pi(s) \quad (11)$$

and *Optimal Action-Value Function*

$$Q_*(s, a) = \sup_{\pi} Q_\pi(s, a) \quad (12)$$

The optimal value functions satisfy the following Bellman equations.

$$V_*(s) = \sup_a Q_*(s, a) = \sup_a \{ \mathcal{R}(s, a) + \gamma T_a V_*(s) \} \quad (13)$$

$$\begin{aligned} Q_*(s, a) &= \mathcal{R}(s, a) + \gamma T_a V_*(s) \\ &= \mathcal{R}(s, a) + \gamma \int_{\mathbb{S}} \mathcal{P}(s, a, s') \sup_{a'} Q_*(s', a') ds' \end{aligned} \quad (14)$$

Again, these are fixed-point equations for which the existence and uniqueness of a solution is guaranteed by the contraction mapping theorem. Given the optimal action-value function Q_* , an optimal policy is obtained by selecting in each state the action with maximizes Q_*

$$a_* = \arg \sup_a Q_*(s, a) \quad (15)$$

This greedy policy is deterministic and only depends on the current state of the system.

2.2 Policy Gradient Methods

The standard way to solve MDPs is through dynamic programming, which simply consists in solving the Bellman fixed-point equations discussed in the previous chapter. Following this approach, the problem of finding the optimal policy is transformed into the problem of finding the optimal value function. However, apart from the simplest cases where the MDP has a limited number of states and actions, dynamic programming becomes computationally infeasible. Moreover, this approach requires complete knowledge of the Markov transition kernel and of the reward function, which in many real-world applications might be unknown or too complex to use. *Reinforcement Learning* (RL) is a subfield of Machine Learning which aims to turn the infeasible dynamic programming methods into practical algorithms that can be applied to large-scale problems. RL algorithms are based on two key ideas: the first is to use samples to compactly represent the unknown dynamics of the controlled system. The second idea is to use powerful function approximation methods to compactly estimate value functions and policies in high-dimensional state and action spaces. In this section we will only focus on a particular class of algorithms called *Policy Gradient Methods*, which have proved successful in many applications. For a more complete introduction to RL, the reader may consult [2], [3] or [4].

In *policy gradient methods* [5], the optimal policy is approximated using a parametrized policy $\pi : \mathbb{S} \times \mathcal{A} \times \Theta \rightarrow \mathbb{R}$ such that, given a parameter vector $\theta \in \Theta \subseteq \mathbb{R}^{D_\theta}$, $\pi(s, B; \theta) = \pi_\theta(s, B)$ gives the probability of selecting an action in $B \in \mathcal{A}$ when the system is in state $s \in \mathbb{S}$. The general goal of policy optimization in reinforcement learning is to optimize the policy parameters $\theta \in \Theta$ so as to maximize a certain objective function $J : \Theta \rightarrow \mathbb{R}$

$$\theta^* = \arg \max_{\theta \in \Theta} J(\theta) \quad (16)$$

In the following, we will focus on gradient-based and model-free methods that exploit the sequential structure of the reinforcement learning problem. The idea of policy gradient algorithms is to update the policy parameters using the gradient ascent direction of the objective function

$$\theta_{k+1} = \theta_k + \alpha_k \nabla_\theta J(\theta_k) \quad (17)$$

where $\{\alpha_k\}_{k \geq 0}$ is a sequence of learning rates. Typically, the gradient of the objective function is not known and its approximation is the key component of every policy gradient algorithm. It is a well-know result from stochastic optimization [6] that, if the gradient estimate is unbiased and the learning rates satisfy the *Robbins-Monro conditions*

$$\sum_{k=0}^{\infty} \alpha_k = \infty \quad \sum_{k=0}^{\infty} \alpha_k^2 < \infty \quad (18)$$

the learning process is guaranteed to converge at least to a local optimum of the objective function. In an episodic environment where the system always starts from an initial state s_0 , the typical objective function is the start value.

$$J_{\text{start}}(\theta) = V_{\pi_\theta}(s_0) = \mathbb{E}_{\pi_\theta} [G_0 | S_0 = s_0] \quad (19)$$

In a continuing environment, where no terminal state exists and the task might go on forever, it is common to use either the average value

$$J_{\text{avV}}(\theta) = \mathbb{E}_{S \sim d^\theta} [V_{\pi_\theta}(S)] = \int_{\mathbb{S}} d^\theta(s) V_{\pi_\theta}(s) ds \quad (20)$$

where d^θ is the stationary distribution of the Markov chain induced by π_θ . Alternatively, one may use the average reward per time step

$$J_{\text{avR}}(\theta) = \rho(\theta) = \mathbb{E}_{\substack{S \sim d^\theta \\ A \sim \pi_\theta}} [\mathcal{R}(S, A)] = \int_{\mathbb{S}} d^\theta(s) \int_{\mathbb{A}} \pi_\theta(s, a) \mathcal{R}(s, a) da ds \quad (21)$$

Luckily, the same methods apply with minor changes to the three objective functions.

2.2.1 Policy Gradient Theorem

The *policy gradient theorem* [7] shows that the gradient can be rewritten in a form suitable for estimation from experience aided by an approximate action-value or advantage function.

Theorem 2.1 (Policy Gradient). *Let π_θ be a differentiable policy. The policy gradient for the average reward formulation is given by*

$$\nabla_\theta \rho(\theta) = \mathbb{E}_{\substack{S \sim d^\theta \\ A \sim \pi_\theta}} [\nabla_\theta \log \pi_\theta(S, A) Q_\theta(S, A)] \quad (22)$$

where d^θ is the stationary distribution of the Markov chain induced by π_θ . The policy gradient for the start value formulation is given by

$$\nabla_\theta J_{\text{start}}(\theta) = \mathbb{E}_{\substack{S \sim d_\gamma^\theta(s_0, \cdot) \\ A \sim \pi_\theta}} [\nabla_\theta \log \pi_\theta(S, A) Q_\theta(S, A)] \quad (23)$$

where $d_\gamma^\theta(s_0, \cdot)$ is the γ -discounted visiting distribution over states starting from the initial state s_0 and following policy π_θ

$$d_\gamma^\theta(s, x) = \sum_{k=0}^{\infty} \gamma^k \mathcal{P}_\theta^{(k)}(s, x) \quad (24)$$

Algorithm 1 GPOMDP

Input:

- Initial policy parameters $\theta_0 = (\theta_0^1, \dots, \theta_0^{D_\theta})^T$
- Learning rate $\{\alpha_k\}$
- Number of trajectories M

Output: Approximation of the optimal policy $\pi_{\theta^*} \approx \pi_*$

1: Initialize $k = 0$

2: **repeat**

3: Sample M trajectories $h^{(m)} = \{(s_t^{(m)}, a_t^{(m)}, r_{t+1}^{(m)})\}_{t=0}^{T^{(m)}}$ of the MDP under policy π_{θ_k}

4: Compute the optimal baseline

$$\hat{b}_k^n = \frac{\sum_{m=1}^M \left[\sum_{i=0}^{T^{(m)}} \partial_{\theta_k} \log \pi_{\theta} \left(s_i^{(m)}, a_i^{(m)} \right) \right]^2 \sum_{j=0}^{T^{(m)}} \gamma^j r_{j+1}^{(m)}}{\sum_{m=1}^M \left[\sum_{i=0}^{T^{(m)}} \partial_{\theta_k} \log \pi_{\theta} \left(s_i^{(m)}, a_i^{(m)} \right) \right]^2} \quad (27)$$

5: Approximate policy gradient

$$\frac{\partial}{\partial \theta^n} J_{\text{start}}(\theta_k) \approx \hat{g}_k^n = \frac{1}{M} \sum_{m=1}^M \sum_{i=0}^{T^{(m)}} \frac{\partial}{\partial \theta^n} \log \pi_{\theta_k} \left(s_i^{(m)}, a_i^{(m)} \right) \left(\sum_{j=i}^{T^{(m)}} \gamma^j r_{j+1}^{(m)} - \hat{b}_k^n \right) \quad (28)$$

6: Update actor parameters $\theta_{k+1} = \theta_k + \alpha_k \hat{g}_k$.

7: $k \leftarrow k + 1$

8: **until** converged

Let us notice that we can subtract a state-dependent baseline from the action-value function without changing the value of the expectation, indeed

$$\begin{aligned} \mathbb{E}_{\substack{S \sim d^\theta \\ A \sim \pi_\theta}} [\nabla_\theta \log \pi_\theta(S, A) B_\theta(S)] &= \int_{\mathbb{S}} d^\theta(s) \int_{\mathbb{A}} \pi_\theta(s, a) \nabla_\theta \log \pi_\theta(s, a) B_\theta(s) da ds \\ &= \int_{\mathbb{S}} d^\theta(s) B_\theta(s) \int_{\mathbb{A}} \nabla_\theta \pi_\theta(s, a) da ds \\ &= \int_{\mathbb{S}} d^\theta(s) B_\theta(s) \nabla_\theta \underbrace{\int_{\mathbb{A}} \pi_\theta(s, a) da}_{=1} ds = 0 \end{aligned}$$

Hence, the policy gradient theorem can be rewritten as

$$\nabla_\theta \rho(\theta) = \mathbb{E}_{\substack{S \sim d^\theta \\ A \sim \pi_\theta}} [\nabla_\theta \log \pi_\theta(S, A) (Q_{\pi_\theta}(S, A) - B_\theta(S))] \quad (25)$$

The baseline can be chosen so as to minimize the variance of the gradient estimate which can prove beneficial for the algorithm convergence [5]. This result can be used as the starting point to derive several policy gradient methods that use different approximation of the action-value function, which is typically unknown. For instance, in an episodic MDP the action-value function can be estimated with the total return obtained on a sample trajectory

$$Q_\theta(s_0, a_0) \approx \sum_{t=0}^{T^{(m)}} \gamma^t r_{t+1}^{(m)} \quad (26)$$

Combining this remark with a Monte Carlo approximation of Eq. (25), we obtain the *Monte Carlo Policy Gradient* algorithm [8] (also known as GPOMDP) for which the pseudocode is reported in Algorithm 1.

2.2.2 Parameter-Based Policy Gradient Methods

In Monte Carlo Policy Gradient, trajectories are generated by sampling at each time step an action according to a stochastic policy π_θ and the objective function gradient is estimated by differentiating

the policy with respect to the parameters. However, sampling an action from the policy at each time step leads to a large variance in the sampled histories and therefore in the gradient estimate, which can in turn slow down the convergence of the learning process. To address this issue, the *policy gradient with parameter-based exploration* (PGPE) method [9] replaces the search in the policy space with a direct search in the model parameter space. Given an episodic MDP, PGPE considers a deterministic controller $F : \mathbb{S} \times \Theta \rightarrow \mathbb{A}$ that, given a set of parameters $\theta \in \Theta \subseteq \mathbb{R}^{D_\theta}$, maps a state $s \in \mathbb{S}$ to an action $a = F(s; \theta) = F_\theta(s) \in \mathbb{A}$. The policy parameters are drawn from a probability distribution p_ξ , with hyper-parameters $\xi \in \Xi \subseteq \mathbb{R}^{D_\xi}$. Combining these two hypotheses, the agent follows a stochastic policy π_ξ defined by

$$\forall B \in \mathcal{A}, \pi_\xi(s, B) = \pi(s, B; \xi) = \int_{\Theta} p_\xi(\theta) \mathbb{1}_{F_\theta(s) \in B} d\theta \quad (29)$$

In this setting, the policy gradient theorem can be reformulated in the following way

Theorem 2.2 (Parameter-Based Policy Gradient). *Let p_ξ be differentiable with respect to ξ , then the gradient of the average reward is given by*

$$\nabla_\xi J(\xi) = \mathbb{E}_{\substack{S \sim d^\xi \\ \theta \sim p_\xi}} [\nabla_\xi \log p_\xi(\theta) Q_{\pi_\xi}(S, \theta)] \quad (30)$$

where we denoted $Q_\xi(S, \theta) = Q_\xi(S, F_\theta(S))$.

This expression is very similar to the original policy gradient theorem, but the expectation is taken over the controller parameters instead of the action space and we have the likelihood score of the controller parameters distribution instead of that of the stochastic policy. Thus, we might interpret this result as if the agent directly selected the parameters θ according to a policy p_ξ , which then lead to an action through the deterministic mapping F_θ . Therefore, it is as if the agent's policy was in the parameters space and not in the control space. As in the standard policy gradient methods, we can subtract a state-dependent baseline $B_\xi(S)$ to the gradient without increasing the bias

$$\nabla_\xi J(\xi) = \mathbb{E} [\nabla_\xi \log p_\xi(\theta) (Q_{\pi_\xi}(S, \theta) - B_\xi(S))] \quad (31)$$

The PGPE algorithm, which is outlined in Algorithm 2, employs a Monte Carlo approximation of this gradient, where the action-value function is estimated using the returns on a sampled trajectory of the MDP. The benefit of this approach is that the controller is deterministic and therefore the actions do not need to be sampled at each time step, with a consequent reduction of the gradient estimate variance. Indeed, It is sufficient to sample the parameters θ once at the beginning of the episode and then generate an entire trajectory following the deterministic policy F_θ . As an additional benefit, the parameter gradient is estimated by direct parameter perturbations, without having to backpropagate any derivatives, which allows to use non-differentiable controllers. Again the baseline can be chosen so as to minimize the gradient estimate variance [10].

3 Reinforcement Learning for Automated Trading

Many financial applications can be seen as sequential decision problems which naturally fall in the stochastic optimal control framework introduced above. In this section we discuss how reinforcement learning algorithms can be applied to the asset allocation problem, where an agent invests his capital on various assets available in the market.

3.1 Asset Allocation With Transaction Costs

The asset allocation problem consists of determining how to dynamically invest the available capital in a portfolio of different assets in order to maximize the expected total return or another relevant performance measure. Let us consider a financial market consisting of $I + 1$ different stocks that are traded only at discrete times $t \in \{0, 1, 2, \dots\}$ and denote by $Z_t = (Z_t^0, Z_t^1, \dots, Z_t^I)^T$ their prices at time t . Typically, Z_t^0 refers to a riskless asset whose dynamic is given by $Z_t^0 = (1 + X)^t$ where X is the deterministic risk-free interest rate. The investment process works as follows: at time t , the investor observes the state of the market S_t , consisting for example of the past asset prices and other relevant economic variables, and subsequently chooses how to rebalance his portfolio, by specifying the units of each stock $n_t = (n_t^0, n_t^1, \dots, n_t^I)^T$ to be held between t and $t + 1$. In doing so, he needs

Algorithm 2 Episodic PGPE algorithm

Input:

- Initial hyper-parameters $\xi_0 = (\xi_0^1, \dots, \xi_0^{D_\xi})^T$
- Learning rate $\{\alpha_k\}$
- Number of trajectories M

Output: Approximation of the optimal policy $F_{\xi^*} \approx \pi_*$

- 1: Initialize $k = 0$
- 2: **repeat**
- 3: **for** $m = 1, \dots, M$ **do**
- 4: Sample controller parameters $\theta^{(m)} \sim p_{\xi_k}$
- 5: Sample trajectory $h^{(m)} = \{(s_t^{(m)}, a_t^{(m)}, r_{t+1}^{(m)})\}_{t=0}^{T^{(m)}}$ under policy $F_{\theta^{(m)}}$
- 6: **end for**
- 7: Compute optimal baseline

$$\hat{b}_k^n = \frac{\sum_{m=1}^M [\partial_{\xi^n} \log p_{\xi_k}(\theta^{(m)})]^2 \sum_{j=0}^{T^{(m)}} \gamma^j r_{j+1}^{(m)}}{\sum_{m=1}^M [\partial_{\xi^n} \log p_{\xi_k}(\theta^{(m)})]^2} \quad (32)$$

- 8: Approximate policy gradient

$$\frac{\partial}{\partial \xi^n} J_{\text{start}}(\xi_k) \approx \hat{g}_k^n = \frac{1}{M} \sum_{m=1}^M \frac{\partial}{\partial \xi^n} \log p_{\xi_k}(\theta^{(m)}) \left(\sum_{j=0}^{T^{(m)}} \gamma^j r_{j+1}^{(m)} - \hat{b}_k^n \right) \quad (33)$$

- 9: Update hyperparameters using gradient ascent $\xi_{k+1} = \xi_k + \alpha_k \hat{g}_k^n$
 - 10: $k \leftarrow k + 1$
 - 11: **until** converged
-

to take into account the transaction costs that he has to pay to the broker to change his position. At time $t + 1$, the investor realizes a profit or a loss from his investment due to the stochastic variation of the stock values. The investor's goal is to maximize a given performance measure. Let W_t denote the wealth of the investor at time t . The profit realized between t and $t + 1$ is simply given by the difference between the trading results and the transaction costs paid to the broker. More formally

$$\Delta W_{t+1} = W_{t+1} - W_t = \text{PNL}_{t+1} - \text{TC}_t$$

where PNL_{t+1} denotes the profit due to the variation of the portfolio asset prices between t and $t + 1$

$$\text{PNL}_{t+1} = n_t \cdot \Delta Z_{t+1} = \sum_{i=0}^I n_t^i (Z_{t+1}^i - Z_t^i)$$

and TC_t denotes the fees paid to the broker to change the portfolio allocation and on the short positions

$$\text{TC}_t = \sum_{i=0}^I \delta_p^i |n_t^i - n_{t-1}^i| Z_t^i - \delta_f W_t \mathbb{1}_{n_t \neq n_{t-1}} - \sum_{i=0}^I \delta_s^i (n_t^i)^- Z_t^i$$

The transaction costs consist of three different components. The first term represent a transaction cost that is proportional to the change in value of the position in each asset. The second term is a fixed fraction of the total value of the portfolio which is paid only if the allocation is changed. The last term represents the fees paid to the broker for the shares borrowed to build a short position. The portfolio return between t and $t + 1$ is thus given by

$$X_{t+1} = \frac{\Delta W_{t+1}}{W_t} = \sum_{i=0}^I \left[a_t^i X_{t+1}^i - \delta_i |a_t^i - \tilde{a}_t^i| - \delta_s (a_t^i)^- \right] - \delta_f \mathbb{1}_{a_t \neq \tilde{a}_t} \quad (34)$$

where

$$X_{t+1}^i = \frac{\Delta Z_{t+1}^i}{Z_t^i}$$

is the return of the i -th stock between t and $t + 1$,

$$a_t^i = \frac{n_t^i Z_t^i}{W_t}$$

is the fraction of wealth invested in the i -th stock between time t and $t + 1$ and finally

$$\tilde{a}_t^i = \frac{n_{t-1}^i Z_t^i}{W_t} = \frac{a_{t-1}^i (1 + X_t^i)}{1 + X_t}$$

is the fraction of wealth invested in the i -th stock just before the reallocation. We assume that the agent invests all his wealth at each step, so that W_t can be also interpreted as the value of his portfolio. This assumption leads to the following constraint on the portfolio weights

$$\sum_{i=0}^I a_t^i = 1 \quad \forall t \in \{0, 1, 2, \dots\} \quad (35)$$

We notice that we are neglecting the typical margin requirements on the short positions, which would reduce the available capital at time t . Considering margin requirements would lead to a more complex constraint on the portfolio weights which would be difficult to treat in the reinforcement learning framework. Plugging this constraint into Eq. (34), we obtain

$$X_{t+1} = X + \sum_{i=1}^I a_t^i (X_{t+1}^i - X) - \sum_{i=0}^I \left[\delta_i |a_t^i - \tilde{a}_t^i| - \delta_s^i (a_t^i)^- \right] - \delta_f \mathbb{1}_{a_t \neq \tilde{a}_{t-1}} \quad (36)$$

which highlights the role of the risk-free asset as a benchmark for the portfolio returns. The total profit realized by the investor between $t = 0$ and T is

$$\Pi_T = W_T - W_0 = \sum_{t=1}^T \Delta W_t = \sum_{t=1}^T W_t X_t$$

The portfolio return between $t = 0$ and T is given by

$$X_{0,T} = \frac{W_T}{W_0} - 1 = \prod_{t=1}^T (1 + X_t) - 1$$

In order to cast the asset allocation problem in the reinforcement learning framework, we consider the log-return of the portfolio between $t = 0$ and T

$$R_{0,T} = \log \frac{W_T}{W_0} = \sum_{t=1}^T \log(1 + X_t) = \sum_{t=1}^T R_t \quad (37)$$

where R_{t+1} is the log-return of the portfolio between t and $t + 1$

$$R_{t+1} = \log \left\{ 1 + \sum_{i=0}^I \left[a_t^i X_{t+1}^i - \delta_i |a_t^i - \tilde{a}_t^i| - \delta_s^i (a_t^i)^- \right] - \delta_f \mathbb{1}_{a_t \neq \tilde{a}_{t-1}} \right\} \quad (38)$$

The portfolio log-return can be used as the reward function of a RL algorithm, either in a offline or in an online approach.

3.2 Reinforcement Learning Application

In the previous section we derived the reward function for the asset allocation problem with transaction costs. In order to apply the policy gradient algorithms discussed in the previous sections we still need to define the state space, the action space and the agent's policy. For simplicity, we limit ourselves to the case of a single risky asset, i.e. $I = 1$, but the discussion could be generalized to the multi-asset case.

We assume that at each time step the agent considers the $P + 1$ past returns of the risky asset, i.e. $\{X_t, X_{t-1}, \dots, X_{t-P}\}$. In order to properly incorporate the effects of transaction costs into his




Figure 2: PyBrain standard architecture for an RL problem.

decision process, the agent must keep track of its current position \tilde{a}_t . The state of the system is thus given by

$$S_t = \{X_t, X_{t-1}, \dots, X_{t-P}, \tilde{a}_t\} \quad (39)$$

We might also include some external variables Y_t that may be relevant to the trader, such as the common technical indicator used in practice. Furthermore, these input variables may be used to construct more complex features for example using some deep learning techniques, such as a deep auto-encoder.

The agent, or trading system, specifies the portfolio weights $a_t = (a_t^0, a_t^1)^T$ according to a long-short strategy, i.e. the agent may be long ($a_t^1 = +1$) or short ($a_t^1 = -1$) on the risky-asset while $a_t^0 = 1 - a_t^1$ since the agents invests all the available capital at each time step. In the GPOMDP framework we assume that the agent selects a_t^1 according to a Boltzmann policy, i.e.

$$\pi_\theta(s, +1) = \frac{e^{\theta^T s}}{1 + e^{\theta^T s}} \quad \pi_\theta(s, -1) = \frac{1}{1 + e^{\theta^T s}} \quad (40)$$

where we included a bias term in the parameters and in the state. In the parameter-based formulation, we assume that agent selects actions according to the binary controller

$$F_\theta(s) = \text{sign}(\theta^T s) \quad (41)$$

where the controller parameters are normally distributed $\theta \sim \mathcal{N}(\mu, \text{diag}(\sigma))$.

4 Python Prototype

In this section, we start discussing the implementation details of this project. The first step of this project has been to implement a prototype in Python, a high-level, general-purpose, interpreted, dynamic programming language which is gaining a widespread popularity both in the academic world and in the industry. Python natively supports the object-oriented paradigm which makes it perfect to quickly develop a prototype of the class architecture, which can then be translated in C++. Moreover, thanks to external libraries such as Numpy, Scipy and Pandas, Python offers an open-source alternative to Matlab for scientific computing applications.

For the basic RL algorithms we exploited PyBrain¹, a modular ML library for Python whose goal is to offer flexible, easy-to-use yet still powerful algorithms for ML tasks and a variety of predefined environments to test and compare different algorithms [11]. An RL task in PyBrain always consists of an Environment, an Agent, a Task and an Experiment interacting with each other as illustrated in Figure 5.

The Environment is the world in which the Agent acts and is characterized by a state

¹<http://pybrain.org/>

which can be accessed through the `getSensors()` method. The Agent receives an observation of this state through the `integrateObservation()` method and selects an action through the `getAction()` method. This action is applied to the Environment with the `performAction()` method. However, the interactions between the Environment and the Agent are not direct but are mediated by the Task. The Task specifies what the goal is in an Environment and how the agent is rewarded for its actions. Hence, the composition of an Environment and a Task fully defines the MDP. An Agent always contains an Actor, which represents the policy used to select actions. Based on the rewards that the Agent receives via the `getReward()` method, the Learner improves the policy via a `learn()` procedure. In this step, an Actor may be used to evaluate a state with the goal of reducing the variance of the learning process. This entire learning process is controlled by an Experiment object.

This structure is quite standard for a RL problem and can be easily adapted to the problem at hand and extended to the learning algorithms developed in this thesis. Based on this architecture, we thus developed a fully-working Python prototype of the asset allocation problem. This prototype yielded some interesting results both on simulated data and on historical data, in particular for the PGPE algorithm. However, the learning process resulted too slow to be run systematically for a large number of time-steps and training epochs. By consequent, we quickly decided to pass to C++.

5 C++ Implementation

Passing from Python to C++ presents some challenges in the design of a suitable architecture for the RL algorithms discussed above. Following the approach of [12], our main goal has been code reusability, which is based on the important attributes of clarity and elegance of design. In addition, we always kept in mind the possibility the our original design might need to be extended in the future. In some cases we thus favored extendability compared to efficiency. First we describe the C++ adaptation of the PyBrain’s Environment, Task, Agent and Experiment objects with a particular attention to their concrete implementations for the asset allocation problem. Secondly, we discuss the design for an Average-Reward Actor-Critic agent (ARAC), which provides a concrete implementation of the Agent interface and can be used to solve the asset allocation problem. Here we only give an overview of the program, addressing the reader to the full documentation which can be found at `Code/Thesis/doc/doc.pdf` for a thorough explanation of all the classes and their methods.

5.1 Environment, Task, Agent and Experiment

Figure 3 schematically represents the design of the base components of an RL application, which closely replicates Pybrain’s architecture. The pure abstract classes Environment, Task, Agent and Experiment define the generic interfaces to which all the concrete implementations of these objects must adhere. To achieve code modularity, we make most of the objects in our design clonable in order to allow for the polymorphic composition of classes. Exploiting this design pattern, we couple a Task with an Environment by storing a `std::unique_ptr<Environment>` as a private member of the class. Similarly, an Experiment is coupled via composition with a Task and an Agent. The methods of these classes are similar to those in Pybrain. For all the linear algebra operations we decided to use Armadillo², a high quality linear algebra library which provides high-level syntax (API) deliberately similar to Numpy and Matlab aiming towards a good balance between speed and ease of use. Therefore, the state of the system and the actions performed by the agent are represented as `arma::vector` objects. Let us now present the concrete implementation of these objects for the asset allocation task. MarketEnvironment implements a financial market by storing the historical daily returns of the risky assets in an `arma::matrix`. These values are read from a given input .csv file, which is generated automatically running a Python script and which either contains real historical returns downloaded from Yahoo Finance³ or synthetic returns generated according to a certain stochastic process. Therefore, we always work on samples of the system without making any assumption on its Markov transition kernel.

²<http://arma.sourceforge.net/>

³<https://uk.finance.yahoo.com/>

The `AssetAllocationTask` implements the asset allocation MDP discussed in Section 3 by providing a method to compute the reward that the `Agent` obtains from investing on the risky assets. Moreover, the `AssetAllocationTask` enlarges the system state so that the `Agent` also observes the past P states and the current portfolio weights. The `AssetAllocationExperiment` handles the interactions between the `Agent` and the `Task`. The learning process is divided in two phases: the training phase consists of a certain number of learning epochs over the same time period during which the `Agent` improves the parameters of its policy via the `learn` method. Some estimates of the objective function are dumped in the `StatisticsExperiment` object and are used in the post-processing phase to plot the learning curves of the algorithm. In the backtest phase, the `Agent` applies the learned policy on the time period which follows the one used for training and the relevant performance measures are stored in the `BacktestLog` for successive analysis and comparison between different learning algorithms.

5.2 ARACAgent


A concrete implementation of the `Agent` interface completes the standard structure of an RL task. In this section we discuss the design for an Average Reward Actor-Critic agent (ARAC), which includes most of the features of the other algorithms tested in this thesis. The full architecture of this agent is illustrated in Figure 4, but for brevity we will only focus on the more interesting aspects.

The `ARACAgent` builds upon a `StochasticActor` and a `Critic` via composition. A `StochasticActor` is simply a wrapper around a `StochasticPolicy`, a pure abstract class defining the generic interface for a stochastic policy used by an agent to select actions. In addition to `getAction` and `get/set` methods for the policy parameters, a `StochasticPolicy` must implement the `likelihoodScore` method which computes the likelihood score $\nabla_{\theta} \log \pi_{\theta}(s, a)$ for a given state and action and which plays a crucial role in any policy gradient algorithm.

We provide two examples of concrete implementations of the `StochasticPolicy`. The first example is the `BoltzmannPolicy` typically used in discrete action spaces. The implementation of this policy is quite straightforward and we won't discuss the details. The second and more interesting stochastic policy implemented is the `PGPEPolicy`. This class is based on the decorator design pattern which is typically used to extend the interface of a certain class. Indeed, the `PGPEPolicy` is a `StochasticPolicy` which contains by polymorphic composition a `Policy`, a pure abstract class which provides the generic interface for a policy, potentially deterministic. This `Policy` object represents the deterministic controller F_{θ} used in the PGPE algorithm. Moreover, the `PGPEPolicy` contains by polymorphic composition a `ProbabilityDistribution`, a pure abstract class defining the generic interface for a probability distribution. This probability distribution represents the hyper-distribution p_{ξ} on the controller parameters. In order to be used in a PGPE algorithm, a `ProbabilityDistribution` must implement a `likelihoodScore` method to compute the likelihood score of the hyper-distribution $\nabla_{\xi} \log p_{\xi}(\theta)$. Hence, the `likelihoodScore` method of `PGPEPolicy` simply broadcasts the call to the `likelihoodScore` method of its underlying `ProbabilityDistribution`. A concrete implementation of a `ProbabilityDistribution` is provided by the `GaussianDistribution`, which implements a multi-variate and axis-aligned normal distribution $\mathcal{N}(\mu, \text{diag}(\sigma))$.

The objects discussed so far are sufficient to implement an actor-only learning algorithm, potentially using a baseline to evaluate the rewards. A more advanced variance reduction technique consists in using a `Critic`, which approximates the value function and provides an evaluation of a given state. The `Critic` class is simply a wrapper around a `FunctionApproximator` which provides a generic interface for a parametric function $B_{\omega}(s)$. The key methods of this class are `evaluate`, which evaluates the approximator at a given point, and `gradient`, which computes its gradient at a given point.

Finally, the `ARACAgent` employs some `LearningRate` to control the speed of the gradient descent optimization algorithm. A naive approach is to use a `ConstantLearningRate`, but this leads to a large-variance in the objective function value attained by the stochastic optimization algorithm. A more sensible choice is to use a `DecayingLearningRate` which decreases with the number of learning epochs performed by the agent according to $\alpha_n = \frac{a}{n^b}$. In this way, the learning process progressively “cools down” (using a simulated annealing terminology) and stabilizes to a given policy. This concludes our quick overview of the class architecture used for this project. In the thesis, other applications and algorithms were considered and we refer the reader to the full document for a more complete discussion.



6 Execution Pipeline

In this section we describe the full pipeline of the program, which is schematically represented in Figure 6. This pipeline allows to run the learning algorithm for the asset allocation problem and automatically determine a trading strategy. The execution consists of the following steps.

Compilation To build the `thesis` library it is sufficient to run the `Makefile` generated with `cmake`. This produces two executables: `main` which is used to debug the program in the `Code::Blocks` IDE and `main_thesis` which is used to run the experiment in the full execution pipeline.

generate_synthetic_series.py This Python script simulates the returns of a synthetic asset and prints them on a `.csv` file which is then read by `main_thesis` and used to initialize the `MarketEnvironment` object. Alternatively, `market_data_collector.py` collects the historical returns for a list of given assets from Yahoo finance.

experiment_launcher.py This Python script manages the execution pipeline. First, the experiment parameters are specified and dumped in a `.pot` file which is then read by `main_thesis`. Given the parameter values, the script determines the folders where the output should be written so that the results can be easily associated to a specific set of parameters. Subsequently, it launches the `main_thesis` executable passing the correct parameters via the command line. Finally it runs the `postprocessing.py` scripts which processes the output files.

main_thesis This executable takes some inputs from the command line, such as the algorithm to use, the paths to the input files and the paths where the output files should be generated. The experiment parameters are then read from the `.pot` file generated by the `experiment_launcher.py` using `GetPot`. The type of learning algorithm is specified by a string passed to the executable via command line and then used by the factory `FactoryOfAgents` to instantiate the corresponding `Agent`. When the `AssetAllocationExperiment` is run, it outputs various statistics to the given destination folders. More in detail, it prints two files for every independent run of the experiment: a `debug.csv` file which contains the learning curves of the algorithm and an `output.csv` file which contains the backtest performance measures for the trading strategy learned by the `Agent` during training.

postprocessing.py This Python script processes the various files produced by `main_thesis` and generates an aggregate analysis of the various learning algorithms, so that they can be easily compared and assessed. In particular, it computes the average and confidence intervals for the learning curves of the algorithms and the backtest cumulative profits of the learned strategies. Moreover, it computes some performance measures typically used in Finance to evaluate a trading strategy, such as the Sharpe ratio and the maximum drawdown. The results of this analysis are stored in some `.csv` files and some images are generated using the Python library `matplotlib`.

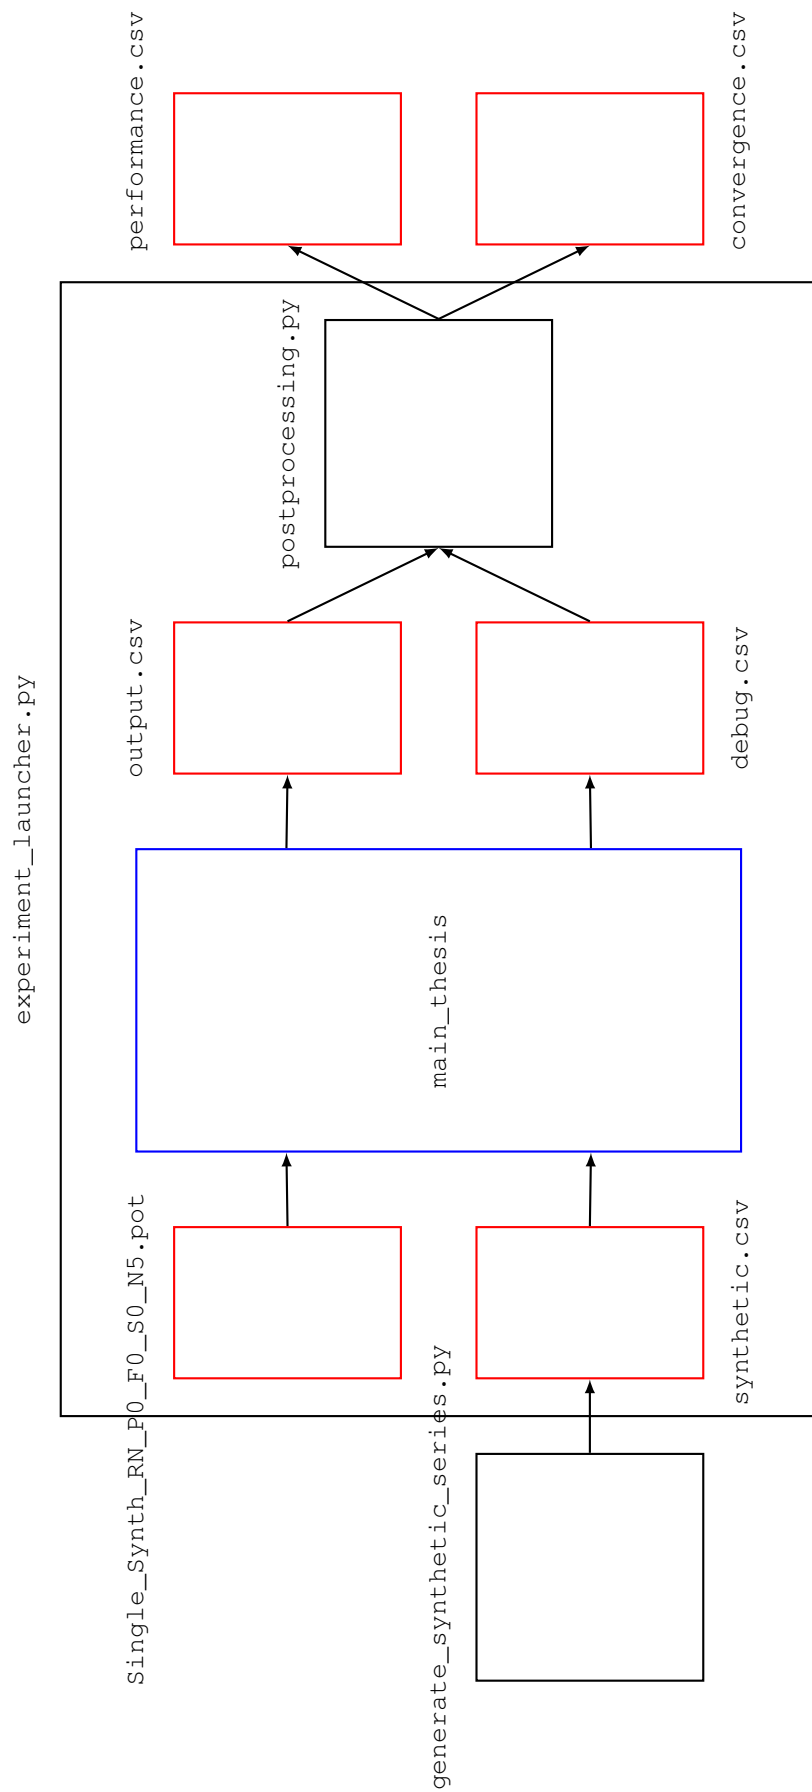


Figure 5: Execution flow of an asset allocation experiment. Black boxes denote Python scripts, blue boxes executables while red boxes input/output files.

7 Numerical Results

In this section we present the numerical results of the online version of the policy gradient algorithms discussed in Section 2 for the asset allocation problem.

7.1 Synthetic Asset

To assess the different reinforcement learning methods in a controlled environment, the algorithms have been tested on a synthetic asset whose behavior presents some features that can be traded profitably. We simulated the log-price series $\{z_t\}$ for the risky asset as a random walk with autoregressive trend $\{\beta_t\}$. The two-parameter model is thus given by

$$\begin{aligned} z_t &= z_{t-1} + \beta_{t-1} + \kappa \epsilon_t \\ \beta_t &= \alpha \beta_{t-1} + \nu_t \end{aligned} \quad (42)$$

We then define the synthetic price series as

$$Z_t = \exp \left(\frac{z_t}{\max_t z_t - \min_t z_t} \right) \quad (43)$$

This model is often taken as a benchmark test in the automated trading literature, see for instance [13]. In addition to presenting some exploitable patterns, the model is stationary and therefore the policy learned on the training set should generalize well on the test set, also known as backtest in the financial jargon. We would thus expect our learning algorithms to perform well on this test case.

7.2 Experimental Setup

All the algorithms were tested on the same price series of size 9000, generated from the process above using $\alpha = 0.9$ and $\kappa = 3$. The learning process consisted of 500 training epochs on the first 7000 days of the series with a learning rate that decreased at each epoch according to a polynomial schedule. The trained agents were subsequently backtested on the final 2000 days, during which the agents kept learning online in order to try to adapt to the changing environment. The results that we present are the average of 10 independent experiments that used slightly different random initialization of the policy parameters.

7.3 Convergence

Let us first discuss the case with no transaction costs. Figure 6 shows the learning curves three algorithms in terms of average daily reward, which is the quantity being maximized by the algorithms, the daily reward standard deviation and the annualized Sharpe ratio. The NPGPE algorithm is an enhancement of the PGPE algorithm based on the natural gradient technique [14]. The first thing we observe is the ARAC algorithm seems not to be improving the trading strategy as the training epochs go by. The average reward obtained is close to zero and will be surely be negative once transaction costs are introduced. On the other hand, NPGPE slowly converges to a profitable strategy which is however suboptimal compared to the one found by PGPE, that is better in all three measures considered. It is interesting to notice that PGPE and NPGPE yield a learning curve for the Sharpe ratio very similar to the one for the average reward. Even if the algorithm is risk-neutral, it manages to improve a risk-sensitive measure at the same time of the average reward. This might be simply a peculiarity of the very simple model assumed for the synthetic risky asset. Moreover, since the price process is stationary, the trading strategy learned on the training set generalizes well to the test set.

7.4 Performances

Figure 7 compares the backtest performances of the three learned policies and a Buy and Hold strategy, which simply consists in investing all the available capital in the risky asset. Let us repeat that the solid lines are the averages of 10 independent experiments, which allows us to determine the 95% confidence intervals represented with the dashed lines. We clearly see that NPGPE and PGPE consistently outperform the market, realizing a total profit of 231.63% and 314.34% respectively against the 7.81% profit of the Buy and Hold strategy over the same period. More statistics of the trading strategies are reported in Table 1.




Figure 6: Risk-neutral learning process for the asset allocation problem with one synthetic risky asset.




Figure 7: Backtest performance of trained trading systems for the asset allocation problem with one synthetic risky asset.

	Buy and Hold	ARAC	NPGPE	PGPE
Total Return	7.81%	-0.86%	231.63%	314.34%
Daily Sharpe	0.27	-0.02	4.13	4.95
Monthly Sharpe	0.19	-0.07	2.90	3.26
Yearly Sharpe	0.23	-0.10	1.55	1.76
Max Drawdown	-22.35%	-12.60%	-3.72%	-3.27%
Avg Drawdown	-1.75%	-1.81%	-0.49%	-0.43%
Avg Up Month	2.87%	1.14%	2.47%	2.74%
Avg Down Month	-2.58%	-1.10%	-0.73%	-0.67%
Win Year %	40.00%	44.00%	98.00%	100.00%
Win 12m %	56.36%	48.00%	100.00%	100.00%
Reallocation Freq	0.00%	50.01%	19.99%	15.43%
Short Freq	0.00%	50.13%	41.59%	44.25%

Table 1: Backtest statistics of the risk-neutral trading strategies for the asset allocation problem with one synthetic risky asset.




Figure 8: Impact of proportional transaction costs on the trading strategies learned by PGPE and NPGPE.

7.5 Impact of Transaction Costs

In the algorithmic trading literature there are many examples of strategies based on the prediction of future returns based on more or less complex indicators [15], [16], [17]. However, as pointed out in [18], the performances of these methods quickly degrade when transaction costs for changing the portfolio composition or for shorting a security are considered. Indeed, these methods simply invest based on the prediction of the future returns, without explicitly taking into account transaction costs. On the other hand, reinforcement learning algorithms should learn to avoid frequent reallocations or shorts thanks to the feedback mechanism between the learning agent and the system, thus generating better trading performances. In this section we analyze how the strategies learned by PGPE and by NPGPE change when gradually increasing the proportional transaction costs and the short-selling fees. Intuitively, we expect a progressive reduction of the frequency of reallocation and of shorting the risky asset.

Figure 8 shows the impact of proportional transaction costs on the trading strategies learned by PGPE and by NPGPE. As expected, the frequency of reallocation for both strategies quickly drops to zero as the transaction costs increase, converging to the profitable buy and hold strategy. It is peculiar that the reallocation frequency for the PGPE strategy initially drops more quickly than for the NPGPE strategy, but then slows down and even increases when $\delta_P = 20$ bps. In summary, both algorithms are able to identify reallocation as the cause for lower rewards and to subsequently reduce the rate of reallocation, converging towards the simple yet profitable buy and hold strategy. Figure 9 shows the impact of short-selling fees on the trading strategies learned by PGPE and NPGPE. Both algorithms behave as expected, displaying a progressive reduction of the frequency of short positions as the fees increase. For large values of short-selling fees, both strategies converge to the profitable buy and hold strategy, which completely avoids paying the fees. In particular, PGPE quickly replicates the buy and hold strategy. On the other hand, NPGPE is not able to exactly reproduce the buy and hold strategy but it seems to converge to it for very large values of the short-selling fee.

8 Conclusion

Reinforcement learning is a general class of algorithms that allow an agent to learn how to behave in a stochastic and possibly unknown environment simply by trial-and-error. In this paper, after recalling the basic concepts of RL, we presented an application of some state-of-the-art learning algorithms to the classical financial problem of determining a profitable long-short trading strategy. For a synthetic asset, the strategies produced outperform the simple Buy & Hold strategy, even when investment decisions are only based on extremely basic autoregressive features. Contrarily to standard prediction-based trading systems, the learning algorithms employed are able to adapt as expected to the introduction of transaction costs by reducing the frequency of reallocation and




Figure 9: Impact of short-selling fees on the trading strategies learned by PGPE and NPGPE.

of short positions. This shows the potential of RL techniques in effectively dealing with complex sequential decision problems that are typical in financial applications.

The algorithms above have also been tested on historical data but we couldn't find a trading strategy that consistently outperformed the simple Buy & Hold strategy. This was partly expected as it is extremely difficult to find profitable daily patterns in highly liquid assets, especially when using such simple features. We believe that developing more complex features for the trading strategy, perhaps employing some deep learning feature extraction techniques such as deep auto-encoders or deep (recurrent) neural networks, would allow us to beat the market. Another possibility we considered is to decrease the sampling frequency of the historical data and let the algorithms look for some profitable patterns. At the moment, these ideas are still work in progress.

Acknowledgments

This project has been realized under the supervision of Prof. Marcello Restelli and of Prof. Carlo Sgarra and I would like to thank them for their support. I would also like to thank Matteo Pirota for his advice on some implementation issues.

References

- [1] David Silver, Aja Huang, Chris J Maddison, Arthur Guez, Laurent Sifre, George Van Den Driessche, Julian Schrittwieser, Ioannis Antonoglou, Veda Panneershelvam, Marc Lanctot, et al. Mastering the game of go with deep neural networks and tree search. *Nature*, 529(7587):484–489, 2016.
- [2] Richard S Sutton and Andrew G Barto. *Introduction to reinforcement learning*, volume 135. MIT Press Cambridge, 1998.
- [3] Csaba Szepesvári. Algorithms for reinforcement learning. *Synthesis lectures on artificial intelligence and machine learning*, 4(1):1–103, 2010.
- [4] Marco Wiering and Martijn Van Otterlo. *Reinforcement Learning: State-of-the-Art*, volume 12 of *Adaptation, Learning, and Optimization*. Springer, 1 edition, 2012.
- [5] Jan Peters and Stefan Schaal. Reinforcement learning of motor skills with policy gradients. *Neural networks*, 21(4):682–697, 2008.
- [6] Harold Kushner and G George Yin. *Stochastic approximation and recursive algorithms and applications*, volume 35. Springer Science & Business Media, 2003.
- [7] Richard S Sutton, David A McAllester, Satinder P Singh, Yishay Mansour, et al. Policy gradient methods for reinforcement learning with function approximation. In *NIPS*, volume 99, pages 1057–1063, 1999.
- [8] Jonathan Baxter and Peter L. Bartlett. Infinite-horizon policy-gradient estimation. *Journal of Artificial Intelligence Research*, 15:319–350, 2001.
- [9] Frank Sehnke, Christian Osendorfer, Thomas Rückstieß, Alex Graves, Jan Peters, and Jürgen Schmidhuber. Policy gradients with parameter-based exploration for control. In *Artificial Neural Networks-ICANN 2008*, pages 387–396. Springer, 2008.
- [10] Tingting Zhao, Hirotaka Hachiya, Gang Niu, and Masashi Sugiyama. Analysis and improvement of policy gradient estimation. In *Advances in Neural Information Processing Systems*, pages 262–270, 2011.
- [11] Tom Schaul, Justin Bayer, Daan Wierstra, Yi Sun, Martin Felder, Frank Sehnke, Thomas Rückstieß, and Jürgen Schmidhuber. PyBrain. *Journal of Machine Learning Research*, 11:743–746, 2010.
- [12] Mark Suresh Joshi. *C++ design patterns and derivatives pricing*, volume 2. Cambridge University Press, 2008.
- [13] John Moody, L. Wu, Y. Liao, and M. Saffell. Performance functions and reinforcement learning for trading systems and portfolios. *Journal of Forecasting*, 17:441–470, 1998.
- [14] Atsushi Miyamae, Yuichi Nagata, Isao Ono, and Shigenobu Kobayashi. Natural policy gradient methods with parameter-based exploration for control tasks. In *Advances in neural information processing systems*, pages 1660–1668, 2010.
- [15] Ken’ichi Kamijo and Tetsuji Tanigawa. Stock price pattern recognition-a recurrent neural network approach. In *IEEE 1990 IJCNN International Joint Conference on Neural Networks*, 1990.
- [16] Emad W. Saad, Danil V. Prokhorov, and Donald C. Wunsch. Comparative study of stock trend prediction using time delay, recurrent and probabilistic neural networks. *IEEE Transactions on Neural Networks*, 9(6):1456–1470, 1998.
- [17] Jiuzhen Liang, Wei Song, and Mei Wang. Stock price prediction based on procedural neural networks. *Adv. Artif. Neu. Sys.*, 2011:1–11, 2011.
- [18] Yue Deng, Feng Bao, Youyong Kong, Zhiquan Ren, and Qionghai Dai. Deep direct reinforcement learning for financial signal representation and trading. *IEEE Transactions on Neural Networks and Learning Systems*, 2016.