



POLITECNICO

MILANO 1863

Scuola di Ingegneria Industriale e dell'Informazione
Corso di Laurea Magistrale in Ingegneria Matematica

POLICY GRADIENT ALGORITHMS FOR THE ASSET ALLOCATION PROBLEM

Relatore: Prof. Carlo SGARRA
Correlatore: Prof. Marcello RESTELLI

Tesi di Laurea di:
Pierpaolo NECCHI
Matr. 837664

Anno Accademico 2015-2016

Abstract

The impact of Automated Trading Systems (ATS) on financial markets is growing every year and the trades generated by an algorithm now account for the majority of orders that arrive at stock exchanges. Historically, these systems were based on advanced statistical methods and signal processing able to extract trading signals from financial data. However, the recent successes of Machine Learning have attracted the interest of the financial community and have raised the question if the techniques that have proven so successful in classifying images or beating Go world champions could perform equally well in detecting patterns in the financial markets. In this thesis, we explore how to find a trading strategy via Reinforcement Learning, a general class of algorithms that allows an agent to find an optimal strategy for a sequential decision problem by directly interacting with the environment. In particular, we focus on Policy Gradient Methods, a class of state-of-the-art techniques which constructs an estimate of the optimal policy for the control problem by iteratively improving a parametric policy.

In the first part of this thesis, we conduct a thorough review of these methods in the traditional risk-neutral framework for reinforcement learning, in which the goal of the agent is to maximize its total returns. In the second part, we propose an original extension of these methods to the risk-sensitive framework, in which the agent wishes instead to optimize the trade-off between the rewards and the risk required to achieve them. We also present an innovative parameter-based reformulation of these algorithms in which the optimal policy is approximated by stochastically perturbing the parameter of a deterministic controller. This extends the well-known Policy Gradient with Parameter-based Exploration (PGPE) method to the risk-sensitive framework. To the best of our knowledge, this is the first time such an algorithm is discussed. In the last part of this thesis, after a review of the financial applications of reinforcement learning that can be found in the literature, we present a numerical application of these algorithms to the traditional asset allocation problem with transaction costs.

Keywords: Automated Trading Systems, Asset Allocation, Machine Learning, Reinforcement Learning, Policy Gradient Methods.

Sommario

L'impatto sui mercati finanziari degli Automated Trading Systems, ovvero sistemi capaci di eseguire ordini in maniera automatica senza il bisogno di supervisione umana, è in costante crescita e ormai la grande maggioranza delle operazioni nelle borse mondiali è generata da algoritmi. Tradizionalmente, questi sistemi erano basati su sofisticati metodi statistici o di analisi dei segnali capaci di estrarre dei segnali dai dati finanziari. Tuttavia, i recenti successi del Machine Learning hanno attratto l'interesse dell'industria finanziaria e in molti si sono chiesti se le tecniche che si sono rivelate così opportune per la classificazione di immagini e per battere il campione del mondo di Go possano essere ugualmente efficaci nell'identificazione di segnali nei mercati finanziari. In questa tesi si esplora come determinare una strategia di trading attraverso Reinforcement Learning, una famiglia di algoritmi che permette ad un agente di trovare una strategia ottima per un problema di controllo ottimo stocastico a tempo discreto attraverso l'interazione diretta con l'ambiente. In particolare, ci si concentra sui metodi Policy Gradient, delle tecniche avanzate che generano un'approssimazione della strategia ottima attraverso il miglioramento iterativo di una strategia parametrica.

Nella prima parte di questa tesi, si conduce uno studio bibliografico di questi metodi per la versione neutrale al rischio del problema di controllo, in cui l'agente è solamente interessato a massimizzare i propri guadagni. Nella seconda parte della tesi, si propone un'estensione originale di questi metodi alla formulazione del problema di controllo in cui l'agente è avverso al rischio e desidera controllare il rischio necessario per realizzare un guadagno. Inoltre, si presenta un'innovativa riformulazione "parameter-based" di questi algoritmi, in cui la ricerca della strategia ottima procede attraverso la perturbazione stocastica dei parametri di un controllore deterministico. Questa idea estende i ben noti metodi "Policy Gradient with Parameter-based Exploration" (PGPE) nel caso di aversione al rischio. Al meglio delle nostre conoscenze, questo è il primo studio in cui un tale algoritmo viene proposto. Nell'ultima parte di questa tesi, dopo aver discusso le applicazioni finanziarie del Reinforcement Learning che si trovano in letteratura, si presenta un'applicazione numerica di questi algoritmi al classico problema di asset allocation con costi di transazione.

Parole chiave: Automated Trading Systems, Asset Allocation, Machine Learning, Reinforcement Learning, Policy Gradient Methods.

Ringraziamenti

Vorrei ringraziare...

Contents

List of Figures	vii
List of Tables	ix
List of Algorithms	xi
Acronyms	xiii
1 Introduction	1
1.1 The Computerization of Finance	1
1.2 The New Dawn of Artificial Intelligence	2
1.3 Structure	3
2 Discrete-Time Stochastic Optimal Control	5
2.1 Markov Decision Processes	5
2.2 Policies	6
2.3 Risk-Neutral Framework	7
2.3.1 Discounted Reward Formulation	8
2.3.2 Average Reward Formulation	10
2.4 Risk-Sensitive Framework	12
2.4.1 Discounted Reward Formulation	12
2.4.2 Average Reward Formulation	15
2.5 Dynamic Programming Algorithms	16
2.5.1 Value Iteration	17
2.5.2 Policy Iteration	17
3 Reinforcement Learning	19
3.1 The Reinforcement Learning Problem	19
3.2 Model-Free RL Methods	20
3.2.1 Model Approximation	21
3.2.2 Value Approximation	22
3.2.3 Policy Approximation	22
4 Risk-Neutral Policy Gradient	23
4.1 Basics of Policy Gradient Methods	23
4.2 Risk-Neutral Objective Functions	24

4.3	Finite Differences	25
4.4	Likelihood Ratio Methods	25
4.4.1	Monte Carlo Policy Gradient	26
4.4.2	GPOMDP	28
4.4.3	Stochastic Policies	29
4.4.4	Policy Gradient with Parameter Exploration	30
4.5	Risk-Neutral Policy Gradient Theorem	33
4.5.1	Theorem Statement and Proof	34
4.5.2	GPOMDP	36
4.5.3	Actor-Critic Policy Gradient	36
4.5.4	Compatible Function Approximation	39
4.5.5	Natural Policy Gradient	39
5	Risk-Sensitive Policy Gradient	43
5.1	Risk-Sensitive Framework	43
5.2	Monte Carlo Policy Gradient	45
5.3	Policy Gradient Theorem	46
5.3.1	Average Reward Formulation	47
5.3.2	Risk-Sensitive Actor-Critic Algorithm	48
5.3.3	Discounted Reward Formulation	49
6	Parameter-Based Policy Gradient	53
6.1	Risk-Neutral Framework	53
6.1.1	Parameter-Based Natural Policy Gradient	55
6.2	Risk-Sensitive Framework	60
6.2.1	Parameter-Based Natural Policy Gradient	60
7	Financial Applications of Reinforcement Learning	63
7.1	Efficient Market Hypothesis	63
7.1.1	Formal Definitions of the EMH	64
7.1.2	Critics to the EMH	65
7.2	Bibliographical Survey	66
7.2.1	Asset Allocation with Transaction Costs	67
7.2.2	Optimal Order Execution in Limit Order Book	68
7.2.3	Smart Order Routing Across Dark Pools	69
7.3	Asset Allocation with Transaction Costs	70
7.3.1	Wealth Dynamics	70
7.3.2	Rewards and Objective Functions	72
7.3.3	States	73
7.3.4	Actions	74
8	Numerical Results for the Asset Allocation Problem	77
8.1	Synthetic Risky Asset	77
8.1.1	Specifications of the Learning Algorithms	78
8.1.2	Experimental Setup	78
8.1.3	Risk-Neutral Framework	79

8.1.4	Risk-Sensitive Framework	82
8.2	Historical Risky Asset	87
8.2.1	Risk-Neutral Framework	87
8.2.2	Risk-Sensitive Framework	89
8.2.3	The Challenge of Historical Data	89
9	Conclusions	91
9.1	What Has Been Achieved	91
9.2	Conclusion	92
9.3	Further Developments	93
	Appendices	95
A	Implementation	95
A.1	Python Prototype	95
A.2	C++ Implementation	96
A.2.1	Environment, Task, Agent and Experiment	97
A.2.2	ARACAgent	99
A.3	Execution Pipeline	102
A.3.1	Compilation	102
A.3.2	<code>generate_synthetic_series.py</code>	102
A.3.3	<code>experiment_launcher.py</code>	102
A.3.4	<code>main_thesis</code>	102
A.3.5	<code>postprocessing.py</code>	103
	Bibliography	105

List of Figures

2.1	Agent-environment interaction in sequential decision problems. .	6
2.2	Policy iteration algorithm	17
3.1	Solution process for the control problem.	21
4.1	“Vanilla” policy gradient vs. natural policy gradient	40
7.1	Episodic formulation for the asset allocation problem.	74
8.1	Risk-neutral learning process for one synthetic risky asset	79
8.2	Backtest performance with one synthetic risky asset	80
8.3	Proportional transaction costs and risk-neutral strategies	82
8.4	Short-selling fees and risk-neutral strategies	82
8.5	Risk-sensitive learning process for one synthetic risky asset . . .	83
8.6	Backtest performance with one synthetic risky asset	84
8.7	Proportional transaction costs and risk-sensitive strategies . . .	86
8.8	Short-selling fees and risk-sensitive strategies	86
8.9	Risk-neutral learning process for the historical risky asset	88
8.10	Backtest performance with one historical risky asset	88
8.11	Risk-sensitive learning process for the historical risky asset . . .	90
8.12	Backtest performance with one historical risky asset	90
A.1	PyBrain standard architecture for an RL problem.	96
A.2	Class architecture for the asset allocation problem.	98
A.3	Class architecture for an ARAC agent	101
A.4	Execution pipeline of an asset allocation experiment	104

List of Tables

8.1	Backtest statistics for risk-neutral learning with one synthetic risky asset	81
8.2	Backtest statistics for risk-sensitive learning with one synthetic risky asset	84
8.3	Comparison of NPGPE and RSNPGPE for the asset allocation problem with one synthetic risky asset.	85
8.4	Backtest statistics for risk-neutral learning with one historical risky asset	88
8.5	Backtest statistics for risk-sensitive learning with one historical risky asset	90

List of Algorithms

4.1	General setup for a policy gradient algorithm.	24
4.2	Episodic REINFORCE policy gradient estimate	28
4.3	Episodic PGPE algorithm	31
4.4	Generic structure for an online actor-critic algorithm.	37
4.5	TD(λ) policy gradient algorithm.	38
5.1	Risk-sensitive REINFORCE policy gradient estimate	46
5.2	Risk-Sensitive Average Reward Actor-Critic algorithm	50
6.1	Actor-Critic PGPE	56
6.2	NPGPE	58
6.3	Natural Actor-Critic PGPE	59
6.4	Risk-Sensitive NPGPE	61

Acronyms

AI	Artificial Intelligence. 1 , 2
ATS	Automated Trading Systems. 1
BPTT	Backpropagation Through Time. 67
DL	Deep Learning. 1 , 2 , 66
EMH	Efficient Market Hypothesis. 63 , 87
HFT	High-Frequency Trading. 1
MDP	Markov Decision Process. 20
ML	Machine Learning. 1 , 2 , 91
NPGPE	Natural Policy Gradient with Parameter Based Exploration. 55
PGPE	Policy Gradient with Parameter-Based Exploration. 3 , 91
RL	Reinforcement Learning. 1–3 , 19 , 91
RRL	Recurrent Reinforcement Learning. 67
RTRL	Real-Time Recurrent Learning. 67
SOR	Smart Order Routing. 69

Chapter 1

Introduction

The impact of [Automated Trading Systems \(ATS\)](#) on financial markets is growing every year and the trades generated by an algorithm now account for the majority of orders that arrive at stock exchanges. Investments banks and hedge funds are investing large amount of resources to develop new systems to maximize their profits and stay ahead in the technological race that is taking place in the financial markets. The goal of developing an automated system able to take financial decisions without the need for human supervision is not new. From the Technical Computer System developed by Commodities Corporation in the 1970, to the algorithm that caused Knight Capital Group to lose \$440 millions over few hours in 2012, many have been the attempts to trade on the markets using an algorithm. Typically, these approaches were based on advanced statistical methods and signal processing. However, the recent successes of [Artificial Intelligence \(AI\)](#), and in particular [Deep Learning \(DL\)](#), have attracted the interest of the financial community. Many quantitative researchers asked themselves if the techniques that have proved so successful in classifying images or beating Go world champions could perform equally well on financial markets. In this thesis, we explore how to find a trading strategy via [Reinforcement Learning \(RL\)](#), a branch of [Machine Learning \(ML\)](#) that allows to find an optimal strategy for a sequential decision problem by directly interacting with the environment. In this introductory chapter, we briefly describe the current state of the financial markets, which are increasingly dominated by algorithms, and provide a quick overview of the recent advances of [AI](#) and the huge impact that these new techniques are having on our day-to-day life. This chapter thus prepare the stage in which we will move during the entire exposition.

1.1 The Computerization of Finance

The trading pit of a stock exchange is often imagined by outsiders as a frenzy place, with telephones constantly ringing and traders shouting orders across the room at a frenetic rhythm. This was probably the reality around thirty years ago, when open outcry was still the main communication system in the pits. Since then the exchanges have become more and more quiet as the majority of

the orders moved to electronic trading platforms, which allow investors to directly execute their orders without passing through the old-fashioned pit traders. In the last decade, the markets have witnessed the widespread adoption of [Automated Trading Systems \(ATS\)](#), that can make investment decisions in a fully automatized way at speeds with orders of magnitude greater than any human equivalent. Of particular interest nowadays is [High-Frequency Trading \(HFT\)](#), a type of algorithmic trading characterized by high speeds, high turnover rates, and high order-to-trade ratios that leverages high-frequency financial data and electronic trading tools. While there is no single definition of [HFT](#), among its key attributes are highly sophisticated algorithms, specialized order types, co-location, very short-term investment horizons, and high cancellation rates of orders. It has been estimated that in 2014 [HFT](#) accounted for more than 75% of the stock shares traded on US exchanges.

The deep transformation of financial markets has brought algorithmic trading and [HFT](#) into the spotlight and raised many questions on the actual utility of these technologies. On one side, computerized trading firms have been blamed of causing bouts of extreme volatility and of systematically exploiting the small investors thanks to superior technology, as was claimed for example in [61]. This is typically the position of regulators and politicians, many of whom put the policing of [HFT](#) at the top of their priorities. On the other hand, the academic literature strongly supports [HFT](#) having a net positive effect on the markets [23]. The usual argument is that these firms make the markets more stable and efficient, even when volatility surges, by quickly bringing prices back into line. By acting as market makers, [HFT](#) provides liquidity to the markets and thus reduces bid-ask spreads and transaction costs.

To develop the algorithms able to extract trading signals from large amount of noisy data and automatically generate financial decisions, investment banks and quantitative hedge funds are employing mathematicians, physicists and computer scientists. These algorithms are typically based on advanced statistics, signal processing, machine learning and other fields of mathematics. However, few of these firms publish their profit-generating “secret ingredient” and not much can be found in the literature. In this thesis we will discuss the application of some state-of-the-art [AI](#) techniques to determine an automatic and potentially profitable trading strategy.

1.2 The New Dawn of Artificial Intelligence

After many false dawns, [AI](#) has made extraordinary progress in the past few years, thanks in particular to a versatile technique called [Deep Learning \(DL\)](#), a form of [Machine Learning \(ML\)](#) that enables computers to learn from experience and understand the world in terms of a hierarchy of concepts. Because the computer gathers knowledge from experience, there is no need for a human computer operator to formally specify all the knowledge that the computer needs. The hierarchy of concepts allows the computer to learn complicated concepts by building them out of simpler ones. A graph of these hierarchies would be many

layers deep [44]. Given enough data, deep neural networks can be trained to successfully complete many different tasks that could previously be done only by humans [1]. For instance, these methods power Google’s search engine, Facebook’s automatic photo tagging, Apple’s voice assistant, Amazon’s shopping recommendations and Tesla’s self-driving cars. These techniques have recently been in the spotlight when AlphaGo, an algorithm developed by DeepMind, defeated the 18-time world champion of Go Lee Sedol in a 5-game series [95]. This ancient board game is so complex that computers had not been expected to master it for at least another decade at least. It is not surprising that this rapid sequence of successes have produce a large hype around ML and DL in particular. Citing John Giannandrea, head of machine-intelligence research at Google, “what got people excited about this field is that one learning technique, deep learning, can be applied to so many different domains”. The financial industry is not excluded.

Hedge funds, stung by eight years of underperformance, are latching onto machine learning as a high-tech answer to their woes. But Wall Street’s heady search for the perfect money machine has collided with a sober reality. The technology, which learns on its own to find investment ideas by hunting through troves of data, requires a heavy commitment of time and money, and a high tolerance for failure, since most algorithms turn out to be duds [56]. Financial markets are a particularly difficult challenge for ML algorithms and there are infinite ways an algorithm can fail spotting the trading “signals” to wager on. While identifying patterns is not particularly hard, finding signals that work reliably in the real world is. Man AHL, a quant unit of Man Group Plc, needed three years of work to gain enough confidence in a machine learning strategy to devote client money to it. This can be hard for the researcher, who will have to stomach and overcome many failures before obtaining positive results. In this thesis, we decided to embark on the difficult journey of developing an automated trading strategy based on Reinforcement Learning (RL), a general class of algorithms in the field of ML that allows an agent to learn how to behave in a stochastic and possibly unknown environment only by trial-and-error.

1.3 Structure

This section outlines the structure of this document, explaining the original contributions that were made in this thesis.

In Chapter 2 we introduce the basic concepts of discrete-time stochastic optimal control, which is the standard theoretical framework used to model sequential decision problems. In particular, we describe the central feedback mechanism between a learning agent and a stochastic environment. Finally, we present the less traditional risk-sensitive framework for discrete-time stochastic optimal control.

In Chapter 3 we present the main ideas of reinforcement learning, a general class of algorithms in the field of machine learning that allows an agent to learn how to behave in a stochastic and possibly unknown environment only by trial-and-

error. We thoroughly discuss the [RL](#) problem and its characteristic features. Finally, we give a high-level overview of the different typologies of reinforcement learning algorithms.

In [Chapter 4](#) we give an in-depth presentation of policy gradient algorithms for the risk-neutral control problem. After introducing the key ideas of these methods, we provide a thorough review of the state-of-the-art algorithms that can be found in the literature. In particular, we introduce the main result that will play a crucial role in the rest of this thesis, the policy gradient theorem.

In [Chapter 5](#) we discuss policy gradient methods for the risk-sensitive control problem, which is still an active field of research. In particular, we provide an extension of the policy gradient theorem to the risk-sensitive framework, both in the average reward and in the discounted reward formulations. To the best of our knowledge, this is the first time that a risk-sensitive policy gradient theorem is derived for a general discounted Markov decision process.

In [Chapter 6](#) we propose an original parameter-based version of the policy gradient theorem, both for the risk-neutral and the risk-sensitive formulation. These theorems allow to derive efficient online learning algorithms similar in spirit to the well-known [Policy Gradient with Parameter-Based Exploration \(PGPE\)](#) algorithm, which was originally conceived only for episodic environments. Moreover, these new algorithms can be easily enhanced using a critic or the natural policy gradient idea. This chapter undoubtedly represents the most innovative contribution of this thesis.

In [Chapter 7](#), after a brief discussion about why finance represents an extremely challenging field of research, we provide a bibliographical survey of successful applications of [RL](#) techniques to financial problems. In particular, we focus on the asset allocation problem with transaction cost, which is used as a test case for the numerical application of the learning algorithms proposed in the previous chapters.

In [Chapter 8](#) we present the numerical results for the asset allocation problem. We show that the learning algorithms proposed in the previous chapters perform extremely well on synthetic data. On the other hand, the algorithms encounter more difficulties on historical data and we try to provide an explanation for this behavior.

[Chapter 9](#) summarizes the contributions of this thesis to the reinforcement learning literature and presents some interesting ideas for further developments and future work.

Given the complexity of financial markets and the ambitious goal of this thesis, success can not be a-priori guaranteed. In this case, the journey matters more than the arrival, as this thesis will allow us to explore these fascinating and current techniques that are revolutionizing the world.

Chapter 2

Discrete-Time Stochastic Optimal Control

In sequential decision problems, an agent interacts with an environment by selecting a series of actions in order to complete a specific task. During this interaction, the agent receives a numerical reward from the environment and its goal is to find the best strategy in order to maximize a certain measure of its performance. The environment evolves stochastically and may be influenced by the interaction with the agent, so that each action taken by the agent may influence the circumstances under which future decisions will be made. Therefore, the agent must balance his desire to obtain a large reward today by acting greedily and the opportunities that will be available in the future. While this setting appears quite simple, it is general enough to encompass a wide range of applications in different fields. A classical example is portfolio management, where an investor must allocate his capital so as to maximize his long-term profits. Another standard example is chess, where two players successively move pieces around the chessboard to checkmate the opponent's king.

The purpose of the following sections is to introduce the notation that will be used in the rest of this work and to recall the fundamental concepts and results of the discrete-time stochastic optimal control theory, which is the standard framework to study sequential decisions problems in mathematical terms. Since our discussion will be far from being comprehensive, we refer the reader to the extensive literature on the subject, such as [16], [84], [15].

2.1 Markov Decision Processes

A sequential decision problem under uncertainty can be schematized as in Figure 2.1: at a given time t , the agent (also known as decision maker or controller) observes the state s_t of the system (also known as environment) and subsequently performs an action a_t . Following this action, the agent receives an immediate reward r_{t+1} (or incurs an immediate cost) and the system evolves to a new state according to a probability distribution which depends on the action selected by the agent. At the subsequent time $t + 1$, the agent selects a new action given




Figure 2.1: Agent-environment interaction in sequential decision problems.

the new state of the system and the process repeats. This interaction can be modeled rigorously using a Markov decision process.

Definition 2.1.1 (Markov Decision Process). *A Markov decision process (MDP) is a stochastic dynamical system specified by the tuple $\langle \mathbb{S}, \mathbb{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$, where*

- i) $(\mathbb{S}, \mathcal{S})$ is a measurable space, called the state space.*
- ii) $(\mathbb{A}, \mathcal{A})$ is a measurable space, called the action space.*
- iii) $\mathcal{P} : \mathbb{S} \times \mathbb{A} \times \mathcal{S} \rightarrow \mathbb{R}$ is a Markov transition kernel, i.e.*
 - a) for every $s \in \mathbb{S}$ and $a \in \mathbb{A}$, $B \mapsto \mathcal{P}(s, a, B)$ is a probability distribution over $(\mathbb{S}, \mathcal{S})$.*
 - b) for every $B \in \mathcal{S}$, $(s, a) \mapsto \mathcal{P}(s, a, B)$ is a measurable function on $\mathbb{S} \times \mathbb{A}$.*
- iv) $\mathcal{R} : \mathbb{S} \times \mathbb{A} \rightarrow \mathbb{R}$ is a reward function.*
- v) $\gamma \in (0, 1)$ is a discount factor.*

Typically, the state space (and similarly the action space) will be either finite, namely $\mathbb{S} = \{s_1, \dots, s_d\}$, or continuous, namely $\mathbb{S} \subseteq \mathbb{R}^{D_s}$. The kernel \mathcal{P} describes the random evolution of the system: suppose that at time t the system is in state s and that the agent takes action a , then, regardless of the previous history of the system, the probability to find the system in a state belonging to $B \in \mathcal{S}$ at time $t + 1$ is given by

$$\mathcal{P}(s, a, B) = \mathbb{P}(S_{t+1} \in B | S_t = s, A_t = a) \quad (2.1)$$

Following this random transition, the agent receives a stochastic reward R_{t+1} . The reward function $\mathcal{R}(s, a)$ gives the expected reward obtained when action a is taken in state s

$$\mathcal{R}(s, a) = \mathbb{E}[R_{t+1} | S_t = s, A_t = a] \quad (2.2)$$

This setting can be easily generalized to the following cases

1. The initial state of the system is a random variable.
2. The actions that an agent can select depend on the state of the system.

2.2 Policies

At any time step, the agent selects its actions according to a certain policy.

Definition 2.2.1 (Policy). *A policy is a function $\pi : \mathbb{S} \times \mathcal{A} \rightarrow \mathbb{R}$ such that*

- i) for every $s \in \mathbb{S}$, $C \mapsto \pi(s, C)$ is a probability distribution over $(\mathbb{A}, \mathcal{A})$.*
- ii) for every $C \in \mathcal{A}$, $s \mapsto \pi(s, C)$ is a measurable function.*

Intuitively, a policy represents a stochastic mapping from the current state of the system to actions. Deterministic policies are a particular case of this general definition. We assumed that the agent's policy is stationary and only depends on the current state of the system. We might in fact consider more general policies that depends on the whole history of the system. However, as we will see, we can always find an optimal policy that depends only on the current state, so that our definition is not restrictive. A policy π and an initial state $s_0 \in \mathbb{S}$ determine a random state-action-reward sequence $\{(S_t, A_t, R_{t+1})\}_{t \geq 0}$ with values on $\mathbb{S} \times \mathbb{A} \times \mathbb{R}$ following the mechanism described above.

Definition 2.2.2 (History). *Given an initial state $s_0 \in \mathbb{S}$ and a policy π , a history (or equivalently trajectory or roll-out) of the system is a random sequence $H_\pi = \{(S_t, A_t)\}_{t \geq 0}$ with values in $\mathbb{S} \times \mathbb{A}$, defined on some probability space $(\Omega, \mathcal{F}, \mathbb{P})$, such that for $t = 0, 1, \dots$*

$$\begin{cases} S_0 = s_0 \\ A_t \sim \pi(S_t, \cdot) \\ S_{t+1} \sim \mathcal{P}(S_t, A_t, \cdot) \end{cases}$$

we will denote by $(\mathbb{H}, \mathcal{H})$ the measurable space of all possible histories.

Moreover, we observe that

- i) the state sequence $\{S_t\}_{t \geq 0}$ is a Markov process $\langle \mathbb{S}, \mathcal{P}_\pi \rangle$.
- ii) the state-reward sequence $\{(S_t, R_{t+1})\}_{t \geq 0}$ is a Markov reward process $\langle \mathbb{S}, \mathcal{P}_\pi, \mathcal{R}_\pi, \gamma \rangle$.

where we denoted

$$\begin{aligned} \mathcal{P}_\pi(s, s') &= \int_{\mathbb{A}} \pi(s, a) \mathcal{P}(s, a, s') da \\ \mathcal{R}_\pi(s) &= \int_{\mathbb{A}} \pi(s, a) \mathcal{R}(s, a) da \end{aligned}$$

In stochastic optimal control, the goal of the agent is to find a policy that maximizes a measure of the agent's long-term performance. In the next sections we discuss some objective functions that are commonly used in the infinite horizon framework.

2.3 Risk-Neutral Framework

In the risk-neutral setting, the agent is only interested in maximizing its reward, without considering the risk it needs to take on to achieve it. In an infinite horizon task, the agent's performance is typically measured either as the total discounted reward or as the average reward obtained at each time step. These two approaches are radically different both from a theoretical and an algorithmic point of view. For this reason, they will be always treated separately.

2.3.1 Discounted Reward Formulation

In the discounted reward formulation, the agent's performance is measured as the expected return obtained following a specific policy.

Definition 2.3.1 (Return). *The return is the total discounted reward obtained by the agent starting from t*

$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

where $0 < \gamma < 1$ is the discount factor.

In some domains, such as economics, discounting can be used to represent interest earned on rewards, so that an action that generates an immediate reward will be preferred over one that generates the same reward some steps into the future. Discounting thus models the trade-off between immediate and delayed reward: if $\gamma = 0$ the agent selects its actions in a myopic way, while if $\gamma \rightarrow 1$ it acts in a far-sighted manner. There are other possible reasons for discounting future rewards. The first is because it is mathematically convenient, as it avoids infinite returns and it solves many convergence issues. Another interpretation is that it models the uncertainty about the future, which may not be fully represented. Indeed, the discount factor could be seen as the probability that the world does not stop at a given time step.

Definition 2.3.2 (State-Value Function). *The state-value function $V_\pi : \mathbb{S} \rightarrow \mathbb{R}$ is the expected return that can be obtained starting from a state and following policy π*

$$V_\pi(s) = \mathbb{E}_\pi [G_t | S_t = s] \quad (2.3)$$

where the subscript in \mathbb{E}_π indicates that all the actions are selected according to policy π . The state-value function measures how good it is for the agent to be in a given state and follow a certain policy. Similarly, we can introduce an action-value function that measures how good it is for the agent to be in a state, take a certain action and then follow the policy.

Definition 2.3.3 (Action-Value Function). *The action-value function $Q_\pi : \mathbb{S} \times \mathbb{A} \rightarrow \mathbb{R}$ is the expected return that can be obtained starting from a state, taking an action and then following policy π*

$$Q_\pi(s, a) = \mathbb{E}_\pi [G_t | S_t = s, A_t = a] \quad (2.4)$$

We have the following relationship between V_π and Q_π

$$V_\pi(s) = \int_{\mathbb{A}} \pi(s, a) Q_\pi(s, a) da \quad (2.5)$$

Almost all reinforcement learning algorithms are designed to estimate these value functions and are typically based on the Bellman equations.

Proposition 2.3.1 (Bellman Expectation Equations).

$$V_\pi(s) = \mathcal{R}_\pi(s) + \gamma T_\pi V_\pi(s) \quad (2.6)$$

$$Q_\pi(s, a) = \mathcal{R}(s, a) + \gamma T_a V_\pi(s) \quad (2.7)$$

where we denoted by T_a (resp. T_π) the transition operator for action a (resp. for policy π)

$$\begin{aligned} T_a F(s) &= \mathbb{E}[F(S_{t+1}) | S_t = s, A_t = a] = \int_{\mathbb{S}} \mathcal{P}(s, a, s') F(s') ds' \\ T_\pi F(s) &= \mathbb{E}_\pi[F(S_{t+1}) | S_t = s] = \int_{\mathbb{A}} \pi(s, a) \int_{\mathbb{S}} \mathcal{P}(s, a, s') F(s') ds' da \end{aligned}$$

If we introduce the Bellman expectation operator B_π , defined as

$$B_\pi V_\pi(s) = \mathcal{R}_\pi(s) + \gamma T_\pi V_\pi(s)$$

Then Eq. (2.6) can be written as a fixed-point equation

$$V_\pi(s) = B_\pi V_\pi(s)$$

which, under some simple assumptions on the reward functions, admits a unique solution by the contraction mapping theorem. A similar argument holds for Eq. (2.7). The goal of the agent is to select a policy π_* that maximizes his expected return in all possible states. Such a policy is called *optimal*.

Definition 2.3.4 (Optimal State-Value Function). *The optimal state-value function $V_* : \mathbb{S} \rightarrow \mathbb{R}$ is the largest expected return that can be obtained starting from a state*

$$V_*(s) = \sup_{\pi} V_\pi(s) \quad (2.8)$$

Definition 2.3.5 (Optimal Action-Value Function). *The optimal action-value function $Q_* : \mathbb{S} \times \mathbb{A} \rightarrow \mathbb{R}$ is the largest expected return that can be obtained starting from a state and taking an action*

$$Q_*(s, a) = \sup_{\pi} Q_\pi(s, a) \quad (2.9)$$

The optimal value functions satisfy the following Bellman equations.

Proposition 2.3.2 (Bellman Optimality Equations).

$$V_*(s) = \sup_a Q_*(s, a) = \sup_a \{\mathcal{R}(s, a) + \gamma T_a V_*(s)\} \quad (2.10)$$

$$\begin{aligned} Q_*(s, a) &= \mathcal{R}(s, a) + \gamma T_a V_*(s) \\ &= \mathcal{R}(s, a) + \gamma \int_{\mathbb{S}} \mathcal{P}(s, a, s') \sup_{a'} Q_*(s', a') ds' \end{aligned} \quad (2.11)$$

Again, these two equations are fixed-point equations and the existence and uniqueness of a solution is guaranteed, under some technical assumptions, by the contraction mapping theorem. Starting from the optimal value functions, we can easily derived an optimal policy. Let us define a partial ordering in the policy space

$$\pi \succeq \pi' \Leftrightarrow V_\pi(s) \geq V_{\pi'}(s) \quad \forall s \in \mathbb{S} \quad (2.12)$$

Then the optimal policy $\pi_* \succeq \pi, \forall \pi$. We have the following results

Theorem 2.3.1 (Optimal Policy). *For any Markov decision process,*

- i) *It exists an optimal policy π_* such that $\pi_* \succeq \pi, \forall \pi$.*
- ii) $V_{\pi_*}(s) = V_*(s)$.
- iii) $Q_{\pi_*}(s, a) = Q_*(s, a)$.

An optimal policy can be found by acting greedily with respect to Q_* , which corresponds to selecting the action that maximizes the action-value function in a given state

$$a_* = \arg \sup_{a \in \mathbb{A}} Q_*(s, a) \quad (2.13)$$

This policy is deterministic and only depends on the current state of the system.

2.3.2 Average Reward Formulation

Most of the research in RL has studied a problem formulation where agents maximize the cumulative sum of rewards. However, this approach cannot handle infinite horizon tasks, where there are no absorbing goal states, without discounting future rewards. Clearly, discounting is only necessary in cyclical tasks, where the cumulative reward sum can be unbounded. More natural long-term measure of optimality exists for such cyclical tasks, based on maximizing the average reward per action. For a more in-depth presentation, the reader may again refer to the extensive literature on the subject, such as [7], [66] and the references therein. In the average reward setting, also known as long-run reward or ergodic reward, the goal of the agent is to find a policy that maximizes the expected reward per step.

Definition 2.3.6 (Average Reward). *The average reward ρ_π associated to a policy π is defined as*

$$\begin{aligned} \rho_\pi &= \lim_{T \rightarrow \infty} \frac{1}{T} \mathbb{E}_\pi \left[\sum_{t=0}^{T-1} R_{t+1} \right] \\ &= \mathbb{E}_{\substack{S \sim d_\pi \\ A \sim \pi}} [\mathcal{R}(S, A)] \\ &= \int_{\mathbb{S}} d_\pi(s) \int_{\mathbb{A}} \pi(s, a) \mathcal{R}(s, a) da ds \end{aligned} \quad (2.14)$$

where d_π is the stationary distribution of the Markov process induced by π .

The agent aims to find an *average optimal* policy

$$\pi_* = \arg \sup_{\pi} \rho_{\pi} \quad (2.15)$$

In this setting, we introduce the *average adjusted* value and action-value functions.

Definition 2.3.7 (Average Adjusted State-Value Function). *The average adjusted state-value function $V_{\pi} : \mathbb{S} \rightarrow \mathbb{R}$ is the expected residual return that can be obtained starting from a state and following policy π*

$$V_{\pi}(s) = \mathbb{E}_{\pi} \left[\sum_{t=0}^{\infty} (R_{t+1} - \rho_{\pi}) \mid S_0 = s \right] \quad (2.16)$$

The term $V_{\pi}(s)$ is usually referred to as the *bias* value, or the *relative* value, since it represents the relative difference in total reward gained starting from a state s as opposed to a generic state. ρ_{π} serves as a baseline that allows to avoid divergence in the value function definition.

Definition 2.3.8 (Average Adjusted Action-Value Function). *The average adjusted action-value function $Q_{\pi} : \mathbb{S} \times \mathbb{A} \rightarrow \mathbb{R}$ is the expected residual return that can be obtained starting from a state, taking an action and then following policy π*

$$Q_{\pi}(s, a) = \mathbb{E}_{\pi} \left[\sum_{t=0}^{\infty} (R_{t+1} - \rho_{\pi}) \mid S_0 = s, A_0 = a \right] \quad (2.17)$$

We have the following relation between the state-value function and the action-value function

$$V_{\pi}(s) = \int_{\mathbb{A}} \pi(s, a) Q_{\pi}(s, a) da \quad (2.18)$$

The value functions satisfy the following Bellman equation

Proposition 2.3.3 (Bellman Expectation Equations).

$$V_{\pi}(s) = \mathcal{R}_{\pi}(s) - \rho_{\pi} + T_{\pi} V_{\pi}(s) \quad (2.19)$$

$$Q_{\pi}(s, a) = \mathcal{R}(s, a) - \rho_{\pi} + T_a V_{\pi}(s) \quad (2.20)$$

Again, by introducing opportune Bellman operators, these equations can be rewritten as fixed-point equations. In the discrete case, where the transition operator correspond to matrices, these Bellman equations become linear systems that can be solved to obtain the value functions.

2.4 Risk-Sensitive Framework

In many application, in addition to maximizing the average reward, the agent may want to control risk by minimizing some measure of variability in rewards. In the risk-sensitive framework, the goal of the agent is to find the policy that optimally solves the trade-off between reward and risk. Although risk-sensitive sequential decision-making has a long history in operations research and finance, it has only recently grabbed the attention of the machine learning community. Hence, the literature offers many reference which approach the risk-sensitive control problem from the traditional stochastic optimal control perspective. On the other hand, there are only few references that attack the problem in the reinforcement learning setting. Again, we can consider the discounted formulation or the average formulation.

2.4.1 Discounted Reward Formulation

A standard way to measure the risk associated with a policy π is the variance of the total discounted reward obtained starting from a given state s

$$\Lambda_\pi(s) = \text{Var}_\pi(G_t | S_t = s) \quad (2.21)$$

This approach is the one considered in [97]. The variance can be decomposed as

$$\Lambda_\pi(s) = U_\pi(s) - V_\pi(s)^2 \quad (2.22)$$

where we denoted by U_π the square state-value function.

Definition 2.4.1 (Square State-Value Function). *The square state-value function $U_\pi(s)$ is the second moment of the return that can be obtained under policy π starting from a state s*

$$U_\pi(s) = \mathbb{E}_\pi[G_t^2 | S_t = s] \quad (2.23)$$

As for the risk-neutral formulation, it comes in handy to introduce a square action-value function.

Definition 2.4.2 (Square Action-Value Function). *The square action-value function $W_\pi(s)$ is the second moment of the return that can be obtained starting from a state s , taking action a and then following policy π*

$$W_\pi(s, a) = \mathbb{E}_\pi[G_t^2 | S_t = s, A_t = a] \quad (2.24)$$

Clearly, the two square value functions are related by the following equation

$$U_\pi(s) = \int_{\mathbb{A}} \pi(s, a) W_\pi(s, a) da \quad (2.25)$$

We would like to obtain a Bellman expectation equation for $U_\pi(s)$ and $W_\pi(s, a)$, in order to piggyback on the discussion about the standard value functions. To do so, let us introduce the square reward function.

Definition 2.4.3 (Square Reward Function). *The square reward function $\mathcal{M}(s, a)$ is the second moment of the reward that can be obtained in state s when taking action a*

$$\mathcal{M}(s, a) = \mathbb{E} [R_{t+1}^2 | S_t = s, A_t = a] \quad (2.26)$$

Let us also define the state-reward function $\mathcal{M}_\pi(s)$

$$\mathcal{M}_\pi(s) = \mathbb{E}_\pi [R_{t+1}^2 | S_t = s] = \int_{\mathbb{A}} \pi(s, a) \mathcal{M}(s, a) da \quad (2.27)$$

Moreover, we will need the reward-return state-covariance function.

Definition 2.4.4 (Reward-Return State-Covariance Function). *The reward-return state-covariance function $C_\pi(s)$ is the covariance between the first day reward and the successive returns starting from state s and then following policy π*

$$C_\pi(s) = \text{Cov}_\pi (R_{t+1}, G_{t+1} | S_t = s) \quad (2.28)$$

Again, it comes in handy to also introduce the reward-return action-covariance function.

Definition 2.4.5 (Reward-Return Action-Covariance Function). *The reward-return action-covariance function $C_\pi(s, a)$ is the covariance between the first reward and the successive return starting from state s , taking action a and then following policy π*

$$C_\pi(s, a) = \text{Cov}_\pi (R_{t+1}, G_{t+1} | S_t = s, A_t = a) \quad (2.29)$$

Then, it is easy to show that the square state-value function satisfies the following Bellman equation

Proposition 2.4.1 (Bellman Expectation Equation).

$$U_\pi(s) = \mathcal{K}_\pi(s) + \gamma^2 T_\pi U_\pi(s) \quad (2.30)$$

$$W_\pi(s, a) = \mathcal{K}_\pi(s, a) + \gamma^2 T_a U_\pi(s) \quad (2.31)$$

where we denoted

$$\mathcal{K}_\pi(s) = \mathcal{M}_\pi(s) + 2\gamma \mathcal{R}_\pi(s) T_\pi V_\pi(s) + 2\gamma C_\pi(s) \quad (2.32)$$

$$\mathcal{K}_\pi(s, a) = \mathcal{M}(s, a) + 2\gamma \mathcal{R}(s, a) T_a V_\pi(s) + 2\gamma C_\pi(s, a) \quad (2.33)$$

Proof. Let us prove the Bellman expectation equation for the square action-value function. The proof for the square state-value function is analogous.

$$\begin{aligned} W_\pi(s, a) &= \mathbb{E}_\pi [G_t^2 | S_t = s, A_t = a] \\ &= \mathbb{E}_\pi [(R_{t+1} + \gamma G_{t+1})^2 | S_t = s, A_t = a] \\ &= \mathcal{M}(s, a) + 2\gamma \mathbb{E}_\pi [R_{t+1} G_{t+1} | S_t = s, A_t = a] + \gamma^2 T_a U_\pi(s) \end{aligned}$$

By the definition of covariance

$$C_\pi(s, a) = \mathbb{E}_\pi [R_{t+1} G_{t+1} | S_t = s, A_t = a] - \mathcal{R}(s, a) T_a V_\pi(s)$$

Plugging in the first equation leads to the result. \square

These equations are analogous to the Bellman expectation equations for the state-value function, with a synthetic reward function $\mathcal{K}_\pi(s)$ and a discount factor γ^2 . It is possible to combine the Bellman expectation equation for $V_\pi(s)$ and for $U_\pi(s)$ to obtain a Bellman equation for the return variance $\Lambda_\pi(s)$.

Proposition 2.4.2 (Bellman Expectation Equation).

$$\Lambda_\pi(s) = \mathcal{V}_\pi(s) + 2\gamma C_\pi(s) + \gamma^2 \text{Var}_\pi(V_\pi(S_{t+1})|S_t = s) + \gamma^2 T_\pi \Lambda_\pi(s) \quad (2.34)$$

where $\mathcal{V}_\pi(s)$ denotes the conditional variance of the reward

$$\mathcal{V}_\pi(s) = \text{Var}_\pi(R_{t+1}|S_t = s) = \mathcal{M}_\pi(s) - \mathcal{R}_\pi(s)^2 \quad (2.35)$$

Proof. The result can be proved starting from Eq. (2.22) and exploiting the Bellman equations for U_π and V_π . Here we follow an alternative way based on the law of total variance.

$$\begin{aligned} \Lambda_\pi(s) &= \text{Var}_\pi(G_t|S_t = s) \\ &= \text{Var}_\pi(R_{t+1} + \gamma G_{t+1}|S_t = s) \\ &= \mathcal{V}_\pi(s) + 2\gamma \text{Cov}_\pi(R_{t+1}, G_{t+1}|S_t = s) + \gamma^2 \text{Var}_\pi(G_{t+1}|S_t = s) \end{aligned}$$

Applying the law of total variance we obtain

$$\begin{aligned} \text{Var}_\pi(G_{t+1}|S_t = s) &= \mathbb{E}_\pi[\text{Var}_\pi(G_{t+1}|S_{t+1})|S_t = s] + \text{Var}_\pi(\mathbb{E}_\pi[G_{t+1}|S_{t+1}]|S_t = s) \\ &= T_\pi \Lambda_\pi(s) + \text{Var}_\pi(V_\pi(S_{t+1})|S_t = s) \end{aligned}$$

Plugging it into the first equality yields the result

$$\Lambda_\pi(s) = \text{Var}_\pi(R_{t+1}|S_t = s) + \text{Var}_\pi(V_\pi(S_{t+1})|S_t = s) + \gamma^2 T_\pi \Lambda_\pi(s)$$

□

Even if appealing from a theoretical point of view, these equations cannot be easily exploited to derive practical algorithms in a non-episodic environment. Indeed, the covariance term between the first day reward and the future return appearing in the synthetic reward would be hard to estimate when the lifespan of the experiment can be infinite. This does not necessarily represent a problem in an episodic environment, where the experiments have a finite (possibly random) lifespan. Even more problematic is the variance of the state-value function, which typically is unknown. For this reason, we will not develop any reinforcement learning algorithm in this framework and we will instead focus on the average reward formulation.

In [104] and [83], the authors circumvent these difficulties by implicitly assuming that the reward R_{t+1} is conditionally independent from the future rewards given the current state. Under this assumption, the covariance terms are null and the previous Bellman equation become

Corollary 2.4.1 (Bellman Equations Under Independence Assumption).

$$U_\pi(s) = \mathcal{M}_\pi(s) + 2\gamma \mathcal{R}_\pi(s) T_\pi V_\pi(s) + \gamma^2 T_\pi U_\pi(s) \quad (2.36)$$

$$W_\pi(s, a) = \mathcal{M}(s, a) + 2\gamma \mathcal{R}(s, a) T_a V_\pi(s) + \gamma^2 T_a U_\pi(s) \quad (2.37)$$

$$\Lambda_\pi(s) = \mathcal{V}_\pi(s) + \gamma^2 \text{Var}_\pi(V_\pi(S_{t+1})|S_t = s) + \gamma^2 T_\pi \Lambda_\pi(s) \quad (2.38)$$

2.4.2 Average Reward Formulation

In [83], the authors consider the long-run variance as a measure of the risk associated to a policy π

Definition 2.4.6 (Long-Run Variance). *The long-run variance Λ_π under policy π is defined as*

$$\Lambda_\pi = \lim_{T \rightarrow \infty} \frac{1}{T} \mathbb{E}_\pi \left[\sum_{t=0}^{T-1} (R_{t+1} - \rho_\pi)^2 \right] \quad (2.39)$$

The long-run variance can be decomposed as follows

$$\Lambda_\pi = \eta_\pi - \rho_\pi^2 \quad (2.40)$$

where η_π is the average square reward per step

Definition 2.4.7 (Average Square Reward).

$$\begin{aligned} \eta_\pi &= \lim_{T \rightarrow \infty} \frac{1}{T} \mathbb{E}_\pi \left[\sum_{t=0}^{T-1} R_{t+1}^2 \right] \\ &= \mathbb{E}_{\substack{S \sim d_\pi \\ A \sim \pi}} [\mathcal{M}(S, A)] \\ &= \int_{\mathbb{S}} d_\pi(s) \int_{\mathbb{A}} \pi(s, a) \mathcal{M}(s, a) da ds \end{aligned} \quad (2.41)$$

where we denoted by $\mathcal{M}(s, a)$ the square reward function

$$\mathcal{M}(s, a) = \mathbb{E} [R_{t+1}^2 | S_t = s, A_t = a] \quad (2.42)$$

As before, we introduce the residual state-value and action-value functions associated with the square reward under policy π

Definition 2.4.8 (Average Adjusted Square State-Value Function). *The average adjusted square state-value function $U_\pi : \mathbb{S} \rightarrow \mathbb{R}$ is the expected square residual return that can be obtained starting from a state and following policy π*

$$U_\pi(s) = \mathbb{E}_\pi \left[\sum_{t=0}^{\infty} (R_{t+1}^2 - \eta_\pi) \middle| S_0 = s \right] \quad (2.43)$$

Definition 2.4.9 (Average Adjusted Square Action-Value Function). *The average adjusted square action-value function $Q_\pi : \mathbb{S} \times \mathbb{A} \rightarrow \mathbb{R}$ is the expected residual square return that can be obtained starting from a state, taking an action and then following policy π*

$$W_\pi(s, a) = \mathbb{E}_\pi \left[\sum_{t=0}^{\infty} (R_{t+1}^2 - \eta_\pi) \middle| S_0 = s, A_0 = a \right] \quad (2.44)$$

The following relation between square state-value function and the square action-value holds

$$U_\pi(s) = \int_{\mathbb{A}} \pi(s, a) W_\pi(s, a) da \quad (2.45)$$

The average adjusted square value functions satisfy the following Bellman equations

Proposition 2.4.3 (Bellman Expectation Equations).

$$U_\pi(s) = \mathcal{M}_\pi(s) - \eta_\pi + T_\pi U_\pi(s) \quad (2.46)$$

$$\begin{aligned} W_\pi(s, a) &= \mathcal{M}(s, a) - \eta_\pi + T_a U_\pi(s) \\ &= \mathcal{M}(s, a) - \eta_\pi + \int_{\mathbb{S}} \mathcal{P}(s, a, s') \int_{\mathbb{A}} \pi(s', a') W_\pi(s', a') da' ds' \end{aligned} \quad (2.47)$$

Now, how can we formally define the agent's objective to take into account the trade-off between rewards and risks? Borrowing from the financial literature, a first alternative is to consider the following mean-variance optimization problem

$$\begin{cases} \max_{\pi} \rho_\pi \\ \text{subject to } \Lambda_\pi \leq \alpha \end{cases} \quad (2.48)$$

for a given $\alpha > 0$. Using the Lagrangian relaxation procedure, we can recast (2.48) to the following unconstrained problem

$$\max_{\lambda} \min_{\pi} L(\pi, \lambda) = -\rho_\pi + \lambda(\Lambda_\pi - \alpha) \quad (2.49)$$

Alternatively, the agent may want to optimize the Sharpe ratio, a commonly used risk-sensitive performance measure defined as

$$\text{Sh}_\pi = \frac{\rho_\pi}{\sqrt{\Lambda_\pi}} \quad (2.50)$$

2.5 Dynamic Programming Algorithms

In discrete-time, the Bellman optimality equations can be seen as an expression of the dynamic programming principle and completely characterize the optimal value functions. In the case of discrete state and action spaces, the value functions can be represented as vectors and the Bellman operators become matrices so that the optimality equations can be directly solved in polynomial time [64]. An alternative approach is to approximate the optimal value functions by using the Bellman optimality equations as the basis of a fixed-point iterative scheme. This leads to the well known *value iteration* and *policy iteration* algorithms [101]. The drawback of these methods is that typically they are only applicable to discrete MDP whose dynamics is perfectly known. Still, these simple algorithms provide some useful insight that can be exploited to develop more advanced reinforcement learning algorithm.




Figure 2.2: Policy iteration algorithm [99].

2.5.1 Value Iteration

Value iteration is an iterative method to compute the optimal state-value function $V_*(s)$. Starting from an arbitrary function $V_0(s)$, the algorithm iteratively updates the approximation by applying the Bellman optimality operator in a fixed-point scheme

$$V_{k+1}(s) = B_* V_k(s) = \sup_a \{ \mathcal{R}(s, a) + \gamma T_a V_k(s) \}$$

This algorithm is guaranteed to converge to the optimal value function V_* by the contraction mapping theorem. Let us notice that the transition operator T_a requires knowledge of the MDP dynamics, which is not available in the typical reinforcement learning framework. Moreover, the update involves an optimization with respect to all possible actions which can be carried out efficiently only in the case of finite action spaces. A drawback of this algorithm is that it does not provide an explicit representation of the optimal policy, which in many control problems is what we are looking for.

2.5.2 Policy Iteration

Policy iteration is an iterative method to approximate both the optimal state-value function V_* and the optimal policy π_* . This algorithm alternates an evaluation step, in which the current policy is evaluated using the state-value function, and an improvement step, in which the policy is improved by acting greedily with respect to the action-value function computed in the evaluation step. In the standard version of policy iteration, the state-value function for the current policy is evaluated starting from an arbitrary function $V_0(s)$ and iteratively applying the Bellman expectation operator in a fixed-point iteration scheme

$$V_{k+1}(s) = B_{\pi_k} V_k(s)$$

This algorithm is guaranteed to converge to V_{π_*} . The new policy is then computed as

$$\pi_{k+1} = \text{greedy}(V_{\pi_k})$$

and we go back to the evaluation step for this new policy. Let us notice that it is not necessary to perfectly evaluate to policy π_n before performing the improvement step and the evaluation procedure can be stopped before convergence. This algorithm suffers from the same problem as above. However it provides the basic structure for most of the value-based reinforcement learning methods. In particular, in *generalized policy iteration* the evaluation step is performed in an arbitrary way which does not necessarily employ the Bellman operator. This scheme is illustrated in Figure 2.2,

Chapter 3

Reinforcement Learning

The standard way to solve Markov decision processes is through dynamic programming, which simply consists in solving the Bellman fixed-point equations discussed in the previous chapter. Following this approach, the problem of finding the optimal policy is transformed into the problem of finding the optimal value function. However, apart from the simplest cases where the MDP has a limited number of states and actions, dynamic programming becomes computationally infeasible. Moreover, this approach requires complete knowledge of the Markov transition kernel and of the reward function, which in many real-world applications might be unknown or too complex to use. [Reinforcement Learning \(RL\)](#) is a subfield of Machine Learning which aims to turn the infeasible dynamic programming methods into practical algorithms that can be applied to large-scale problems. [RL](#) algorithms are based on two key ideas: the first is to use samples to compactly represent the unknown dynamics of the controlled system. The second idea is to use powerful function approximation methods to compactly estimate value functions and policies in high-dimensional state and action spaces. In the following sections, we present the [RL](#) problem in more depth and discuss how it relates to the standard discrete-time stochastic optimal control theory and to the other subfields of machine learning, such as supervised learning. In particular, we will see how RL extends ideas from optimal control theory and stochastic approximation to address the broader goal of artificial intelligence. This quick overview of RL is propaedeutic to the following chapters, where we will present in more detail a particular class of algorithms, called policy gradient methods, which are well suited for continuous action spaces. For a more thorough presentation, the reader may consult [\[99\]](#), [\[101\]](#) or [\[110\]](#).

3.1 The Reinforcement Learning Problem

Reinforcement Learning (RL) is a general class of algorithms in the field of machine learning that allows an agent to learn how to behave in a stochastic and possibly unknown environment, where the only feedback consists of a scalar reward signal. In order to maximize the long-run reward, the agent must learn which actions are the most profitable by trial-and-error. Therefore, RL

algorithms can be seen as computational methods to solve Markov decision processes by directly interacting with the environment, for which a model may or may not be available. Trial-and-error search and a delayed reward signal can be seen as the most characteristic features of reinforcement learning.

Compared to supervised learning, one of the main branches of machine learning, the feedback the learner receives is much less. In supervised learning, the agent is provided with examples of the correct or expected behavior by a knowledgeable external supervisor and the agent's goal is to learn how to replicate these examples as well as possible and possibly generalize this knowledge to new examples. In reinforcement learning, the agent only receives a numerical reward that only gives a partial feedback of the goodness of actions taken. Therefore, this feedback system is evaluative rather than instructive and it is much more difficult for the agent to learn how to behave in uncharted territory without any external guidance.

This particular framework generates some challenges that are not present in other kinds of learning. The first one is the trade-off between exploration and exploitation. In order to maximize his reward, an agent would greedily select actions that have already been tried in the past and found to be effective in producing rewards. However, to find these actions, the agent must also test actions that have not been chosen before in order to evaluate their potential. Clearly, this might result in worse performance in the short-term because the actions might be suboptimal. However, without trying them, the agent might not be able to find possible improvements. Thus, an agent must exploit what is known to obtain rewards but also needs to explore to select better actions in the future. A second challenge is the credit assignment problem. Since rewards might be delayed in time, it will be difficult for the agent to understand which actions are mostly responsible for the outcome.

3.2 Model-Free RL Methods

In Section 2.5, we discussed the policy iteration method for computing an optimal policy for an [Markov Decision Process \(MDP\)](#) in a finite state and action spaces. This algorithm belongs to the class of *model-based* methods, since it requires perfect knowledge of the Markov transition kernel and reward function, which consists of a model of the MDP. RL is primarily concerned with how to obtain an optimal policy when such a model is not available. In this section, we discuss some classes of *model-free* methods which do not rely on the MDP model. The lack of a model generates the need to sample the MDP to gather statistical knowledge about this unknown model. In the control setting, the goal is to approximate the optimal policy, which depends on the optimal value, which in turn depends on the model of the MDP, as shown in Figure 3.1. Indeed, we have already seen that a policy which is greedy with respect to the optimal action-value function Q_* , namely for which holds

$$\int_{\mathbb{A}} \pi_*(s, a) Q_*(s, a) da = \sup_a Q_*(s, a) \quad (3.1)$$




Figure 3.1: Solution process for the control problem.

is optimal. Therefore, we can derive the following three methodologies that differ in which part of the solution process is approximated

- i) *Model-approximation* algorithms approximate the MDP model and compute an estimate of the optimal policy by dynamic programming.
- ii) *Value-approximation* algorithms use samples to directly approximate V_* or Q_* , from which an estimate of π_* can be derived by acting greedily.
- iii) *Policy-approximation* algorithms directly try to estimate the optimal policy.

It should be noticed that these approaches are not mutually exclusive and can be combined to derive hybrid algorithms. In the following sections we discuss these three classes of algorithm in more detail, following closely [110].

3.2.1 Model Approximation

Model-approximation algorithms approximate the MDP model and compute an optimal policy by dynamic programming, using the techniques discussed in the previous chapter. Since \mathbb{S} , \mathbb{A} and γ are assumed to be known, these methods are based on learning an approximation of the Markov transition kernel \mathcal{P} and the reward function \mathcal{R} . Thanks to the Markov property, these quantities only depend on the current state and action, so that their approximation corresponds to a density estimation problem and a regression problem respectively, which are fairly standard supervised learning problems. Learning the model may not be trivial, but it is in general easier than learning the value of a policy or optimizing the policy directly. The major drawback of model-based algorithms in continuous-state MDPs is that, even if the model is available, it is in general infeasible to compute the value functions by dynamic programming and to extract an optimal policy for all states by acting greedily. Alternatively, a transition model estimate may be used to generate sample trajectories from the MDP, which can then be used to estimate the value function or directly improve

the policy. However, the value function and the policy estimated using these samples can only be as accurate as the learned model, so that in many cases it may be easier to directly approximate the value function or the policy using the methods described below.

3.2.2 Value Approximation

Value-approximation algorithms use samples from the MDP to approximate V_* or Q_* directly and then derive an estimate of the optimal policy by acting greedily with respect to Q_* . Typically, when the state and action spaces are large, the value functions are estimated using a parametric function approximator, whose parameters are iteratively updated given the observed samples. Many reinforcement learning algorithms fall into this category and they can be distinguished based on whether they are on-policy or off-policy and whether they update the value function estimates online or offline. *On-policy* algorithms approximate the state-value function V_π or the action-value function Q_π from samples of the MDP obtained by following the same policy π to be evaluated. Although the optimal policy π_* is initially unknown, such algorithms can eventually approximate the optimal value functions V_* or Q_* from which an approximation of the optimal policy can be derived. On the other hand, *off-policy* algorithms can learn the value of a policy different from the one used for obtaining the MDP samples. *Online* algorithms adapt their value approximation after each observed sample while *Offline* algorithms operate on batches of samples. Online algorithms typically require less computation per sample but their convergence is slower.

3.2.3 Policy Approximation

Policy-approximation algorithms only store a policy and update this policy to maximize a given performance measure and eventually converge to the optimal policy. Since these algorithms only use an estimate of the optimal policy and do not rely on a value function approximation, they are also referred to as *direct policy-search* or *actor-only* algorithms. Algorithms that store both a policy and a value function are commonly known as *actor-critic* algorithms [100], [55]. Policy gradient algorithms, the most studied policy-approximation methods, will be discussed in much more detail in the next chapter.

Chapter 4

Risk-Neutral Policy Gradient

Policy gradient methods directly store and iteratively improve a parametric approximation of the optimal policy. This policy is commonly referred to as an *actor* and methods that directly approximate the policy without exploiting an approximation of the optimal value function are called *actor-only*. Algorithms that combine an approximation of the optimal policy with an approximation of the value function are commonly called *actor-critic* methods.

As discussed in Chapter 2, an optimal policy can be obtained by simply acting greedily with respect to the optimal action-value function. However, in large or continuous action spaces, this leads to a complex optimization problem that is computationally expensive to solve. Therefore, it can be beneficial to store an explicit estimation of the optimal policy from which we can select actions. Policy gradient methods have other advantages compared to standard value-based approaches. In many applications, a good policy has a more compact representation than the value function so that it may be easier to approximate. Moreover, the policy parameterization can be chosen so as to be relevant for the task and to directly incorporate prior knowledge. Another advantage is that these methods can learn stochastic policies and not only deterministic ones, which might be useful in multi-player frameworks or in partially observable environments. Finally, these methods are guaranteed to converge at least to a local optimum, which may be good enough in practice. On the other hand, policy gradient methods are typically characterized by a large variance which may hinder the converge speed.

In the following sections, we present the basics of policy gradient methods closely following the thorough overview provided in [82]. Then, we discuss some state-of-the-art policy gradient algorithms. In particular, we focus on the risk-neutral framework postponing the analysis of the risk-sensitive framework to the next chapter.

4.1 Basics of Policy Gradient Methods

In policy gradient methods, the optimal policy is approximated using a parametrized policy $\pi : \mathcal{S} \times \mathcal{A} \times \Theta \rightarrow \mathbb{R}$ such that, given a parameter vector $\theta \in \Theta \subseteq \mathbb{R}^{D_\theta}$,

Algorithm 4.1 General setup for a policy gradient algorithm.

Input: Initial parameters θ_0 , learning rate $\{\alpha_k\}$

Output: Approximation of the optimal policy π_*

- 1: **repeat**
 - 2: Approximate policy gradient $\hat{g} \approx \nabla_{\theta} J(\theta_k)$
 - 3: Update parameters using gradient ascent $\theta_{k+1} = \theta_k + \alpha_k \hat{g}$
 - 4: $k \leftarrow k + 1$
 - 5: **until** converged
-

$\pi(s, B; \theta) = \pi_{\theta}(s, B)$ gives the probability of selecting an action in $B \in \mathcal{A}$ when the system is in state $s \in \mathbb{S}$. The general goal of policy optimization in reinforcement learning is to optimize the policy parameters $\theta \in \Theta$ so as to maximize a certain objective function $J : \Theta \rightarrow \mathbb{R}$

$$\theta^* = \arg \max_{\theta \in \Theta} J(\theta)$$

In the following, we will focus on gradient-based and model-free methods that exploit the sequential structure of the reinforcement learning problem. The idea of policy gradient algorithms is to update the policy parameters using the gradient ascent direction of the objective function

$$\theta_{k+1} = \theta_k + \alpha_k \nabla_{\theta} J(\theta_k) \quad (4.1)$$

where $\{\alpha_k\}_{k \geq 0}$ is a sequence of learning rates. Typically, the gradient of the objective function is not known and its approximation is the key component of every policy gradient algorithm, a general setup of which is outlined in Algorithm 4.1. It is a well-known result from stochastic optimization [57] that, if the gradient estimate is unbiased and the learning rates satisfy the *Robbins-Monro conditions*

$$\sum_{k=0}^{\infty} \alpha_k = \infty \quad \sum_{k=0}^{\infty} \alpha_k^2 < \infty \quad (4.2)$$

the learning process is guaranteed to converge at least to a local optimum of the objective function. Before describing various methods to approximate the gradient, let us give an overview of some standard objective functions and the situation in which they are commonly used.

4.2 Risk-Neutral Objective Functions

In the risk-neutral setting, the agent is only interested in maximizing his reward, without considering the risk he needs to take on to achieve it. In an episodic environment where the system always starts from an initial state s_0 , the typical objective function is the start value.

Definition 4.2.1 (Start Value). *The start value is the expected return that can be obtained starting from the start state $s_0 \in \mathbb{S}$ and following policy π_{θ}*

$$J_{start}(\theta) = V_{\pi_{\theta}}(s_0) = \mathbb{E}_{\pi_{\theta}} [G_0 | S_0 = s_0] \quad (4.3)$$

In a continuing environment, where no terminal state exists and the task might go on forever, it is common to use either the average value or the average reward per time step.

Definition 4.2.2 (Average Value). *The average value is the expected value that can be obtained following policy π_θ*

$$J_{avV}(\theta) = \mathbb{E}_{S \sim d^\theta} [V_{\pi_\theta}(S)] = \int_{\mathbb{S}} d^\theta(s) V_{\pi_\theta}(s) ds \quad (4.4)$$

where d^θ is the stationary distribution of the Markov chain induced by π_θ .

Definition 4.2.3 (Average Reward per Time Step). *The average reward per time step is the expected reward that can be obtained over a single time step by following policy π_θ*

$$J_{avR}(\theta) = \rho(\theta) = \mathbb{E}_{\substack{S \sim d^\theta \\ A \sim \pi_\theta}} [\mathcal{R}(S, A)] = \int_{\mathbb{S}} d^\theta(s) \int_{\mathbb{A}} \pi_\theta(s, a) \mathcal{R}(s, a) da ds \quad (4.5)$$

where d^θ is the stationary distribution of the Markov chain induced by π_θ .

4.3 Finite Differences

The most straightforward way to estimate the objective function gradient consists in replacing the partial derivatives with the corresponding finite differences. In other words, the k -th gradient component is approximated by

$$\frac{\partial J(\theta)}{\partial \theta_k} \approx \frac{J(\theta + \epsilon e_k) - J(\theta)}{\epsilon} \quad (4.6)$$

where ϵ is a small constant and $e_k \in \mathbb{R}^{D_\theta}$ is the k -th element of the canonical basis. Hence, to compute the gradient it is sufficient to estimate the objective function for $D_\theta + 1$ different parameters combinations. Since the objective function is unknown, it must be estimated from sample trajectories simulated following the policies associated to the parameterizations appearing in the finite differences. For this reason, this method is computationally demanding and the gradient estimate obtained in this way is extremely noisy, which may slow down or even prevent the convergence of the algorithm. Still, this approach is sometimes effective and works for arbitrary, even non-differentiable, policies. Moreover, finite differences are often used to check gradient estimates during debugging.

4.4 Likelihood Ratio Methods

In this section, we describe several methods based on the *likelihood ratio* technique of stochastic optimization. Let Z be a random variable with a parametric probability density p_θ and assume that p_θ is known, explicitly computable and

differentiable with respect to the parameters. The likelihood ratio technique is used to compute the following gradient

$$\nabla_{\theta} \mathbb{E}_{Z \sim p_{\theta}} [f(Z)] = \nabla_{\theta} \int p_{\theta}(z) f(z) dz$$

This type of problem appears in many domains, such as when computing the greeks of a derivative product in computational finance [79]. Under some regularity assumptions on the function f and on the probability distribution p_{θ} , this gradient can be rewritten in a more amenable way

$$\begin{aligned} \nabla_{\theta} \mathbb{E}_{Z \sim p_{\theta}} [f(Z)] &= \int \nabla_{\theta} p_{\theta}(z) f(z) dz \\ &= \int p_{\theta}(z) \frac{\nabla_{\theta} p_{\theta}(z)}{p_{\theta}(z)} f(z) dz \\ &= \int p_{\theta}(z) \nabla_{\theta} \log p_{\theta}(z) f(z) dz \\ &= \mathbb{E}_{Z \sim p_{\theta}} [\nabla_{\theta} \log p_{\theta}(Z) f(Z)] \end{aligned}$$

Thus, the likelihood ratio technique simply consists in the following equality

Proposition 4.4.1 (Likelihood Ratio Technique).

$$\nabla_{\theta} \mathbb{E}_{Z \sim p_{\theta}} [f(Z)] = \mathbb{E}_{Z \sim p_{\theta}} [\nabla_{\theta} \log p_{\theta}(Z) f(Z)] \quad (4.7)$$

The quantity $\nabla_{\theta} \log p_{\theta}(Z)$ is usually called likelihood score. This result provides a simple way to approximate the gradient using a Monte Carlo estimate: let $\{Z^{(m)}\}_{m=1}^M$ be i.i.d. samples from p_{θ} , then

$$\nabla_{\theta} \mathbb{E}_{Z \sim p_{\theta}} [f(Z)] \approx \frac{1}{M} \sum_{m=1}^M \nabla_{\theta} \log p_{\theta}(Z^{(m)}) f(Z^{(m)}) \quad (4.8)$$

As we will see in the following sections, this technique provides the basis for some of the most common policy gradient algorithms in the reinforcement learning literature.

4.4.1 Monte Carlo Policy Gradient

Let $h = \{(s_t, a_t)\}_{t \geq 0} \in \mathbb{H}$ be a given trajectory of the MDP and let us denote by $p_{\theta}(h) = \mathbb{P}_{\pi_{\theta}}(H = h)$ the probability of obtaining this trajectory under policy π_{θ} . Let $G(h)$ denote the expected return obtained on trajectory h

$$G(h) = \mathbb{E}[G_0 | H = h] = \sum_{t=0}^{\infty} \gamma^t \mathcal{R}(s_t, a_t)$$

Let us consider the start value objective function, which can be rewritten as an expectation over all possible trajectories

$$J_{\text{start}}(\theta) = \mathbb{E}_{H \sim p_{\theta}} [G(H)]$$

Applying the likelihood ratio technique, we obtain

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{H \sim p_{\theta}} [\nabla_{\theta} \log p_{\theta}(H) G(H)] \quad (4.9)$$

The crucial point is that $\nabla_{\theta} \log p_{\theta}(H)$ can be computed without knowledge of the transition probability kernel \mathcal{P} . Indeed, by a recursive application of the Markov property, we have

$$p_{\theta}(h) = \mathbb{P}(S_0 = s_0) \prod_{t=0}^{\infty} \pi_{\theta}(s_t, a_t) \mathcal{P}(s_t, a_t, s_{t+1})$$

taking the logarithm yields

$$\log p_{\theta}(h) = \log \mathbb{P}(S_0 = s_0) + \sum_{t=0}^{\infty} \log \pi_{\theta}(s_t, a_t) + \sum_{t=0}^{\infty} \log \mathcal{P}(s_t, a_t, s_{t+1})$$

Since the only term depending on the parameters θ is the policy term,

$$\nabla_{\theta} \log p_{\theta}(H) = \sum_{t=0}^{\infty} \nabla_{\theta} \log \pi_{\theta}(s_t, a_t) \quad (4.10)$$

Therefore, we do not need the transition model to compute the $\nabla_{\theta} \log p_{\theta}(H)$. Moreover, it is easy to prove that

$$\mathbb{E}_{H \sim p_{\theta}} [\nabla_{\theta} \log p_{\theta}(H)] = 0 \quad (4.11)$$

Hence, a constant baseline $b \in \mathbb{R}$ can always be subtracted in Eq. 4.9 without changing the gradient value

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{H \sim p_{\theta}} [\nabla_{\theta} \log p_{\theta}(H) (G(H) - b)] \quad (4.12)$$

Intuitively, the baseline b should measure the expected return under the policy, so that better (resp. worse) than average returns would produce a positive (resp. negative) gradient. A simple approach is to use a moving average of the returns observed while improving the policy. A more elaborate approach is to compute the baseline that minimizes the variance of the gradient estimator.

In an episodic environment, we can derive an estimate of the gradient by sampling M trajectories $h^{(m)} = \{(s_t^{(m)}, a_t^{(m)})\}_{t=0}^{T^{(m)}}$ under policy π_{θ} and by approximating the expected value via Monte Carlo

$$\hat{g}_{\text{RF}} = \frac{1}{M} \sum_{m=1}^M \left[\sum_{i=0}^{T^{(m)}} \nabla_{\theta} \log \pi_{\theta}(s_i^{(m)}, a_i^{(m)}) \right] \left[\sum_{j=0}^{T^{(m)}} \gamma^j r_{j+1}^{(m)} - b \right] \quad (4.13)$$

This method, synthetized in Algorithm 4.2, is known in the literature as the REINFORCE algorithm and is guaranteed to converge to the true gradient at a

Algorithm 4.2 Episodic REINFORCE policy gradient estimate**Input:** Policy parameterization θ , number of trajectories M **Output:** REINFORCE policy gradient estimate $\hat{g}_{RF} \approx \nabla_{\theta} J(\theta)$

- 1: Sample M trajectories of the MDP following policy π_{θ}
- 2: For all k , Compute the optimal baseline b_k^* according to Eq. (4.15)
- 3: Compute \hat{g}_{RF} according to Eq. (4.13)

pace of $O(M^{-1/2})$. In practice, we can obtain an approximation of the gradient using only one sample which leads to a stochastic gradient ascent method

$$\hat{g}_{SRF} = \left[\sum_{i=0}^T \nabla_{\theta} \log \pi_{\theta}(s_i, a_i) \right] \left[\sum_{j=0}^T \gamma^j r_{j+1} - b \right] \quad (4.14)$$

This method is very easy and works well on many problems. However, the gradient estimate is characterized by a large variance which can hamper the convergence rate of the algorithm. A first approach to address this issue is to optimally set the baseline to reduce the gradient variance.

4.4.1.1 Optimal Baseline

A standard variance reduction technique consists in setting the baseline so as to minimize the gradient estimate variance. More in detail, the optimal baseline for the k -th gradient component \hat{g}_k solves

$$b_k^* = \arg \min_b \text{Var}(\hat{g}_k)$$

It is easy to show that

$$b_k^* = \frac{\mathbb{E} [G(H) (\partial_{\theta_k} \log p_{\theta}(H))^2]}{\mathbb{E} [(\partial_{\theta_k} \log p_{\theta}(H))^2]}$$

which can be approximated by

$$\hat{b}_k^* = \frac{\sum_{m=1}^M \left[\sum_{i=0}^{T^{(m)}} \partial_{\theta_k} \log \pi_{\theta} \left(s_i^{(m)}, a_i^{(m)} \right) \right]^2 \sum_{j=0}^{T^{(m)}} \gamma^j r_{j+1}^{(m)}}{\sum_{m=1}^M \left[\sum_{i=0}^{T^{(m)}} \partial_{\theta_k} \log \pi_{\theta} \left(s_i^{(m)}, a_i^{(m)} \right) \right]^2} \quad (4.15)$$

4.4.2 GPOMDP

The Monte Carlo policy gradient estimate is typically characterized by a large variance, which may slow down the method's convergence. To improve the estimate, it is sufficient to notice that future actions and past rewards are independent, unless the policy has been changed. Therefore, combining this observation with Eq. (4.11) yields

$$\mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(S_t, A_t) \mathcal{R}(S_s, A_s)] = 0 \quad \forall t > s$$

Plugging this equation in Eq. (4.12) leads to the well-known GPOMDP algorithm [10] for generating an estimate of the gradient of the objective function

$$\hat{g}_{\text{GPOMDP}} = \frac{1}{M} \sum_{m=1}^M \sum_{i=0}^{T^{(m)}} \nabla_{\theta} \log \pi_{\theta}(s_i^{(m)}, a_i^{(m)}) \left(\sum_{j=i}^{T^{(m)}} \gamma^j r_{j+1}^{(m)} - b \right) \quad (4.16)$$

By removing almost half of the cross-products, this estimate typically has a smaller variance than the trivial REINFORCE policy gradient and can be further reduced by computing an optimal baseline. In the following sections we will see how this algorithm can be easily derived from a more general result: the policy gradient theorem.

4.4.3 Stochastic Policies

The likelihood ratio technique is a powerful tool and it allows to derive some powerful policy gradient algorithms. However, this approach only works if the policy is stochastic. In most cases this is not a big problem, since stochastic policies are needed anyway to ensure sufficient exploration of the state-action space. Moreover, stochastic policies might be beneficial in partially observable MDP to avoid state aliasing and in adversarial learning where the Pareto optimal strategies of different players are randomized, such as in the classical game “rock-paper-scissors”. We introduce now two standard examples of stochastic policies for discrete and continuous action spaces respectively.

4.4.3.1 Boltzmann Exploration Policy

In discrete action spaces, the Boltzmann exploration policy, also known as soft-max policy, is a common choice. In state $s \in \mathbb{S}$, this policy selects an action $a \in \mathbb{A}$ with probability

$$\pi_{\theta}(s, a) = \frac{e^{\theta^T \Phi(s, a)}}{\sum_{b \in \mathbb{A}} e^{\theta^T \Phi(s, b)}} \quad (4.17)$$

where $\Phi(s, a) \in \mathbb{R}^{D_{\theta}}$ is a given feature vector corresponding to state s and action a . the likelihood score for this policy is thus given by

$$\nabla_{\theta} \log \pi_{\theta}(s, a) = \Phi(s, a) - \sum_{b \in \mathbb{A}} \pi_{\theta}(s, b) \Phi(s, b) \quad (4.18)$$

4.4.3.2 Gaussian Exploration Policy

In continuous action spaces, a Gaussian exploration policy is commonly used. According to this policy, in a state s , actions are sampled from a Gaussian distribution with a parametric state-dependent mean $\mu_{\psi}(s) \in \mathbb{R}^{D_a}$, with $\psi \in \mathbb{R}^{D_{\psi}}$, and a covariance matrix $\Sigma \in \mathbb{R}^{D_a \times D_a}$. The policy parameters consist of $\theta = \{\psi, \Sigma\}$ and the likelihood score are thus given by

$$\nabla_{\psi} \log \pi_{\theta}(s, a) = \left(\frac{\partial \mu_{\psi}(s)}{\partial \psi} \right)^T \Sigma^{-1} (a - \mu_{\psi}(s)) \quad (4.19)$$

$$\nabla_{\Sigma} \log \pi_{\theta}(s, a) = \frac{1}{2} \left[\Sigma^{-1} (a - \mu_{\psi}(s)) (a - \mu_{\psi}(s))^T \Sigma^{-1} - \Sigma^{-1} \right] \quad (4.20)$$

where $\frac{\partial \mu_{\psi}(s)}{\partial \psi}$ denotes the Jacobian matrix of μ_{ψ} with respect to ψ .

4.4.4 Policy Gradient with Parameter Exploration

In Monte Carlo Policy Gradient, trajectories are generated by sampling at each time step an action according to a stochastic policy π_{θ} and the objective function gradient is estimated by differentiating the policy with respect to the parameters. However, sampling an action from the policy at each time step leads to a large variance in the sampled histories and therefore in the gradient estimate, which can in turn slow down the convergence of the learning process. To address this issue, in [92] the authors propose the *policy gradient with parameter-based exploration* (PGPE) method, in which the search in the policy space is replaced with a direct search in the model parameter space. We start by presenting the episodic case and we will later extend this approach to the infinite horizon setting.

4.4.4.1 Episodic PGPE

Given an episodic MDP, PGPE considers a deterministic controller $F : \mathbb{S} \times \Theta \rightarrow \mathbb{A}$ that, given a set of parameters $\theta \in \Theta \subseteq \mathbb{R}^{D_{\theta}}$, maps a state $s \in \mathbb{S}$ to an action $a = F(s; \theta) = F_{\theta}(s) \in \mathbb{A}$. The policy parameters are drawn from a probability distribution p_{ξ} , with hyper-parameters $\xi \in \Xi \subseteq \mathbb{R}^{D_{\xi}}$. Combining these two hypotheses, the agent follows a stochastic policy π_{ξ} defined by

$$\forall B \in \mathcal{A}, \pi_{\xi}(s, B) = \pi(s, B; \xi) = \int_{\Theta} p_{\xi}(\theta) \mathbb{1}_{F_{\theta}(s) \in B} d\theta$$

The advantage of this approach is that the controller is deterministic and therefore the actions do not need to be sampled at each time step, with a consequent reduction of the gradient estimate variance. Indeed, It is sufficient to sample the parameters θ once at the beginning of the episode and then generate an entire trajectory following the deterministic policy F_{θ} . As an additional benefit, the parameter gradient is estimated by direct parameter perturbations, without having to backpropagate any derivatives, which allows to use non-differentiable controllers.

The hyper-parameters ξ will be updated by following the gradient ascent direction of the gradient of the expected reward, which can be rewritten as

$$J(\xi) = \mathbb{E}_{\substack{\theta \sim p_{\xi} \\ H \sim p_{\theta}}} [G(H)] = \int_{\Theta} \int_{\mathbb{H}} p_{\xi}(\theta, h) G(h) dh d\theta \quad (4.21)$$

By remarking that h is conditionally independent from ξ given θ , so that $p_{\xi}(\theta, h) = p_{\xi}(\theta)p_{\theta}(h)$, and applying the likelihood ratio technique, we obtain

$$\nabla_{\xi} J(\xi) = \mathbb{E}_{\substack{\theta \sim p_{\xi} \\ H \sim p_{\theta}}} [\nabla_{\xi} \log p_{\xi}(\theta) G(H)] \quad (4.22)$$

Algorithm 4.3 Episodic PGPE algorithm**Input:** Initial hyper-parameters ξ_0 , learning rate $\{\alpha_k\}$ **Output:** Approximation of the optimal policy $F_{\xi^*} \approx \pi_*$

```

1: repeat
2:   for  $m = 1, \dots, M$  do
3:     Sample controller parameters  $\theta^{(m)} \sim p_{\xi_k}$ 
4:     Sample trajectory  $h^{(m)} = \{(s_t^{(m)}, a_t^{(m)})\}_{t \geq 0}$  under policy  $F_{\theta^{(m)}}$ 
5:   end for
6:   Approximate policy gradient  $\nabla_{\xi} J(\xi_k) \approx \hat{g}_{\text{PGPE}}$  using Eq. (4.24)
7:   Update hyperparameters using gradient ascent  $\xi_{k+1} = \xi_k + \alpha_k \hat{g}_{\text{PGPE}}$ 
8:    $k \leftarrow k + 1$ 
9: until converged

```

Again, we can subtract a constant baseline $b \in \mathbb{R}$ from the total return

$$\nabla_{\xi} J(\xi) = \mathbb{E} [\nabla_{\xi} \log p_{\xi}(\theta) (G(H) - b)] \quad (4.23)$$

In an episodic environment, the gradient can be approximated via Monte Carlo by first drawing M samples $\theta^{(m)} \sim p_{\xi}$ and then, for each combination of parameters, generating a trajectory $h^{(m)} = \{(s_t^{(m)}, a_t^{(m)})\}_{t \geq 0}$ where actions are selected according to the deterministic controller $F_{\theta^{(m)}}$. This leads to the following estimate

$$\hat{g}_{\text{PGPE}} = \frac{1}{M} \sum_{m=1}^M \nabla_{\xi} \log p_{\xi}(\theta^{(m)}) [G(h^{(m)}) - b] \quad (4.24)$$

In order to further reduce the estimate variance, an optimal baseline can be computed similarly to the REINFORCE case [113]. The episodic PGPE algorithm obtained in this way is reported in Algorithm 4.3. In the following paragraphs, we discuss some possible choices for the policy parameters distribution p_{ξ} , which is the last component of the algorithm.

Independent Gaussian Parameter Distribution A simple approach is to assume that all the components of the parameter vector θ are independent and normally distributed with mean μ_i and variance σ_i^2 , i.e. $\theta_i \sim \mathcal{N}(\mu_i, \sigma_i^2)$, the gradient with respect to the hyper-parameters $\xi = (\mu_1, \dots, \mu_{D_{\theta}}, \sigma_1, \dots, \sigma_{D_{\theta}})^T$ is given by

$$\begin{aligned} \frac{\partial \log p_{\xi}(\theta)}{\partial \mu_i} &= \frac{\theta_i - \mu_i}{\sigma_i^2} \\ \frac{\partial \log p_{\xi}(\theta)}{\partial \sigma_i} &= \frac{(\theta_i - \mu_i)^2 - \sigma_i^2}{\sigma_i^3} \end{aligned} \quad (4.25)$$

Using a constant learning rate $\alpha_i = \alpha \sigma_i^2$, the gradient updates takes the following form

$$\begin{aligned} \mu_i^{k+1} &= \mu_i^k + \alpha [G(h) - b] (\theta_i - \mu_i) \\ \sigma_i^{k+1} &= \sigma_i^k + \alpha [G(h) - b] \frac{(\theta_i - \mu_i)^2}{\sigma_i} \end{aligned} \quad (4.26)$$

where b can be computed as a moving average of the past returns. Intuitively, if $G(h) > b$ we adjust ξ so as to increase the probability of θ while if $G(h) < b$ we do the opposite.

Gaussian Parameter Distribution A more elaborate approach is to assume a generic dependence among the controller parameters, i.e. $\theta \sim \mathcal{N}(\mu, \Sigma)$. However, if we directly used the covariance matrix Σ as an hyper-parameter, it would be computationally difficult to enforce that it remains well-defined, i.e. symmetric and semidefinite positive, during the gradient ascent iterations. A simpler approach is to parameterize the distribution using the Cholesky factor Σ , i.e. the matrix C such that $\Sigma = C^T C$. This choice has two advantages: first, C makes explicit the $n(n+1)/2$ independent parameters determining the covariance matrix Σ ; in addition, $C^T C$ is by construction a well-defined covariance matrix. Hence, the hyper-parameters are $\xi = \{\mu, C\}$. The likelihood score for this distribution doesn't have a simple expression, but it becomes extremely easy in the Natural version of PGPE which will be discussed in the next sections.

Symmetric Sampling and Gain Normalization In some settings, comparing the gain with a baseline can be misleading. In their original work, the authors propose a symmetric sampling technique similar to antithetic variates that further improves the convergence of the method. More in detail, a more robust gradient estimate can be obtained by measuring the difference in reward between two symmetric samples on either side of the current mean. That is, we sample a random perturbation $\epsilon \sim \mathcal{N}(0, \Sigma)$, where $\Sigma = \text{diag}(\sigma_1^2, \dots, \sigma_{D_\theta}^2)$, and we define $\theta^+ = \mu + \epsilon$ and $\theta^- = \mu - \epsilon$. Denoting by G^+ (resp. G^-) the gains obtained on the trajectory associated to θ^+ (resp. θ^-), the objective function gradient can be approximated with

$$\begin{aligned}\nabla_{\mu_i} J(\xi) &\approx \frac{\epsilon_i (G^+ - G^-)}{2\sigma_i^2} \\ \nabla_{\sigma_i} J(\xi) &\approx \frac{\epsilon_i^2 - \sigma_i^2}{\sigma_i^3} \left(\frac{G^+ + G^-}{2} - b \right)\end{aligned}\tag{4.27}$$

Hence, by choosing $\alpha_i^k = 2\alpha\sigma_i^2$, we have the following update rules

$$\begin{aligned}\mu_i^{k+1} &= \mu_i^k + \alpha\epsilon_i (G^+ - G^-) \\ \sigma_i^{k+1} &= \sigma_i^k + \alpha \frac{\epsilon_i^2 - \sigma_i^2}{\sigma_i} (G^+ + G^- - 2b)\end{aligned}\tag{4.28}$$

In addition, the authors propose to normalize the gains in order to make the updates independent of the scale of the rewards. For instance, we could modify the hyperparameters updates as follows

$$\begin{aligned}\mu_i^{k+1} &= \mu_i^k + \alpha\epsilon_i \frac{(G^+ - G^-)}{2m - G^+ - G^-} \\ \sigma_i^{k+1} &= \sigma_i^k + \alpha \frac{\epsilon_i^2 - \sigma_i^2}{\sigma_i} \frac{(G^+ + G^- - 2b)}{m - b}\end{aligned}\tag{4.29}$$

where m might be the maximum gain the agent can receive, if known, or alternatively the maximum gain achieved so far. Symmetric sampling and gain normalization can drastically improve the gradient estimate quality and consequently the convergence time.

4.4.4.2 Infinite Horizon PGPE

While in the episodic PGPE the parameters θ are sampled only at the beginning of each episode, in *Infinite Horizon PGPE* (IHPGPE) [91] the parameters and learning are carried out simultaneously, while interacting with the environment. Let $0 < \varepsilon < 1$ the probability of updating the policy parameters, the parameters θ_t can be sampled consecutively as follows

$$p_{\xi}(\theta_{i,t+1}) = \varepsilon \mathcal{N}(\mu_{i,t}, \sigma_{i,t}^2) + (1 - \varepsilon) \delta_{\theta_{i,t}} \quad (4.30)$$

In practice, ε should be chosen so that the expected frequency of changing a single parameter is coherent with the typical episode length in the episodic framework. Alternatively, one could sample all the parameters at a certain time step simultaneously

$$p_{\xi}(\theta_{t+1}) = \varepsilon \mathcal{N}(\mu_t, \Sigma_t) + (1 - \varepsilon) \delta_{\theta_t} \quad (4.31)$$

This is equivalent to splitting the state-action space into artificial episodes. However, updating parameters asynchronously changes the policy only slightly thus introducing less noise in the process. Again, parameters can be updated at every time step by gradient ascent

$$\begin{aligned} \mu_{i,t+1} &= \mu_{i,t} + \alpha [G_t(h) - b] (\theta_{i,t} - \mu_{i,t}) \\ \sigma_{i,t+1} &= \sigma_{i,t} + \alpha [G_t(h) - b] \frac{(\theta_{i,t} - \mu_{i,t})^2}{\sigma_{i,t}} \end{aligned} \quad (4.32)$$

Similarly to the episodic case, we can improve the gradient estimate by symmetric sampling and gain normalization.

4.5 Risk-Neutral Policy Gradient Theorem

In the previous sections, we saw that the core of a policy gradient algorithm consist in the approximation of the objective function gradient. In this section we present the *policy gradient theorem* [100], which shows that the gradient can be rewritten in a form suitable for estimation from experience aided by an approximate action-value or advantage function. First, we will prove the theorem in the risk-neutral setting for which it was originally proposed. In particular, we will see how the GPOMDP algorithm can be easily derived by this result via a Monte Carlo approximation. Moreover, we will derive a class of actor-critic algorithms [55] that, in addition to a parametric approximation of the policy, also exploit an approximation of the action-value function or of an

advantage function to reduce the variance of the gradient estimate. In particular, we review the powerful idea of compatible function approximation [100] which assures the convergence to a local optimum of the objective function. Finally, we discuss the natural policy gradient idea [51], which forms the basis of many state-of-the-art algorithms. The extension to the risk-sensitive setting will be done in the next section.

4.5.1 Theorem Statement and Proof

Theorem 4.5.1 (Risk-Neutral Policy Gradient). *Let π_θ be a differentiable policy. The policy gradient for the average reward formulation is given by*

$$\nabla_\theta \rho(\theta) = \mathbb{E}_{\substack{S \sim d^\theta \\ A \sim \pi_\theta}} [\nabla_\theta \log \pi_\theta(S, A) Q_\theta(S, A)] \quad (4.33)$$

where d^θ is the stationary distribution of the Markov chain induced by π_θ . The policy gradient for the start value formulation is given by

$$\nabla_\theta J_{\text{start}}(\theta) = \mathbb{E}_{\substack{S \sim d_\gamma^\theta(s_0, \cdot) \\ A \sim \pi_\theta}} [\nabla_\theta \log \pi_\theta(S, A) Q_\theta(S, A)] \quad (4.34)$$

where $d_\gamma^\theta(s_0, \cdot)$ is the γ -discounted visiting distribution over states starting from the initial state s_0 and following policy π_θ

$$d_\gamma^\theta(s, x) = \sum_{k=0}^{\infty} \gamma^k \mathcal{P}_\theta^{(k)}(s, x) \quad (4.35)$$

Proof. We first prove the result for the average-reward formulation and then for the start state formulation. From the basic relation between state-value function and action-value function, we have

$$\begin{aligned} \nabla_\theta V_\theta(s) &= \nabla_\theta \int_{\mathbb{A}} \pi_\theta(s, a) Q_\theta(s, a) da \\ &= \int_{\mathbb{A}} [\nabla_\theta \pi_\theta(s, a) Q_\theta(s, a) + \pi_\theta(s, a) \nabla_\theta Q_\theta(s, a)] da \end{aligned}$$

Using the Bellman expectation equation for Q_θ

$$\begin{aligned} \nabla_\theta Q_\theta(s, a) &= \nabla_\theta \left[\mathcal{R}(s, a) - \rho_\theta + \int_{\mathbb{S}} \mathcal{P}(s, a, s') V_\theta(s') ds' \right] \\ &= -\nabla_\theta \rho_\theta + \int_{\mathbb{S}} \mathcal{P}(s, a, s') \nabla_\theta V_\theta(s') ds' \end{aligned}$$

Hence, plugging in the first equation

$$\nabla_\theta V_\theta(s) = \int_{\mathbb{A}} \nabla_\theta \pi_\theta(s, a) Q_\theta(s, a) da - \nabla_\theta \rho_\theta + \int_{\mathbb{A}} \pi_\theta(s, a) \int_{\mathbb{S}} \mathcal{P}(s, a, s') \nabla_\theta V_\theta(s') ds'$$

Integrating both sides with respect to the stationary distribution d^θ and noting that, because of stationarity,

$$\int_{\mathbb{S}} d^\theta(s) \int_{\mathbb{A}} \pi(s, a) \int_{\mathbb{S}} \mathcal{P}(s, a, s') \nabla_\theta V(s') ds' da ds = \int_{\mathbb{S}} d^\theta(s) \nabla_\theta V_\theta(s) ds$$

we obtain the result

$$\begin{aligned} \nabla_\theta \rho_\theta &= \int_{\mathbb{S}} d^\theta(s) \int_{\mathbb{A}} \nabla_\theta \pi_\theta(s, a) Q_\theta(s, a) da ds \\ &= \int_{\mathbb{S}} d^\theta(s) \int_{\mathbb{A}} \pi_\theta(s, a) \nabla_\theta \log \pi_\theta(s, a) Q_\theta(s, a) da ds \\ &= \mathbb{E}_{\substack{S \sim d^\theta \\ A \sim \pi_\theta}} [\nabla_\theta \log \pi_\theta(S, A) Q_\theta(S, A)] \end{aligned}$$

Let us now prove the theorem for the start state formulation. The first step is exactly the same. Hence, using the Bellman expectation equation for Q_θ

$$\begin{aligned} \nabla_\theta Q_\theta(s, a) &= \nabla_\theta \left[\mathcal{R}(s, a) + \gamma \int_{\mathbb{S}} \mathcal{P}(s, a, s') V_\theta(s') ds' \right] \\ &= \gamma \int_{\mathbb{S}} \mathcal{P}(s, a, s') \nabla_\theta V_\theta(s') ds' \end{aligned}$$

we obtain

$$\begin{aligned} \nabla_\theta V_\theta(s) &= \int_{\mathbb{A}} \left[\nabla_\theta \pi_\theta(s, a) Q_\theta(s, a) + \gamma \int_{\mathbb{S}} \mathcal{P}(s, a, s') \nabla_\theta V_\theta(s') ds' \right] da \\ &= \int_{\mathbb{S}} \sum_{k=0}^{\infty} \gamma^k \mathcal{P}_\theta^{(k)}(s, x) \int_{\mathbb{A}} \nabla_\theta \pi_\theta(x, a) Q_\theta(x, a) da dx \end{aligned}$$

after unrolling $\nabla_\theta V_\theta$ infinite times and denoting by $\mathcal{P}_\theta^{(k)}(s, x)$ the probability of going from state s to state x in k steps under policy π_θ . Defining the γ -discounted visiting distribution of state x starting from state s as

$$d_\gamma^\theta(s, x) = \sum_{k=0}^{\infty} \gamma^k \mathcal{P}_\theta^{(k)}(s, x)$$

we have the result

$$\begin{aligned} \nabla_\theta V_\theta(s) &= \int_{\mathbb{S}} d_\gamma^\theta(s, x) \int_{\mathbb{A}} \nabla_\theta \pi_\theta(x, a) Q_\theta(x, a) da dx \\ &= \mathbb{E}_{\substack{S \sim d_\gamma^\theta(s_0, \cdot) \\ A \sim \pi_\theta}} [\nabla_\theta \log \pi_\theta(S, A) Q_{\pi_\theta}(S, A)] \end{aligned}$$

□

The action-value function is typically unknown and needs to be approximated. As for the REINFORCE algorithm, we can subtract a state-dependent

baseline from the action-value function without changing the value of the expectation. Indeed

$$\begin{aligned}
\mathbb{E}_{\substack{S \sim d^\theta \\ A \sim \pi_\theta}} [\nabla_\theta \log \pi_\theta(S, A) B_\theta(S)] &= \int_{\mathbb{S}} d^\theta(s) \int_{\mathbb{A}} \pi_\theta(s, a) \nabla_\theta \log \pi_\theta(s, a) B_\theta(s) da ds \\
&= \int_{\mathbb{S}} d^\theta(s) B_\theta(s) \int_{\mathbb{A}} \nabla_\theta \pi_\theta(s, a) da ds \\
&= \int_{\mathbb{S}} d^\theta(s) B_\theta(s) \nabla_\theta \underbrace{\int_{\mathbb{A}} \pi_\theta(s, a) da}_{=1} ds = 0
\end{aligned}$$

Hence, we can rewrite the policy gradient for the average reward formulation as

$$\nabla_\theta \rho(\theta) = \mathbb{E}_{\substack{S \sim d^\theta \\ A \sim \pi_\theta}} [\nabla_\theta \log \pi_\theta(S, A) (Q_{\pi_\theta}(S, A) - B_\theta(S))] \quad (4.36)$$

and similarly for the start state formulation. This result can be used as the starting point to derive several policy gradient methods that use different approximation of the action-value function.

4.5.2 GPOMDP

For an episodic MDP, the action-value function can be estimated with the total return obtained on a sample trajectory

$$Q_\theta(s_0, a_0) \approx \sum_{t=0}^{T^{(m)}} \gamma^t r_{t+1}^{(m)}$$

Combining this remark with a Monte Carlo approximation of Eq. (4.36), we obtain the GPOMDP gradient estimate (4.16). Therefore, the GPOMDP algorithm is a simple application of the policy gradient algorithm, which is much more general than the likelihood ratio technique used to derive the algorithm in the previous sections. In the next subsections we present other ways to exploit the policy gradient theorem to design efficient learning algorithms.

4.5.3 Actor-Critic Policy Gradient

A baseline should ideally measure the typical return obtained by an agent in a certain state when following a certain policy. Therefore, it becomes natural to use the state-value function as a benchmark for the action-value function. Eq. (4.36) thus becomes

$$\nabla_\theta \rho(\theta) = \mathbb{E}_{\substack{S \sim d^\theta \\ A \sim \pi_\theta}} [\nabla_\theta \log \pi_\theta(S, A) A_\theta(S, A)] \quad (4.37)$$

where we introduced the advantage function

$$A_\theta(s, a) = Q_\theta(s, a) - V_\theta(s) \quad (4.38)$$

Algorithm 4.4 Generic structure for an online actor-critic algorithm.

Input:

- Initial actor parameters θ_0 ,
- Initial critic parameters ψ_0 ,
- Learning rate $\{\alpha_k\}$

Output: Approximation of the optimal policy $\pi_{\theta^*} \approx \pi_*$

- 1: **repeat**
 - 2: Observe tuple $\langle s_k, a_k, r_{k+1}, s_{k+1} \rangle$ sampled from the MDP.
 - 3: Update critic parameters ψ_{k+1} using a value-based method.
 - 4: Estimate policy gradient as $\hat{g}_k^{\text{AC}} = \nabla_{\theta} \log \pi_{\theta_k}(s_k, a_k) \hat{A}_{\psi_{k+1}}(s_k, a_k)$
 - 5: Update actor by gradient ascent $\theta_{k+1} = \theta_k + \alpha_k \hat{g}_k^{\text{AC}}$.
 - 6: $k \leftarrow k + 1$
 - 7: **until** converged
-

which measures how good is to take action a in state s compared to simply following the policy π_{θ} . The advantage function is unknown and should be estimated from samples of the MDP. *Actor-Critic* algorithms consist in all those methods that employ an approximation of the advantage function or of the value function, also known as critic, to estimate the policy gradient. These methods thus maintain two sets of parameters: a *critic* that updates the action-value function parameters ψ and an *actor* that updates the policy parameters θ in the direction suggested by the critic. The general structure for an online actor-critic algorithm is reported in Algorithm 4.4.

There are two possible approaches to estimate the advantage function: the first one is to estimate both the state-value function $\hat{V}_{\psi} \approx V_{\theta}$ and the action-value function $\hat{Q}_{\xi} \approx Q_{\theta}$ and derive an estimate for the advantage function $\hat{A}_{\psi, \xi} = \hat{Q}_{\xi} - \hat{V}_{\psi} \approx A_{\theta}$. The second approach directly approximates the advantage function $\hat{A}_{\psi} \approx A_{\theta}$ and is usually preferred, since it allows us to maintain only one critic. Actor-critic algorithms typically use a temporal-difference algorithm to update an approximation of the value function $\hat{V}_{\psi} \approx V_{\theta}$, from which we can derive an approximation for the advantage function. Assume that the true state-value function V_{θ} is given. Then the TD error

$$\delta_{\theta} = R - \rho_{\theta} + V_{\theta}(S') - V_{\theta}(S) \quad (4.39)$$

is an unbiased estimate of the advantage function. Indeed

$$\begin{aligned} \mathbb{E} [\delta_{\theta} | S = s, A = a] &= \mathbb{E} [R - \rho_{\theta} + V_{\theta}(S') | S = s, A = a] - V_{\theta}(s) \\ &= Q_{\theta}(s, a) - V_{\theta}(s) \\ &= A_{\theta}(s, a) \end{aligned}$$

By a simple conditioning argument, the policy gradient can be rewritten as

$$\nabla_{\theta} \rho(\theta) = \mathbb{E}_{\substack{S \sim d^{\theta} \\ A \sim \pi_{\theta}}} [\nabla_{\theta} \log \pi_{\theta}(S, A) \delta_{\theta}] \quad (4.40)$$

Algorithm 4.5 TD(λ) policy gradient algorithm.

Input:

- Initial actor parameters θ_0 ,
- Initial critic parameters ψ_0 ,
- Learning rate $\{\alpha_k\}$

Output: Approximation of the optimal policy $\pi_{\theta^*} \approx \pi_*$

- 1: Initialize $k = 0$ and the eligibility trace $e_{-1} = 0$
 - 2: **repeat**
 - 3: Observe tuple $\langle s_k, a_k, r_{k+1}, s_{k+1} \rangle$ sampled from the MDP.
 - 4: Compute TD error $\hat{\delta}_k = r_{k+1} - \hat{\rho}_k + \hat{V}_{\psi_k}(s_{k+1}) - \hat{V}_{\psi_k}(s_k)$
 - 5: Update critic parameters $\psi_{k+1} = \psi_k + \alpha_k \hat{\delta}_k \nabla_{\psi} \hat{V}_{\psi_k}(s_k)$
 - 6: Update eligibility trace $e_{k+1} = \lambda e_k + \nabla_{\theta} \log \pi_{\theta_k}(s_k, a_k)$
 - 7: Update actor parameters $\theta_{k+1} = \theta_k + \alpha_k \hat{\delta}_k e_k$.
 - 8: $k \leftarrow k + 1$
 - 9: **until** converged
-

In practice, we can use an approximate TD error

$$\hat{\delta}_k = r_{k+1} - \hat{\rho}_k + \hat{V}_{\psi_k}(s_{k+1}) - \hat{V}_{\psi_k}(s_k) \quad (4.41)$$

where $\hat{\rho}_k$ is an estimate of the average reward. Hence, we can easily obtain an approximation of the policy gradient by

$$\hat{g}_{\text{TD}(0)}^k = \nabla_{\theta} \log \pi_{\theta}(s_k, a_k) \hat{\delta}_k \quad (4.42)$$

which can be used to update the critic parameter in the gradient ascent direction. On the other hand, the critic parameters can be updated using a TD(0) temporal difference scheme

$$\psi_{k+1} = \psi_k + \alpha_k \hat{\delta}_k \nabla_{\psi} \hat{V}_{\psi_k}(s_k) \quad (4.43)$$

The discussion can be easily extended to more complex value-based methods for estimating the critic, such as the backward-view TD(λ). This method employs eligibility traces to assign credit for the rewards obtained by the agent to all previous states and actions. More formally, the policy parameters update rule becomes

$$\theta_{k+1} = \theta_k + \alpha_k \hat{\delta}_k e_k \quad (4.44)$$

where the eligibility trace e_k is defined by the following recursive equation

$$e_{k+1} = \lambda e_k + \nabla_{\theta} \log \pi_{\theta_k}(s_k, a_k) \quad (4.45)$$

$0 \leq \lambda \leq 1$ is a parameter that manages the amount of bootstrap: for $\lambda = 0$ the method is the standard temporal difference scheme while for $\lambda = 1$ the method is equivalent to Monte Carlo. The complete scheme is reported in Algorithm 4.5.

4.5.4 Compatible Function Approximation

In the previous sections we saw that a critic may reduce the variance of the policy gradient estimate. However, this is achieved at the cost of introducing a bias in the approximation which, in certain cases, may endanger the convergence of the method to a good solution. Luckily, the *compatible function approximation* [100] avoids introducing any bias by a careful choice of the critic. This technique consists in estimating the advantage function with a linear regression on a family of suitable basis functions consisting of the gradient of the likelihood score, i.e.

$$\hat{A}_\theta(s, a) = \psi^T \nabla_\theta \log \pi_\theta(s, a) \quad (4.46)$$

The power of these basis functions becomes apparent in the following theorem

Theorem 4.5.2 (Compatible Function Approximation). *If*

i) the advantage function approximator is compatible to the policy, i.e.

$$\nabla_\psi A_\psi(s, a) = \nabla_\theta \log \pi_\theta(s, a) \quad (4.47)$$

ii) the advantage function parameters ψ minimize the mean square error

$$\theta = \arg \min_{\psi} \mathbb{E}_{\pi_\theta} \left[\left(A_{\pi_\theta}(S, A) - \hat{A}_\psi(S, A) \right)^2 \right] \quad (4.48)$$

Then the policy gradient is exact

$$\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta} \left[\nabla_\theta \log \pi_\theta(S, A) \hat{A}_\psi(S, A) \right] \quad (4.49)$$

Hence, the true advantage function in the policy gradient formula can be replaced by the compatible function approximation without changing the policy gradient value. A linear regression may seem a too simple representation of the advantage function. However, one must bear in mind that this approximation is only used to reduce the variance of the policy gradient estimate and should not be expected to provide an accurate representation of the true-value function. This technique has proved useful in many application.

4.5.5 Natural Policy Gradient

Despite all the advances in the variance reduction techniques, these methods still tend to perform surprisingly poorly. Even when applied to simple examples with rather few states, where the gradient can be determined very accurately, they turn out to be quite inefficient. Typically, one of the reasons of this behavior is the presence of large plateaus in the expected return landscape where the gradients are small and often do not point directly towards the optimal solution. In this context, the steepest ascent with respect to the Fisher information metric, called the *natural policy gradient*, turns out to be significantly more efficient than normal gradients for such plateaus. This technique was first proposed in the reinforcement learning setting in [51] and later applied to actor-critic algorithms in [82]. The following properties of natural policy gradient make it one of the more reliable policy-based methods,

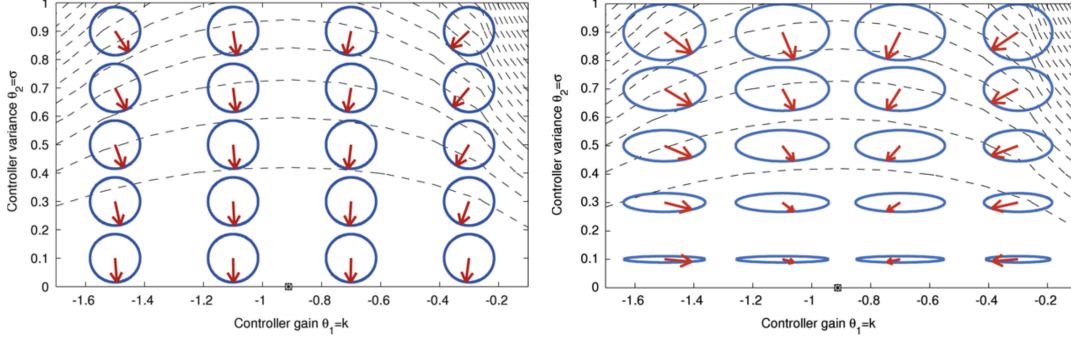


Figure 4.1: “Vanilla” policy gradient (left) vs. natural policy gradient (right). The main difference is how the two approaches punish the change in parameters. This distance is indicated by the blue ellipses in the contour plot while the dashed lines show the expected return. Obtaining a gradient then corresponds to finding a vector pointing from the center of the ellipses to the location with maximum expected return on the ellipse. A vanilla policy gradient (a) considers a change in all parameters as equally distant, thus, it is a search for a maximum on a circle while the natural gradient (b) uses scales determined by the Fisher information which results in a reduction in exploration. The slower reduction in exploration results into a faster convergence to the optimal policy [82].

- i) Convergence to a local minimum is guaranteed.
- ii) By choosing a more direct path to the optimal solution in parameter space, the natural gradient typically has faster convergence and avoids premature convergence of “vanilla gradients”.
- iii) The natural policy gradient can be shown to be covariant, i.e. independent of the coordinate frame chosen for expressing the policy parameters.
- iv) As the natural gradient analytically averages out the influence of the stochastic policy (including the baseline of the function approximator), it requires fewer data points for a good gradient estimate than “vanilla gradients”.

4.5.5.1 Formalism of Natural Policy Gradients

Policy gradient methods improve the policy π_θ by iteratively applying “small” changes $\Delta\theta$ to the policy parameters θ . However, the meaning of “small” is ambiguous. For instance, when working with an Euclidean metric, the size of this update $\|\Delta\theta\| = \sqrt{\Delta\theta^T \Delta\theta}$ and therefore the update size depends on the parameterization of the policy, which often results in unnaturally slow learning even if higher-order gradient methods were employed. This problem poses the question whether we can achieve a covariant gradient descent, i.e. a gradient descent with respect to an invariant measure of the closeness between the current policy and the updated policy based upon the distribution of the paths generated by each of these. Standard measures of distance between probability distributions are the Kullback-Leibler divergence $d_{KL}(p_\theta(h) \parallel p_{\theta+\Delta\theta}(h))$ and the Hellinger distance. These two distances can be approximation in first instance

by a second-order Taylor expansion

$$d_{\text{KL}}(p_\theta(h) \parallel p_{\theta+\Delta\theta}(h)) \approx \frac{1}{2} \Delta\theta^T F_\theta \Delta\theta$$

where F_θ is the Fischer information matrix

$$\begin{aligned} F_\theta &= \int_{\mathbb{H}} p_\theta(h) \nabla_\theta \log p_\theta(h) \nabla_\theta \log p_\theta(h)^T dh \\ &= \mathbb{E}_{H \sim p_\theta} [\nabla_\theta \log p_\theta(H) \nabla_\theta \log p_\theta(H)^T] \end{aligned} \quad (4.50)$$

The goal is to find the optimal update $\Delta\theta$ for the policy parameters so as to maximize the objective function, under the constraint that the new policy must be in a radius ϵ from the previous policy with respect to the Kullback-Leibler divergence, i.e.

$$\begin{cases} \max_{\Delta\theta} J(\theta + \Delta\theta) \approx J(\theta) + \Delta\theta^T \nabla_\theta J(\theta) \\ \text{s.t. } d_{\text{KL}}(p_\theta(h) \parallel p_{\theta+\Delta\theta}(h)) \approx \frac{1}{2} \Delta\theta^T F_\theta \Delta\theta < \epsilon \end{cases}$$

The optimal solution is given by

$$\Delta\theta = \alpha_n F_\theta^{-1} \nabla_\theta J(\theta)$$

with

$$\alpha_n = \sqrt{\frac{\epsilon}{\nabla_\theta J(\theta)^T F_\theta^{-1} \nabla_\theta J(\theta)}}$$

The direction $\tilde{\nabla}_\theta J(\theta) = \Delta\theta/\alpha_n$ is called the natural gradient and learning algorithms that use this gradient instead of the standard one are called natural policy gradient algorithms. The strongest theoretical advantage of this approach is that its performance no longer depends on the parameterization of the policy and it is therefore safe to use for arbitrary policies. In practice, the learning process converges significantly faster in most practical cases and requires less manual parameter tuning of the learning algorithm. The only remaining point to discuss is how to compute the inverse of the Fischer information matrix. This is not an easy task and the matrix will usually need to be estimated from sample trajectories. However, we will see that this matrix can be computed analytically after reformulating the policy gradient theorem for the parameter-based search methods such as PGPE. The formal derivation will be presented in Chapter 6 and this result will lead to the *Natural PGPE* (NPGPE) algorithm [71].

Chapter 5

Risk-Sensitive Policy Gradient

Risk aware decision making plays a crucial role in many fields, such as finance and process control. In this chapter we discuss some policy gradient methods for the risk-sensitive formulation of a sequential decision problem. This extension presents some difficulties which have attracted the interest of many researchers in the last years. In particular, while in the risk-neutral framework the policy gradient theorem represents the keystone for all the learning algorithms, in the risk-sensitive framework the approaches for the episodic setting, the discounted reward and the average reward formulations are quite different. The goal of this chapter is to give an overview of the methods found in the literature and to try to unify them in a way similar to the risk-neutral framework. This chapter and the following represent the main contribution of this thesis to the policy gradient literature.

5.1 Risk-Sensitive Framework

In the risk-sensitive framework, in addition to maximizing the rewards, the agent also wants to control the risk necessary to achieve it. It is thus necessary to introduce a function $\Lambda : \Theta \rightarrow \mathbb{R}$ such that $\Lambda(\theta)$ measures the risk associated with the policy π_θ . In an episodic framework where the system always starts from the same initial state s_0 , the variance of the total return can be used [104].

Definition 5.1.1 (Start Variance). *The start variance is the variance of the return that can be obtained starting from the start state $s_0 \in \mathbb{S}$ and following policy π_θ*

$$\Lambda_{start}(\theta) = \text{Var}_{\pi_\theta}(G_0 \mid S_0 = s_0) = U_{\pi_\theta}(s_0) - V_{\pi_\theta}(s_0)^2 \quad (5.1)$$

In many application one may want control only the downside risk, that is the risk of the actual return being below the expected return or the uncertainty about the magnitude of that difference. This risk may be measured by the semivariance.

Definition 5.1.2 (Semivariance).

$$\Lambda_{down}(\theta) = SVar(\theta) = \mathbb{E}_{\pi_\theta} [\min \{G_0 - \mathbb{E}_{\pi_\theta}[G_0], 0\}^2 \mid S_0 = s_0] \quad (5.2)$$

The square root of this quantity is called semideviation. In a continuing environment, we might consider the long-run variance [83] defined in Section 2.4.

Definition 5.1.3 (Long-Run Variance). *The long-run variance Λ_π under policy π is defined as*

$$\Lambda_{long-run}(\theta) = \lim_{T \rightarrow \infty} \frac{1}{T} \mathbb{E}_{\pi_\theta} \left[\sum_{t=0}^{T-1} (R_{t+1} - \rho(\theta))^2 \right] \quad (5.3)$$

In order to formalize the trade-off between reward and risk from a mathematical point of view, we borrow two standard risk-sensitive performance measures from the financial literature: the mean-variance criterion [70] and the Sharpe ratio [94].

Definition 5.1.4 (Mean-Variance Criterion). *The mean-variance criterion is defined as*

$$J_{MV}(\theta) = J(\theta) - \chi \Lambda(\theta) \quad (5.4)$$

where $\chi > 0$ is a constant that controls the trade-off between reward and risk.

Definition 5.1.5 (Sharpe Ratio). *The Sharpe ratio is defined as*

$$\text{Sh}(\theta) = \frac{J(\theta)}{\sqrt{\Lambda(\theta)}} \quad (5.5)$$

Let us remark that in the financial literature the ratio of the expected return and the semideviation is called Sortino ratio. In a risk-sensitive policy gradient algorithm, we try to approximate the optimal parameters that maximize these objective functions by updating the parameters following the gradient ascent directions. In the average-reward formulation, the ascent directions are given by

$$\nabla_\theta J_{MV}(\theta) = \nabla_\theta \rho(\theta) - \chi \nabla_\theta \Lambda(\theta) \quad (5.6)$$

for the mean-variance criterion, where

$$\nabla_\theta \Lambda(\theta) = \nabla_\theta \eta(\theta) - 2\rho(\theta) \nabla_\theta \rho(\theta) \quad (5.7)$$

and by

$$\nabla_\theta \text{Sh}(\theta) = \frac{\eta(\theta) \nabla_\theta \rho(\theta) - \frac{1}{2} \rho(\theta) \nabla_\theta \eta(\theta)}{\Lambda(\theta) \sqrt{\Lambda(\theta)}} \quad (5.8)$$

for the Sharpe ratio. Hence, in the risk-sensitive framework it is necessary to estimate to different gradients: $\nabla_\theta \rho(\theta)$ and $\nabla_\theta \eta(\theta)$. The equivalent expressions for the episodic setting can be obtained by replacing $\rho(\theta)$ (resp. $\eta(\theta)$) with $V_\theta(s_0)$ (resp. $U_\theta(s_0)$). The estimation of $\nabla_\theta \rho(\theta)$ (resp. $V_\theta(s_0)$) has been the focus of the last chapter. In the next sections, we discuss instead several techniques to estimate the new gradient $\eta(\theta)$ (resp. $U_\theta(s_0)$).

5.2 Monte Carlo Policy Gradient

For an episodic environment, the REINFORCE algorithm can be easily extended to the risk-sensitive framework described above [104]. Indeed, it is sufficient to adapt its derivation for the average return to the second-moment of return

$$U(\theta) = \mathbb{E}_{H \sim p_\theta} [G(H)^2] \quad (5.9)$$

Applying the likelihood ratio technique, we obtain

$$\nabla_\theta U(\theta) = \mathbb{E}_{H \sim p_\theta} [\nabla_\theta \log p_\theta(H) G(H)^2] \quad (5.10)$$

Similarly to the risk-neutral framework, we can introduce a baseline without affecting the value of the gradient

$$\nabla_\theta U(\theta) = \mathbb{E}_{H \sim p_\theta} [\nabla_\theta \log p_\theta(H) (G(H)^2 - b)] \quad (5.11)$$

In an episodic environment, we can estimate the gradient via its Monte Carlo estimate

$$\nabla_\theta U(\theta) \approx \frac{1}{M} \sum_{m=0}^M \nabla_\theta \log p_\theta(h^{(m)}) [G(h^{(m)})^2 - b] \quad (5.12)$$

where $h^{(m)} = \{(s_t^{(m)}, a_t^{(m)})\}_{t=0}^{T^{(m)}}$ are M trajectories of the MDP sampled under policy π_θ . When using a single trajectory, we obtain the following stochastic gradient estimate

$$\nabla_\theta U(\theta) \approx \nabla_\theta \log p_\theta(h) [G(h)^2 - b] \quad (5.13)$$

Again the baseline can be set so as to minimize the variance of the gradient estimate

$$\hat{b}_k^* = \frac{\sum_{m=1}^M [\partial_{\theta_k} \log p_\theta(h^{(m)})]^2 G(h^{(m)})^2}{\sum_{m=1}^M [\partial_{\theta_k} \log p_\theta(h^{(m)})]^2} \quad (5.14)$$

Combining this estimate with the one for the average return discussed in Section 4.4.1 yields the risk-sensitive Monte Carlo Policy Gradient method, which is outlined in Algorithm 5.1.

In an episodic setting, as long as we can rewrite the objective function as expected values on all possible trajectories of the MDP, the likelihood ratio technique yields an analytical expression for its gradient. Hence, this approach can be generalized to more complex measures of risk commonly used in finance, such as the semivariance introduced above. Indeed, the semivariance can be rewritten as

$$\text{SVar}(\theta) = \mathbb{E}_{H \sim p_\theta} [\min \{G(H) - \mathbb{E}_{H \sim p_\theta} [G(H)], 0\}^2] \quad (5.15)$$

Hence, by the likelihood ratio technique

$$\nabla_\theta \text{SVar}(\theta) = \mathbb{E}_{H \sim p_\theta} [\nabla_\theta \log p_\theta(H) \min \{G(H) - \mathbb{E}_{H \sim p_\theta} [G(H)], 0\}^2] \quad (5.16)$$

From which we can easily derive a Monte Carlo estimate. Since the problems we will consider in the next chapters are not episodic. In the episodic setting, the extension of policy gradient algorithms to the risk-sensitive formulation doesn't present particular difficulties. For a more thorough presentation as well as some more advanced learning algorithms, we refer the interested reader to the extensive work of Tamar et Al. [102], [105], [103], [31].

Algorithm 5.1 Risk-sensitive REINFORCE policy gradient estimate

Input: Policy parameterization θ , number of trajectories M

Output: Risk-sensitive REINFORCE policy gradient estimate

- 1: Sample M trajectories of the MDP following policy π_θ
 - 2: Compute the optimal baseline for the return via Eq. (4.15)
 - 3: Compute the optimal baseline for the squared return via Eq. (5.14)
 - 4: Compute the gradient of the expected return via Eq. (4.13)
 - 5: Compute the gradient of the expected squared return via Eq. (5.12)
 - 6: Compute the risk-sensitive policy gradient via either Eq. (5.6) or (5.8)
-

5.3 Policy Gradient Theorem

In this section the policy gradient theorem is extended to the risk-sensitive framework. In the average reward formulation of the control problem, the theorem and its derivation are analogous to those for the risk-neutral framework. This will allow us to derive in a trivial way the risk-sensitive versions of all the learning algorithms seen in the previous chapter. The risk-sensitive policy gradient theorem for the average-reward formulation was first derived in [83] and the presentation of this section closely follows the original article.

On the other hand, obtaining an equivalent theorem for the discounted reward formulation is more challenging. In the article cited above, the authors prove the theorem under a very strong assumption on the dependence of rewards obtained at different time steps which is not verified in many applications, among which is the asset allocation problem that we will consider for the numerical applications. In the following sections, we will generalize this result by assuming that the reward can depend on both the initial state and the final state of the system. Furthermore, we will discuss the problems arising in the discounted setting and why it is not easy to derive an online policy gradient theorem in this case.

In the original article the authors mostly considered the mean-variance criterion since all the algorithms they propose are easily adapted to the Sharpe ratio criterion. Here we take the opposite direction and present the algorithms for the Sharpe ratio, referring to their article for the mean-variance counterparts.

Let us consider a family of parametrized policies π_θ , with $\theta \in \Theta \subseteq \mathbb{R}^{D_\theta}$. The optimization problem then becomes

$$\max_{\theta} \text{Sh}(\theta) = \frac{\rho(\theta)}{\sqrt{\Lambda(\theta)}} \quad (5.17)$$

where we denoted by $\rho(\theta) = \rho_{\pi_\theta}$ and similarly for the other quantities. Using a policy gradient approach, the policy parameters are updated following the gradient ascent direction.

5.3.1 Average Reward Formulation

In the average reward formulation the gradient of the Sharpe ratio is

$$\nabla_{\theta} \text{Sh}(\theta) = \frac{\eta(\theta) \nabla_{\theta} \rho(\theta) - \frac{1}{2} \rho(\theta) \nabla_{\theta} \eta(\theta)}{\Lambda(\theta) \sqrt{\Lambda(\theta)}} \quad (5.18)$$

Hence, to compute the update direction, it is sufficient to estimate the various quantities appearing in this formula. For instance, the average reward $\rho(\theta)$, the average square reward $\eta(\theta)$ and the reward variance $\Lambda(\theta)$ can be approximated using exponentially weighted moving averages. On the other hand, the gradient of the average reward $\nabla_{\theta} \rho(\theta)$ is given by the standard policy gradient theorem 4.5.1. The only term remaining is the gradient of the average square reward $\nabla_{\theta} \eta(\theta)$, which is provided by the risk-sensitive policy gradient theorem

Theorem 5.3.1 (Risk-Sensitive Policy Gradient). *Let π_{θ} be a differentiable policy. The policy gradient for the average square reward is given by*

$$\nabla_{\theta} \eta(\theta) = \mathbb{E}_{\substack{S \sim d^{\theta} \\ A \sim \pi_{\theta}}} [\nabla_{\theta} \log \pi_{\theta}(S, A) W_{\theta}(S, A)] \quad (5.19)$$

where d^{θ} is the stationary distribution of the Markov chain induced by π_{θ} .

Proof. The proof is analogous to that of the risk-neutral version of theorem. From the basic relation between state-value function and action-value function, we have

$$\begin{aligned} \nabla_{\theta} U_{\theta}(s) &= \nabla_{\theta} \int_{\mathbb{A}} \pi_{\theta}(s, a) W_{\theta}(s, a) da \\ &= \int_{\mathbb{A}} [\nabla_{\theta} \pi_{\theta}(s, a) W_{\theta}(s, a) + \pi_{\theta}(s, a) \nabla_{\theta} W_{\theta}(s, a)] da \end{aligned}$$

Using the Bellman expectation equation for W_{θ}

$$\begin{aligned} \nabla_{\theta} W_{\theta}(s, a) &= \nabla_{\theta} \left[\mathcal{M}(s, a) - \eta_{\theta} + \int_{\mathbb{S}} \mathcal{P}(s, a, s') U_{\theta}(s') ds' \right] \\ &= -\nabla_{\theta} \eta_{\theta} + \int_{\mathbb{S}} \mathcal{P}(s, a, s') \nabla_{\theta} U_{\theta}(s') ds' \end{aligned}$$

Hence, plugging in the first equation

$$\nabla_{\theta} U_{\theta}(s) = \int_{\mathbb{A}} \nabla_{\theta} \pi_{\theta}(s, a) W_{\theta}(s, a) da - \nabla_{\theta} \eta_{\theta} + \int_{\mathbb{A}} \pi_{\theta}(s, a) \int_{\mathbb{S}} \mathcal{P}(s, a, s') \nabla_{\theta} U_{\theta}(s') ds' da$$

Integrating both sides with respect to the stationary distribution d^{θ} and noting that, because of stationarity,

$$\int_{\mathbb{S}} d^{\theta}(s) \int_{\mathbb{A}} \pi(s, a) \int_{\mathbb{S}} \mathcal{P}(s, a, s') \nabla_{\theta} U(s') ds' da ds = \int_{\mathbb{S}} d^{\theta}(s) \nabla_{\theta} U_{\theta}(s) ds$$

we obtain the result

$$\begin{aligned}
\nabla_{\theta} \eta_{\theta} &= \int_{\mathbb{S}} d^{\theta}(s) \int_{\mathbb{A}} \nabla_{\theta} \pi_{\theta}(s, a) W_{\theta}(s, a) da ds \\
&= \int_{\mathbb{S}} d^{\theta}(s) \int_{\mathbb{A}} \pi_{\theta}(s, a) \nabla_{\theta} \log \pi_{\theta}(s, a) W_{\theta}(s, a) da ds \\
&= \mathbb{E}_{\substack{S \sim d^{\theta} \\ A \sim \pi_{\theta}}} [\nabla_{\theta} \log \pi_{\theta}(S, A) W_{\theta}(S, A)]
\end{aligned}$$

□

The result is identical to that for the average reward, provided that we replace the action-value function $Q_{\theta}(s, a)$ by the square action-value function $W_{\theta}(s, a)$. As in the standard risk-neutral case, a state-dependent baseline can be introduced in the gradient without changing the result. In particular, by using the average-adjusted square value function as baseline, we can replace the average adjusted action-value functions with the following square advantage function

$$B_{\theta}(s, a) = W_{\theta}(s, a) - U_{\theta}(s) \quad (5.20)$$

Thus, the gradients can be written as

$$\nabla_{\theta} \eta(\theta) = \mathbb{E}_{\substack{S \sim d_{\pi} \\ A \sim \pi}} [\nabla_{\theta} \log \pi_{\theta}(S, A) B_{\theta}(S, A)] \quad (5.21)$$

From this result, following the exact same reasoning used in the risk-neutral framework, we can design a variety of risk-sensitive actor-critic algorithms, which employ an approximation of the standard and of the square advantage functions to obtain a more accurate estimate of the objective function.

5.3.2 Risk-Sensitive Actor-Critic Algorithm

In [83], starting from Eqs. (4.37) and (5.21), the authors propose a TD(0) risk-sensitive actor-critic algorithm for the average reward setting. The algorithm maintains two critics that estimate the average adjusted value-functions $V_{\theta}(S)$ and $U_{\theta}(S)$ respectively and are updated via a TD(0) temporal difference scheme. Let δ_n^A and δ_n^B be the TD errors for residual value and square value functions

$$\begin{aligned}
\delta_t^A &= R_{t+1} - \hat{\rho}_{t+1} + \hat{V}(S_{t+1}) - \hat{V}(S_t) \\
\delta_t^B &= R_{t+1}^2 - \hat{\eta}_{t+1} + \hat{U}(S_{t+1}) - \hat{U}(S_t)
\end{aligned} \quad (5.22)$$

where \hat{V} , \hat{U} , $\hat{\rho}$ and $\hat{\eta}$ are unbiased estimate of V_{θ} , U_{θ} , $\rho(\theta)$ and $\eta(\theta)$ respectively. It is easy to show that δ_t^A and δ_t^B are unbiased estimates of the advantage functions.

$$\begin{aligned}
\mathbb{E}_{\theta} [\delta_t^A | S_t = s, A_t = a] &= A_{\theta}(s, a) \\
\mathbb{E}_{\theta} [\delta_t^B | S_t = s, A_t = a] &= B_{\theta}(s, a)
\end{aligned}$$

An estimate of the gradients can be obtained by replacing the advantage functions with the TD errors

$$\begin{aligned}\nabla_{\theta}\rho(\theta) &\approx \nabla_{\theta}\log \pi_{\theta}(S_t, A_t)\delta_t^A \\ \nabla_{\theta}\eta(\theta) &\approx \nabla_{\theta}\log \pi_{\theta}(S_t, A_t)\delta_t^B\end{aligned}\tag{5.23}$$

The value functions are linearly approximated using some feature vectors $\Phi_V : \mathbb{S} \rightarrow \mathbb{R}^{D_V}$ and $\Phi_U : \mathbb{S} \rightarrow \mathbb{R}^{D_U}$ as follows

$$\begin{aligned}\widehat{V}(s) &= \psi_V^T \Phi_V(s) \\ \widehat{U}(s) &= \psi_U^T \Phi_U(s)\end{aligned}\tag{5.24}$$

Combining all these ingredients lead to the Risk-Sensitive Average Reward Actor-Critic algorithm (RSARAC), the pseudocode of which is reported in Algorithm 5.2. Let us notice that the algorithm is a three time-scale stochastic approximation algorithm, where the learning rates, in addition to the usual Robbins-Monro conditions, should satisfy

$$\alpha_k < \beta_k < \gamma_k$$

5.3.3 Discounted Reward Formulation

In the discounted reward formulation, we need to adapt the definition of the Sharpe ratio associated to a policy. Let us suppose that the system always start in the same initial state s_0 , then we can introduce the start-state Sharpe ratio that can be achieved following policy π_{θ} as

$$\text{Sh}(\theta) = \frac{V_{\theta}(s_0)}{\sqrt{\Lambda_{\theta}(s_0)}}\tag{5.25}$$

where V_{θ} and Λ_{θ} are the state-value function and the variance-function introduced in Chapter 2. Hence, the gradient of the Sharpe ratio is given by

$$\nabla_{\theta}\text{Sh}(\theta) = \frac{U_{\theta}(s_0)\nabla_{\theta}V_{\theta}(s_0) - \frac{1}{2}V_{\theta}(s_0)\nabla_{\theta}U_{\theta}(s_0)}{\Lambda_{\theta}(s_0)\sqrt{\Lambda_{\theta}(s_0)}}\tag{5.26}$$

The gradient of the state-value function $\nabla_{\theta}V_{\theta}(s_0)$ is given by the risk-neutral policy gradient theorem. Therefore, in order to approximate the gradient of the Sharpe ratio, we need to estimate the value functions $V_{\theta}(s_0)$, $U_{\theta}(s_0)$, the variance $\Lambda_{\theta}(s_0)$, and the gradient of the square state-value function $U_{\theta}(s_0)$. The first three terms might be easily approximated using moving averages. On the other hand, estimating $\nabla_{\theta}U_{\theta}(s_0)$ in the same spirit of the policy gradient theorem is much more delicate.

Theorem 5.3.2 (Risk-Sensitive Policy Gradient). *Let π_{θ} be a differentiable policy. The policy gradient for the square state-value function is given by*

$$\begin{aligned}\nabla_{\theta}U_{\theta}(s_0) &= \mathbb{E}_{\substack{S \sim d_{\gamma^2}^{\theta}(s_0, \cdot) \\ A \sim \pi_{\theta}(S, \cdot) \\ S' \sim \mathcal{P}(S, A, \cdot)}} [\nabla_{\theta}\log \pi_{\theta}(S, A)W_{\theta}(S, A) \\ &\quad + 2\gamma\mathcal{R}(S, A)\nabla_{\theta}V_{\theta}(S') + 2\gamma C_{\theta}(S, A)]\end{aligned}\tag{5.27}$$

Algorithm 5.2 Risk-Sensitive Average Reward Actor-Critic algorithm

Input:

- Initial actor parameters θ^0
- Initial critics parameters ψ_V^0 and ψ_U^0
- Actor learning rate $\{\alpha_k\}$
- Critics learning rate $\{\beta_k\}$
- Averages learning rate $\{\gamma_k\}$

Output: Approximation of the optimal policy $\pi_{\theta^*} \approx \pi_*$ 1: **repeat**2: Observe tuple $\langle s_k, a_k, r_{k+1}, s_{k+1} \rangle$ sampled from the MDP.

3: Update averages

$$\begin{aligned}\hat{\rho}_{k+1} &= (1 - \gamma_k)\hat{\rho}_k + \gamma_k r_{k+1} \\ \hat{\eta}_{k+1} &= (1 - \gamma_k)\hat{\eta}_k + \gamma_k r_{k+1}^2\end{aligned}$$

4: Compute TD errors

$$\begin{aligned}\delta_k^A &= r_{k+1} - \hat{\rho}_{k+1} + (\psi_V^k)^T \Phi_V(s_{k+1}) - (\psi_V^k)^T \Phi_V(s_k) \\ \delta_k^B &= r_{k+1}^2 - \hat{\eta}_{k+1} + (\psi_U^k)^T \Phi_U(s_{k+1}) - (\psi_U^k)^T \Phi_U(s_k)\end{aligned}$$

5: Update critic parameters

$$\begin{aligned}\psi_V^{k+1} &= \psi_V^k + \beta_k \delta_k^A \Phi_V(s_k) \\ \psi_U^{k+1} &= \psi_U^k + \beta_k \delta_k^B \Phi_U(s_k)\end{aligned}$$

6: Update actor parameters $\theta^{k+1} = \theta^k + \alpha_k \widehat{\text{Sh}}(\theta_k)$.7: $k \leftarrow k + 1$ 8: **until** converged

where $d_{\gamma^2}^\theta(s_0, \cdot)$ is the γ^2 -discounted visiting distribution over states starting from the initial state s_0 and following policy π_θ

$$d_{\gamma^2}^\theta(s_0, x) = \sum_{k=0}^{\infty} \gamma^{2k} \mathcal{P}_\theta^{(k)}(s_0, x) \quad (5.28)$$

Proof. From the basic relation between square state-value function and the square action-value function, we have

$$\begin{aligned}\nabla_\theta U_\theta(s) &= \nabla_\theta \int_{\mathbb{A}} \pi_\theta(s, a) W_\theta(s, a) da \\ &= \int_{\mathbb{A}} [\nabla_\theta \pi_\theta(s, a) W_\theta(s, a) + \pi_\theta(s, a) \nabla_\theta W_\theta(s, a)] da\end{aligned}$$

Hence, using the Bellman expectation equation for W_θ

$$\begin{aligned}\nabla_\theta W_\theta(s, a) &= \nabla_\theta [\mathcal{M}(s, a) + 2\gamma\mathcal{R}(s, a)T_a V_\pi(s) + 2\gamma C_\theta(s, a) + \gamma^2 T_a U_\theta(s)] \\ &= 2\gamma\mathcal{R}(s, a)\nabla_\theta T_a V_\theta(s) + 2\gamma\nabla_\theta C_\theta(s, a) + \gamma^2\nabla_\theta T_a U_\theta(s) \\ &= \int_{\mathbb{S}} \mathcal{P}(s, a, s') [2\gamma\mathcal{R}(s, a)\nabla_\theta V_\theta(s') + 2\gamma\nabla_\theta C_\theta(s, a) + \nabla_\theta U_\theta(s')] ds'\end{aligned}$$

where we assumed to be able to exchange the gradient and the integral. Plugging in the first equation and exploiting the fact that $\int_{\mathbb{S}} \mathcal{P}(s, a, s') ds' = 1$, we obtain

$$\begin{aligned}\nabla_\theta U_\theta(s) &= \int_{\mathbb{A}} \pi_\theta(s, a) \int_{\mathbb{S}} \mathcal{P}(s, a, s') [\nabla_\theta \log \pi_\theta(s, a) W_\theta(s, a) \\ &\quad + 2\gamma\mathcal{R}(s, a)\nabla_\theta V_\theta(s') + 2\gamma\nabla_\theta C_\theta(s, a) + \nabla_\theta U_\theta(s')] ds' da\end{aligned}$$

Unrolling $\nabla_\theta U_\theta$ infinite times and denoting by $\mathcal{P}_\theta^{(k)}(s, x)$ the probability of going from state s to state x in k steps under policy π_θ , we obtain

$$\begin{aligned}\nabla_\theta U_\theta(s) &= \sum_{k=0}^{\infty} \gamma^{2k} \int_{\mathbb{S}} \mathcal{P}_\theta^{(k)}(s, x) \int_{\mathbb{A}} \pi_\theta(x, a) \int_{\mathbb{S}} \mathcal{P}(x, a, x') [\nabla_\theta \log \pi_\theta(x, a) W_\theta(x, a) \\ &\quad + 2\gamma\mathcal{R}(x, a)\nabla_\theta V_\theta(x') + 2\gamma\nabla_\theta C_\theta(x, a)] dx' dadx\end{aligned}$$

Defining the γ^2 -discounted visiting distribution of state x starting from state s as

$$d_{\gamma^2}^\theta(s, x) = \sum_{k=0}^{\infty} \gamma^{2k} \mathcal{P}_\theta^{(k)}(s, x)$$

we have the result

$$\begin{aligned}\nabla_\theta U_\theta(s) &= \int_{\mathbb{S}} d_{\gamma^2}^\theta(s, x) \int_{\mathbb{A}} \pi_\theta(x, a) \int_{\mathbb{S}} \mathcal{P}(x, a, x') [\nabla_\theta \log \pi_\theta(x, a) W_\theta(x, a) \\ &\quad + 2\gamma\mathcal{R}(x, a)\nabla_\theta V_\theta(x') + 2\gamma\nabla_\theta C_\theta(x, a)] dx' dadx\end{aligned}$$

□

Compared to risk-sensitive policy gradient theorem in the average reward formulation, we have two additional terms: one term depending on the gradient of the state-value function at the next state and another term depending on the covariance between the one-step reward and the successive return. These terms are very difficult to approximate online in a continuing environment. Therefore, the result is of practical interest only for an episodic environment where the experiments have a finite (possibly random) lifespan. Since the application we considered does not fall in this category, we will not discuss any risk-sensitive algorithm for the discounted formulation.

Chapter 6

Parameter-Based Policy Gradient

The learning algorithms derived from the policy gradient theorems look for a set of optimal controller parameters by perturbing the control signal. Indeed, the policy gradient theorem holds only for stochastic policies, which explore in the action space. However, a significant problem with this sampling strategy is that the high variance in their gradient estimates leads to slow convergence. On the other hand, parameter-based algorithms such as PGPE look for an optimal solution to the control problem by directly perturbing the controller parameters. These algorithms thus employ a deterministic controller and explore in the parameters space. However, in its original version, PGPE was conceived for episodic environments and extended to continuing environment by artificially truncating the experiments lifespan. In this section we propose a parameter-based version of the policy gradient theorem, which can be used to derive online learning algorithms both in the risk-neutral and the risk-sensitive frameworks. The following discussion starts from and generalizes the results presented in [71], where the authors propose an online natural PGPE algorithm for risk-neutral problems. Given the difficulties to extend the policy gradient theorem to the risk-sensitive discounted formulation, here we consider only the average reward formulation.

6.1 Risk-Neutral Framework

Let us consider a deterministic controller $F : \mathbb{S} \times \Theta \rightarrow \mathbb{A}$ that, given a set of parameters $\theta \in \Theta \subseteq \mathbb{R}^{D_\theta}$, maps a state $s \in \mathbb{S}$ to an action $a = F(s; \theta) = F_\theta(s) \in \mathbb{A}$. The policy parameters are drawn from a probability distribution p_ξ , with hyper-parameters $\xi \in \Xi \subseteq \mathbb{R}^{D_\xi}$. Combining these two hypotheses, the agent follows a stochastic policy π_ξ defined by

$$\forall B \in \mathcal{A}, \pi_\xi(s, B) = \pi(s, B; \xi) = \int_{\Theta} p_\xi(\theta) \mathbb{1}_{F_\theta(s) \in B} d\theta \quad (6.1)$$

In the risk-neutral setting, the goal of the agent is to find a set of hyperparameters ξ_* that maximizes the average reward obtained over a single time-step

$$\xi_* = \arg \max_{\xi} \rho(\xi)$$

Following the policy gradient approach, we iteratively update the parameters in the gradient ascent direction

$$\xi_{k+1} = \xi_k + \alpha_k \nabla_{\xi} \rho(\xi_k)$$

Theorem 6.1.1 (Risk-Neutral Parameter-Based Policy Gradient). *Let p_{ξ} be differentiable with respect to ξ , then the gradient of the average reward is given by*

$$\nabla_{\xi} \rho(\xi) = \mathbb{E}_{\substack{S \sim d^{\xi} \\ \theta \sim p_{\xi}}} [\nabla_{\xi} \log p_{\xi}(\theta) Q_{\pi_{\xi}}(S, \theta)] \quad (6.2)$$

where we denoted $Q_{\xi}(S, \theta) = Q_{\xi}(S, F_{\theta}(S))$.

Proof. Thanks to the likelihood ratio technique, we have

$$\begin{aligned} \nabla_{\xi} \pi_{\xi}(s, a) &= \int_{\Theta} \nabla_{\xi} p_{\xi}(\theta) \mathbb{1}_{F_{\theta}(s)=a} d\theta \\ &= \int_{\Theta} p_{\xi}(\theta) \nabla_{\xi} \log p_{\xi}(\theta) \mathbb{1}_{F_{\theta}(s)=a} d\theta \\ &= \mathbb{E}_{\theta \sim p_{\xi}} [\nabla_{\xi} \log p_{\xi}(\theta) \mathbb{1}_{F_{\theta}(s)=a}] \end{aligned}$$

hence, the policy π_{ξ} is differentiable with respect to ξ and the standard policy gradient theorem holds

$$\begin{aligned} \nabla_{\xi} \rho(\xi) &= \mathbb{E}_{\substack{S \sim d^{\xi} \\ A \sim \pi_{\xi}}} [\nabla_{\theta} \log \pi_{\theta}(S, A) Q_{\theta}(S, A)] \\ &= \int_{\mathbb{S}} d^{\xi}(s) \int_{\mathbb{A}} \pi_{\xi}(s, a) \nabla_{\xi} \log \pi_{\xi}(s, a) Q_{\pi_{\xi}}(s, a) da ds \\ &= \int_{\mathbb{S}} d^{\xi}(s) \int_{\mathbb{A}} \nabla_{\xi} \pi_{\xi}(s, a) Q_{\pi_{\xi}}(s, a) da ds \end{aligned}$$

Plugging the first equation in the inner integral and exchanging the integral over the action space with the expectation yields

$$\begin{aligned} \int_{\mathbb{A}} \nabla_{\xi} \pi_{\xi}(s, a) Q_{\pi_{\xi}}(s, a) da &= \mathbb{E}_{\theta \sim p_{\xi}} \left[\nabla_{\xi} \log p_{\xi}(\theta) \int_{\mathbb{A}} \mathbb{1}_{F_{\theta}(s)=a} Q_{\pi_{\xi}}(s, a) da \right] \\ &= \mathbb{E}_{\theta \sim p_{\xi}} [\nabla_{\xi} \log p_{\xi}(\theta) Q_{\pi_{\xi}}(s, F_{\theta}(s))] \\ &= \mathbb{E}_{\theta \sim p_{\xi}} [\nabla_{\xi} \log p_{\xi}(\theta) Q_{\pi_{\xi}}(s, \theta)] \end{aligned}$$

Finally, plugging this equation in the policy gradient theorem, we obtain the result

$$\nabla_{\xi} \rho(\xi) = \mathbb{E}_{\substack{S \sim d^{\xi} \\ \theta \sim p_{\xi}}} [\nabla_{\xi} \log p_{\xi}(\theta) Q_{\pi_{\xi}}(S, \theta)]$$

□

This expression is very similar to the original policy gradient theorem, but the expectation is taken over the controller parameters instead of the action space and we have the likelihood score the controller parameters distribution instead of that of the stochastic policy. Thus, we might interpret this result as if the agent directly selected the parameters θ according to a policy p_ξ , which then lead to an action through the deterministic mapping F_θ . Therefore, it is as if the agent's policy was in the parameters space and not in the control space. As in the standard policy gradient methods, we can add a state-dependent baseline $B_\xi(S)$ to the gradient without increasing the bias

$$\nabla_\xi \rho(\xi) = \mathbb{E} [\nabla_\xi \log p_\xi(\theta) (Q_{\pi_\xi}(S, \theta) - B_\xi(S))] \quad (6.3)$$

Indeed,

$$\begin{aligned} \mathbb{E} [\nabla_\xi \log p_\xi(\theta) B(S)] &= \int_{\mathbb{S}} d^\xi(s) \int_{\Theta} p_\xi(\theta) \nabla_\xi \log p_\xi(\theta) B_\xi(s) d\theta ds \\ &= \int_{\mathbb{S}} d^\xi(s) B_\xi(s) ds \underbrace{\int_{\Theta} \nabla_\xi p_\xi(\theta) d\theta}_{\nabla_\xi 1=0} = 0 \end{aligned}$$

This result can be used to design several actor-only or actor-critic algorithms that are the parameter-based equivalents of the traditional control-based policy-gradient algorithms discussed in the previous sections. Algorithm 6.1 reports the pseudocode for a parameter-based actor-critic algorithm, which is one of the learning methods that will be used in the numerical applications.

6.1.1 Parameter-Based Natural Policy Gradient

Given the interpretation of p_ξ as a parameter-based stochastic policy, we can easily generalize the natural policy gradient methods by introducing the parameter-based Fischer information matrix

$$F_\xi = \mathbb{E}_{\theta \sim p_\xi} [\nabla_\xi \log p_\xi(\theta) \nabla_\theta \log p_\xi(\theta)^T] \quad (6.4)$$

Natural parameter-based policy gradient methods update the hyper-parameters following the natural gradient direction

$$\tilde{\nabla}_\xi \rho(\xi) = F_\xi^{-1} \nabla_\xi \rho(\xi) \quad (6.5)$$

The advantage of this approach is that, for some particular probability distributions, the inverse of the Fisher information matrix can be computed analytically. This is for instance the case if the controller parameters are normally distributed. This choice leads to an efficient online algorithm called [Natural Policy Gradient with Parameter Based Exploration \(NPGPE\)](#) [71]. More in detail, suppose that

$$\theta \sim \mathcal{N}(\mu, \Gamma \Gamma^T)$$

Algorithm 6.1 Actor-Critic PGPE**Input:**

- Deterministic controller $F : \mathbb{S} \times \Theta \rightarrow \mathbb{A}$
- Initial hyper-parameters ξ_0
- Initial critic parameters ψ_0
- Learning rate $\{\alpha_k\}$
- Momentum parameter λ

Output: Approximation of the optimal hyper-parameters ξ_*

- 1: Initialize $k = 0$, average reward $\hat{\rho}_0 = 0$ and eligibility trace $e_{-1} = 0$
- 2: **repeat**
- 3: Observe current state s_k
- 4: Simulate controller parameters $\theta_k \sim p_{\xi_k}$
- 5: Perform action $a_k = F_{\theta_k}(s_k)$ and receive reward r_{k+1}
- 6: Update average reward estimate $\hat{\rho}_{k+1} = \hat{\rho}_k + \alpha_k(r_{k+1} - \hat{\rho}_k)$
- 7: Compute TD error $\hat{\delta}_k = r_{k+1} - \hat{\rho}_{k+1} + \hat{V}_{\psi_k}(s_{k+1}) - \hat{V}_{\psi_k}(s_k)$
- 8: Update critic parameters $\psi_{k+1} = \psi_k + \alpha_k \hat{\delta}_k \nabla_{\psi} \hat{V}_{\psi_k}(s_k)$
- 9: Update eligibility trace $e_k = \lambda e_{k-1} + \nabla_{\xi} \log p_{\xi}(\theta_k)$
- 10: Update hyper-parameters $\xi_{k+1} = \xi_k + \alpha_k \hat{\delta}_{\psi_k}^k e_k$
- 11: $k \leftarrow k + 1$
- 12: **until** converged

where $\mu \in \mathbb{R}^n$ is the controller parameters mean and $\Gamma \in \mathbb{R}^{n \times n}$ denotes the Cholesky factor of the distribution covariance matrix Σ . The hyper-parameters are thus given by

$$\xi = (\mu, \Gamma_1, \Gamma_2, \dots, \Gamma_n) \in \mathbb{R}^{\frac{n^2+3n}{2}}$$

where we denoted $\Gamma_k = \Gamma_{k,k:n}^T$ the k -th column of the Cholesky factor. For this distribution, the parameter-based policy gradient is given by

$$\nabla_{\mu} \log p_{\xi}(\theta) = \Sigma^{-1}(\theta - \mu) \quad (6.6)$$

$$\nabla_{\Gamma_k} \log p_{\xi}(\theta) = \begin{bmatrix} 0 & I_{\bar{k}} \end{bmatrix} (\Gamma^{-1} Y - \text{diag}(\Gamma^{-1})) e_k \quad (6.7)$$

where $Y = \Gamma^{-T}(\theta - \mu)(\theta - \mu)^T \Gamma^{-1}$. These formulas might be used in a standard PGPE algorithm. However, the inversion of Γ is computationally expensive except for some very simple cases, for instance when Σ is a diagonal matrix. In [98], the authors proved that the Fisher information matrix for this distribution is a block diagonal matrix $F_{\xi} = \text{diag}\{B_0, B_1, \dots, B_n\}$ where

$$\begin{cases} B_0 &= \Sigma^{-1} \\ B_k &= \begin{bmatrix} 0 & I_{\bar{k}} \end{bmatrix} \Gamma^{-1} (e_k e_k^T + I) \Gamma^{-T} \begin{bmatrix} 0 \\ I_{\bar{k}} \end{bmatrix} \end{cases}$$

where e_k is the k -th element of the canonical basis of \mathbb{R}^n and $I_{\bar{k}}$ is $n - k + 1$ -dimensional identity matrix. In [4], the authors found the following analytical

expression for the inverse of each block

$$\begin{cases} B_0^{-1} &= \Sigma \\ B_k^{-1} &= \begin{bmatrix} 0 & I_{\bar{k}} \end{bmatrix} \Gamma^T \left(\begin{bmatrix} 0 & 0 \\ 0 & I_{\bar{k}} \end{bmatrix} - \frac{1}{2} e_k e_k^T \right) \Gamma \begin{bmatrix} 0 \\ I_{\bar{k}} \end{bmatrix} \end{cases}$$

Multiplying the policy gradients by the inverse Fisher informatio matrix $F_\xi^{-1} = \text{diag}\{B_0^{-1}, B_1^{-1}, \dots, B_n^{-1}\}$, we obtain the following natural policy gradients

$$\tilde{\nabla}_\mu \log p_\xi(\theta) = \theta - \mu \quad (6.8)$$

$$\tilde{\nabla}_\Gamma \log p_\xi(\theta) = \left(\text{triu}(Y) - \frac{1}{2} \text{diag}(Y) - \frac{1}{2} I \right) \Gamma \quad (6.9)$$

The natural policy gradients can be computed in $O(n^3)$, which is a sensible improvement compared to the $O(n^6)$ complexity for the standard PGPE updates. Let us remark that for an independent multi-variate Gaussian distribution, i.e. a diagonal covariance matrix Σ , the updates can be performed in $O(n)$. However, this approach neglects the dependences among parameters and could lead to suboptimal performances. On the other hand. updating the full Cholesky factor introduces many more parameters to be learned which may slow down the convergence of the method. The resulting NPGPE algorithm is reported in Algorithm 6.2. Given the parameter-based policy gradient algorithm, we can easily combine the actor-critic approach and the natural gradient technique in a parameter-based natural actor-critic algorithm, which is outlined in Algorithm 6.3.

Algorithm 6.2 NPGPE

Input:

- Deterministic controller $F : \mathbb{S} \times \Theta \rightarrow \mathbb{A}$
- Initial hyper-parameters ξ^0
- Learning rate $\{\alpha_k\}$
- Momentum parameter λ

Output: Approximation of the optimal hyper-parameters ξ_*

- 1: Initialize $k = 0$, average reward $\hat{\rho}_0 = 0$ and eligibility trace $e_{-1} = 0$
- 2: **repeat**
- 3: Observe current state s_k
- 4: Draw $\zeta_k \sim \mathcal{N}(0, I_n)$
- 5: Compute controller parameters $\theta_k = \mu_k + \Gamma^T \zeta_k$
- 6: Perform action $a_k = F_{\theta_k}(s_k)$ and receive reward r_{k+1}
- 7: Update average reward estimate $\hat{\rho}_{k+1} = \hat{\rho}_k + \alpha_k(r_{k+1} - \hat{\rho}_k)$
- 8: Compute natural policy gradients

$$\tilde{\nabla}_\mu \log p_{\xi_k}(\theta_k) = \theta_k - \mu_k$$

$$\tilde{\nabla}_\Gamma \log p_{\xi_k}(\theta_k) = \left(\text{triu}(\zeta_k \zeta_k^T) - \frac{1}{2} \text{diag}(\zeta_k \zeta_k^T) - \frac{1}{2} I \right) \Gamma$$

- 9: Update eligibility trace $e_k = \lambda e_{k-1} + \nabla_\xi \log p_{\xi_k}(\theta_k)$
 - 10: Update hyper-parameters $\xi_{k+1} = \xi_k + \alpha_k(r_{k+1} - \hat{\rho}_k)e_k$
 - 11: $k \leftarrow k + 1$
 - 12: **until** converged
-

Algorithm 6.3 Natural Actor-Critic PGPE

Input:

- Deterministic controller $F : \mathbb{S} \times \Theta \rightarrow \mathbb{A}$
- Initial hyper-parameters ξ^0
- Initial critic parameters ψ_0
- Actor learning rate $\{\alpha_k\}$
- Critics learning rate $\{\beta_k\}$
- Momentum parameter λ

Output: Approximation of the optimal hyper-parameters ξ_*

- 1: Initialize $k = 0$, average reward $\hat{\rho}_0 = 0$ and eligibility trace $e_{-1} = 0$
- 2: **repeat**
- 3: Observe current state s_k
- 4: Draw $\zeta_k \sim \mathcal{N}(0, I_n)$
- 5: Compute controller parameters $\theta_k = \mu_k + \Gamma^T \zeta_k$
- 6: Perform action $a_k = F_{\theta_k}(s_k)$ and receive reward r_{k+1}
- 7: Update average reward estimate $\hat{\rho}_{k+1} = \hat{\rho}_k + \alpha_k(r_{k+1} - \hat{\rho}_k)$
- 8: Compute TD error $\hat{\delta}_k = r_{k+1} - \hat{\rho}_{k+1} + \hat{V}_{\psi_k}(s_{k+1}) - \hat{V}_{\psi_k}(s_k)$
- 9: Update critic parameters $\psi_{k+1} = \psi_k + \alpha_k \hat{\delta}_k \nabla_{\psi} \hat{V}_{\psi_k}(s_k)$.
- 10: Compute natural policy gradients

$$\tilde{\nabla}_{\mu} \log p_{\xi_k}(\theta_k) = \theta_k - \mu_k$$

$$\tilde{\nabla}_{\Gamma} \log p_{\xi_k}(\theta_k) = \left(\text{triu}(\zeta_k \zeta_k^T) - \frac{1}{2} \text{diag}(\zeta_k \zeta_k^T) - \frac{1}{2} I \right) \Gamma$$

- 11: Update eligibility trace $e_k = \lambda e_{k-1} + \tilde{\nabla}_{\xi} \log p_{\xi_k}(\theta_k)$
 - 12: Update hyper-parameters $\xi_{k+1} = \xi_k + \alpha_k \hat{\delta}_k e_k$.
 - 13: $k \leftarrow k + 1$
 - 14: **until** converged
-

6.2 Risk-Sensitive Framework

In the risk-sensitive setting, the agent tries to find a policy that maximizes the Sharpe ratio. Following the same reasoning of the control-based approach, we simply need to estimate the gradient of the average square reward $\eta(\xi)$. The parameter-based policy gradient generalizes without any additional effort

Theorem 6.2.1 (Risk-Sensitive Parameter-Based Policy Gradient). *Let p_ξ be differentiable with respect to ξ , then the gradient of the average square reward is given by*

$$\nabla_\xi \eta(\xi) = \mathbb{E}_{\substack{S \sim d^\xi \\ \theta \sim p_\xi}} [\nabla_\xi \log p_\xi(\theta) W_{\pi_\xi}(S, \theta)] \quad (6.10)$$

where we denoted $W_\xi(S, \theta) = W_\xi(S, F_\theta(S))$.

Again, we can add a state-dependent baseline $B_\xi(S)$ to the gradient without increasing the bias

$$\nabla_\xi \eta(\xi) = \mathbb{E} [\nabla_\xi \log p_\xi(\theta) (W_{\pi_\xi}(S, \theta) - B_\xi(S))] \quad (6.11)$$

6.2.1 Parameter-Based Natural Policy Gradient

The natural policy gradient idea can be trivially extended to the gradient of the average square reward. This leads to the following natural gradient estimate

$$\begin{aligned} \tilde{\nabla}_\xi \eta(\xi) &= F_\xi^{-1} \nabla_\xi \eta(\xi) \\ &= F_\xi^{-1} \mathbb{E}_{\substack{S \sim d^\xi \\ \theta \sim p_\xi}} [\nabla_\xi \log p_\xi(\theta) (W_{\pi_\xi}(S, \theta) - B_\xi(S))] \\ &= \mathbb{E}_{\substack{S \sim d^\xi \\ \theta \sim p_\xi}} [\tilde{\nabla}_\xi \log p_\xi(\theta) (W_{\pi_\xi}(S, \theta) - B_\xi(S))] \end{aligned} \quad (6.12)$$

which can be easily estimated via Monte Carlo. Combining this natural gradient with the one for the average reward, we obtain the natural policy gradient for the Sharpe ratio

$$\tilde{\nabla}_\xi \text{Sh}(\xi) = F_\xi^{-1} \nabla_\xi \text{Sh}(\xi) = \frac{\eta(\xi) \tilde{\nabla}_\xi \rho(\xi) - \frac{1}{2} \rho(\xi) \tilde{\nabla}_\xi \eta(\xi)}{\Lambda(\xi) \sqrt{\Lambda(\xi)}} \quad (6.13)$$

Hence, it is trivial to derive a risk-sensitive version of the NPGPE learning algorithm. The pseudocode is reported in Algorithm 6.4. Similarly to the risk-neutral case, it is simple to introduce a critic both for the average reward and for the average square reward.

Algorithm 6.4 Risk-Sensitive NPGPE

Input:

- Deterministic controller $F : \mathbb{S} \times \Theta \rightarrow \mathbb{A}$
- Initial hyper-parameters ξ^0
- Learning rate $\{\alpha_k\}$
- Momentum parameter λ

Output: Approximation of the optimal hyper-parameters ξ_*

- 1: Initialize $k = 0$, average reward $\hat{\rho}_0 = 0$, average square reward $\hat{\eta}_0 = 0$ and eligibility trace $e_{-1} = 0$
- 2: **repeat**
- 3: Observe current state s_k
- 4: Draw $\zeta_k \sim \mathcal{N}(0, I_n)$
- 5: Compute controller parameters $\theta_k = \mu_k + \Gamma^T \zeta_k$
- 6: Perform action $a_k = F_{\theta_k}(s_k)$ and receive reward r_{k+1}
- 7: Update average reward estimate $\hat{\rho}_{k+1} = \hat{\rho}_k + \alpha_k(r_{k+1} - \hat{\rho}_k)$
- 8: Update average square reward estimate $\hat{\eta}_{k+1} = \hat{\eta}_k + \alpha_k(r_{k+1}^2 - \hat{\eta}_k)$
- 9: Compute natural policy gradients

$$\tilde{\nabla}_{\mu} \log p_{\xi_k}(\theta_k) = \theta_k - \mu_k$$

$$\tilde{\nabla}_{\Gamma} \log p_{\xi_k}(\theta_k) = \left(\text{triu}(\zeta_k \zeta_k^T) - \frac{1}{2} \text{diag}(\zeta_k \zeta_k^T) - \frac{1}{2} I \right) \Gamma$$

- 10: Update eligibility trace $e_k = \lambda e_{k-1} + \nabla_{\xi} \log p_{\xi_k}(\theta_k)$
- 11: Compute natural gradient $\tilde{\nabla}_{\xi} \rho(\xi_k) = e_k (r_{k+1} - \hat{\rho}_k)$
- 12: Compute natural gradient $\tilde{\nabla}_{\xi} \eta(\xi_k) = e_k (r_{k+1}^2 - \hat{\eta}_k)$
- 13: Compute Sharpe ratio natural gradient

$$\tilde{\nabla}_{\xi} \text{Sh}(\xi_k) = \frac{\hat{\eta}_k \tilde{\nabla}_{\xi} \rho(\xi_k) - \frac{1}{2} \hat{\rho}_k \tilde{\nabla}_{\xi} \eta(\xi_k)}{\hat{\Lambda}_k \sqrt{\hat{\Lambda}_k}}$$

- 14: Update hyper-parameters $\xi_{k+1} = \xi_k + \alpha_k \tilde{\nabla}_{\xi} \text{Sh}(\xi_k)$
 - 15: $k \leftarrow k + 1$
 - 16: **until** converged
-

Chapter 7

Financial Applications of Reinforcement Learning

Both in the academia and in the financial industry there has always been a strong interest in developing automated systems able to take financial decisions in the place of humans. The advent of computers made it possible to analyze the huge amount of data available in the markets and to discover patterns that can be exploited automatically by a program, in what is known as a statistical arbitrage. However, finding these signals and transform them in a profit is not straightforward. In Section 7.1 we briefly discuss the *Efficient Market Hypothesis* (EMH) which states that it is impossible to "beat the market" because stock market efficiency causes existing share prices to always incorporate and reflect all relevant information. However, the success of many quantitative hedge-funds over the years provides strong evidence that real markets are not so efficient as they are often considered in the literature. We thus discuss some weaker versions of market efficiency which should always be kept in mind when trying to develop automated trading systems able to outperform the market. In this chapter, out of the many techniques that can be tested to achieve this goal, we will focus on reinforcement learning. In Section 7.2, we describe the relevant works that can be found in the literature. This presentation will give an overview of the spectrum of financial problems that can be tackled with RL techniques. In Section 7.3, we present in detail the problem of asset allocation with transaction costs, which will be used for the numerical applications.

7.1 Efficient Market Hypothesis

When trying to forecast future returns of speculative assets, one should always keep in mind the [Efficient Market Hypothesis \(EMH\)](#) [107]. In its most basic form, the EMH states that asset returns cannot be predicted, otherwise many investors would exploit them to generate unlimited profits. Such a "money-machine" producing unlimited wealth is impossible in a stable economy. This hypothesis seems intuitively sensible because as soon as an inconsistency appears in the market, there will be an agent who exploits it and makes the opportunity

disappear. Hence, forecasting models “self-destructs” in an efficient market which auto-corrects the emerging anomalies. That might appear as the end of the forecasters’ ambitions. However this idea is not completely convincing and papers continue to appear attempting to forecast stock returns, usually with very little success. On the other hand, a number of successful stories of quantitative hedge-funds which consistently made profits looking for statistical arbitrages make us doubt the EMH. In the next sections we give more formal definitions of the EMH in its various forms and discuss some of the critics.

7.1.1 Formal Definitions of the EMH

Definition 7.1.1 (Efficient Market Hypothesis). *A market is efficient with respect to information set \mathcal{F}_t if it is impossible to make economic profits by trading on the basis of information set \mathcal{F}_t .*

In the previous definition, economic profits stand for risk adjusted returns net of all costs. In other words, financial markets are efficient if they do not allow investors to earn above-average risk-adjusted returns. The EMH is in essence an extension of the zero-profit competitive equilibrium condition from the certainty world of classical price theory to the dynamic behavior of prices in speculative markets under conditions of uncertainty [48]. The application of the zero profit condition to speculative markets under the assumption of zero storage costs and zero transactions costs gives us the result that asset prices (after adjustment for required returns) will behave as a martingales with respect to the information set $\{\mathcal{F}_t\}$.

Several versions of the EMH have been proposed in the literature depending on the information set considered.

Definition 7.1.2 (Weak Form of the EMH). *\mathcal{F}_t represents only the information contained in the past price history of the market as of time t .*

Definition 7.1.3 (Semi-Strong Form of the EMH). *\mathcal{F}_t represents all information publicly available at time t .*

Definition 7.1.4 (Strong Form of the EMH). *\mathcal{F}_t represents all information (public and private) known to anyone at time t .*

Let us notice that it is not usually asserted that a market is efficient with respect to inside information since this information is not widely accessible and hence cannot be expected to be fully incorporated in the current price. The empirical evidence presented in [68] is largely supportive of weak form and semi-strong form efficiency, while [39] reports stronger evidence of predictability in returns based both on lagged values of returns and publicly available information.

7.1.2 Critics to the EMH

The intellectual dominance of the efficient-market revolution has been strongly challenged by economists who stress psychological and behavioral elements of stock-price determination and by econometricians who argue that stock returns are, to a considerable extent, predictable. Still, in [67], the author expresses the following pro-EMH view

What I do not argue is that the market pricing is always perfect. After the fact, we know that markets have made egregious mistakes as I think occurred during the recent Internet bubble. Nor do I deny that psychological factors influence securities prices. But I am convinced that Benjamin Graham was correct in suggesting that while the stock market in the short run may be a voting mechanism, in the long run it is a weighing mechanism. True value will win out in the end. And before the fact, there is no way in which investors can reliably exploit any anomalies or patterns that might exist. I am skeptical that any of the “predictable patterns” that have been documented in the literature were ever sufficiently robust so as to have created profitable investment opportunities and after they have been discovered and publicized, they will certainly not allow investors to earn excess returns.

This seems to be contradicted by the numerous examples of successful quantitative hedge funds that would suggest that it is possible, even if extremely difficult, to identify consistent and profitable patterns in market dynamics. Medallion, the flagship fund of Renaissance Technologies, one of the most successful hedge fund ever, returned 39 percent per year on average between the end of 1989 and 2006. Its founder Jim Simons, a mathematician who gave notable contributions to string theory, a code breaker who worked at the Pentagon’s Institute for Defense Analyses during the cold war, a lifelong speculator and entrepreneur, assembled one of the most successful and secretive group of quantitative researchers ever. The following is the only information available on its website¹

Renaissance Technologies LLC is an investment management company dedicated to producing superior returns for its clients and employees by adhering to mathematical and statistical methods

This does not shed any light on how RenTech achieves its extraordinary results and the fact that almost nobody leaves the firm made its strategies one of the best-kept secret in the world. Some may argue that using successful hedge funds as a counterexample to market efficiency is affected by the well-known *survivorship bias*, that is the logical error of concentrating on the people or things that “survived” some process and inadvertently overlooking those that did not because of their lack of visibility. This bias can lead to overly optimistic

¹<https://www.rentec.com/>

beliefs because failures are ignored and to the false belief that the successes in a group have some special property, rather than just coincidence. In our opinion, RenTech's success is not a coincidence and perfectly shows how difficult it is to identify and profitably exploit market inefficiencies. For a richer account of this (and other) legendary hedge fund, the reader may refer to [69].

There is another important point that is worth mentioning. It is often argued that there is a "file drawer" bias in published studies due to the difficulty associated with publishing empirical studies that find insignificant effects. In studies of market efficiency, a reverse bias may be present. A researcher who genuinely believes he or she has identified a method for predicting the market has little incentive to publish the method in an academic journal and would presumably be tempted to sell it to an investment bank or to an hedge fund.

7.2 Bibliographical Survey

Many financial applications can be represented as sequential decision problems and can be naturally tackled with the reinforcement learning techniques presented in the previous chapters. Consider for instance the process of trading, in which an investor tries to select the assets that will perform the best in the future and will allow him to realize a profit. This activity is well depicted as an online decision problem involving the two critical steps of *market representation* and *optimal action execution*.

Financial markets are extremely complex systems, typically characterized by a large degree of non-stationarity, non-linearity and low signal-to-noise ratios. The first challenge is thus how to summarize the financial environment or, put in more statistical terms, how to design powerful features to be used as input of predictive models. The search for good indicators has been extensively studied in quantitative finance and econometrics. In particular, *technical analysis* is a security analysis methodology for forecasting the direction of prices through the study of past market data, primarily price and volume [65]. Whether technical analysis actually works is a matter of controversy. Methods vary greatly, and different technical analysts can sometimes make contradictory predictions from the same data. Many investors claim that they experience positive returns, but academic appraisals often find that it has little predictive power. Another issue is that designing hand-crafted features requires a deep knowledge of the financial markets. A solution to this drawback is offered by **DL**, a branch of machine learning based on a set of algorithms that attempt to model high level abstractions in data by using a deep graph with multiple processing layers, composed of multiple linear and non-linear transformations [44]. In the last years, deep learning has been in the spotlight as its application to fields like computer vision, automatic speech recognition, natural language processing, audio recognition and bioinformatics have produced state-of-the-art results on various tasks. These successes have attracted the interest of the financial community and the literature offers many examples of how deep neural network can be used

to predict future prices and invest on these forecasts. For example, the interested reader may refer to [53], [86], [63].

The second challenge is how to select the best possible action based on this representation of the market. This is not a trivial task, as the action selected by the investor may influence the market and modify the opportunities that will be available to him in the future. A simple example of this situation is the phenomenon of slippage, which consists in a loss due to the difference between the expected price of a trade and the price at which the trade is executed. This is common in limit-order books when there is no sufficient liquidity to complete a market order at the best available price and the order “walks the book”, obtaining progressively worse prices. Another difficulty is that the investor’s actions may impact its profits in a non-trivial way. For example, frequently changing its positions or entering in short positions will lead to large transaction costs which will undermine the performance of even the most accurate predictive trading strategy, such as those based on deep neural networks. Hence, these approaches are typically viable only on long investment horizons as their performances quickly degrade on shorter horizons because of transaction costs. Therefore, an investor should take into account the dynamic behavior of the market and of the returns he obtains. Reinforcement learning appears as the natural approach to determine the optimal strategy to be employed. In the following sections we briefly present some of the financial applications of reinforcement learning that can be found in the literature. This will give us a better feeling for the intrinsic difficulties of the financial markets and for how they can be tackled.

7.2.1 Asset Allocation with Transaction Costs

One of the first financial applications of reinforcement learning has been the problem of asset allocation with transaction costs. As discussed in the previous section, transaction costs introduce a dependence between the sequence of positions selected by the investor and its future returns. In this setting, predictive trading systems typically perform poorly as they do not explicitly take this mechanism into account. Therefore, a profitable trading system requires prior decisions as input in order to properly take into account the effects of transactions costs, market impact, and taxes. For this reason, in [74] the authors propose to learn a trading strategy that maximizes various performance measures, such as profit, economic utility, the Sharpe ratio and a differential Sharpe ratio, by [Recurrent Reinforcement Learning \(RRL\)](#) [47]. This consists in a policy gradient algorithm where the investor’s policy is recurrent, since it considers the action selected at the previous step as an input. This introduces a recursive dependence on the parameters of the policy and the policy is updated iteratively using standard techniques for recurrent neural networks, such as [Backpropagation Through Time \(BPTT\)](#) [109] or [Real-Time Recurrent Learning \(RTRL\)](#) [111]. The trading strategy learned in this way results profitable both on simulated and historical data. Moreover, the strategy reacts as

expected to transaction costs by reducing the reallocation frequency. In [72], the authors extend their previous work by comparing the performance of recurrent reinforcement learning with two standard value-based algorithms: TD-learning and Q-learning. These papers are among the first applications of reinforcement learning to finance to be published. Very similar approaches are presented in [30], [29], [43], [28], [34] and [62]. In [36], the authors improve the parametric strategy used by the trading system by introducing a deep neural network that uses some fuzzy representation of past returns as input. This approach greatly improves the representation ability of the policy which is able to extract much more complex features from historical data. This learned strategy is shown to be profitable both on historical data and outperforms purely predictive trading systems when transaction costs are considered. The asset allocation problem will be investigated further in the last section and will be used to test the state-of-the-art methods presented in the previous chapters.

7.2.2 Optimal Order Execution in Limit Order Book

Another problem that has been tackled with reinforcement learning is the *optimal order execution*, which consists in determining how to modify a portfolio from a given starting composition to a specified final composition within a specified period of time so as to maximize the expected revenue of trading (or equivalently minimizing the costs), with a suitable penalty for the uncertainty of revenue (or cost) [5]. In its simplest form, the problem can be formulated as follows: how to buy (or symmetrically sell) V shares of a given stock withing a time horizon T so as to minimize the total cost payed. The optimal execution problem can be viewed as an important horizontal aspect of quantitative trading, since virtually every strategy's profitability will depend on how well it is implemented.

This problem is of particular interest for a brokerage firm which may be asked by a client to execute an order for a certain number of shares before a given deadline. The broker will try to fulfill the client's order so as to maximize its revenue. Let us consider some naive strategies: the broker may decide to buy all the shares required straight away and wait until the deadline. Without considering market impact, the price to pay is known and no uncertainty is involved. However, the broker might obtain a better price before the client's deadline. This strategy involves more risk, since the price could move in the wrong direction and the broker could end up paying more than what it could have by following the previous strategy. If the order is sufficiently large compared to the liquidity available for a stock, market impact cannot be neglected and the broker will obtain progressively worse prices as the order is filled. Therefore, it is common to break large orders into smaller chunks so as to minimize slippage. In practice, the quality of execution is typically measured by the *implementation shortfall*, that is the difference between the decision price and the final execution price (including commissions, taxes, etc.) [49]. It seems natural to model the optimal execution problem in the stochastic optimal control framework: the goal of the

broker is to determine the execution schedule that minimize the implementation shortfall. Let us observe that the basic formulation of the problem given above is finite-horizon and therefore the optimal policy will be time-dependent.

In [76], the authors present the first large-scale application of reinforcement learning to the optimal execution problem in an electronic limit order book and show that a custom version of Q-learning results in substantial improvements in performance with respect to other usual execution policies. This approach is extended in [46], where the authors propose a hybrid approach which enhances a given analytical solution with attributes from the market microstructure. Using the Almgren-Chriss (AC) model [5] as a base, for a finite liquidation horizon with discrete trading periods, the algorithm determines the proportion of the AC-suggested trajectory to trade based on prevailing volume/spread attributes.

7.2.3 Smart Order Routing Across Dark Pools

Reinforcement learning has also been applied to the [Smart Order Routing \(SOR\)](#) problem. While the optimal execution problem consists in how to split an order across time, in the smart order routing problem one looks for the optimal split of an order across different trading venues. An increasingly popular source of liquidity for traders is represented by *dark pools*, a recent type of stock exchange in which information about outstanding orders is deliberately hidden in order to minimize the market impact of large trades [27]. In a typical dark pool, buyers and sellers submit unpriced orders to buy and sell securities, with the price derived exogenously from another market at a specified time. Upon submitting an order to buy (or sell) V shares, a trader is put in a queue of buyers (or sellers) awaiting transaction. Matching between buyers and sellers occurs in sequential arrival of orders, similar to a lit exchange. However, unlike a lit exchange, no information is provided to traders about how many parties or shares might be available in the pool at any given moment. Thus in a given time period, a submission of V shares results only in a report of how many shares up to V were executed.

The SOR problem can be formulated as follows: given K dark pools, what is the optimal way of splitting an order to buy V shares of a stock across these venues so as to maximize the completion rate. In [41], the authors propose to determine the optimal allocation using an approach similar to the reinforcement learning techniques used for *multi-armed bandits* [42]. What differentiates the SOR problem from most standard learning settings is that if an order for V_k shares is submitted to venue k and the order is completely filled, we only learn that at venue k the sell capacity is at least V_k , not the precise number of shares that could have been bought there. This mechanism is known as censoring in the statistics literature. The good performances of this algorithm demonstrate once again the potential of reinforcement learning in the field of algorithmic trading, where one seeks to optimize properties of a specified trade rather than decide what to trade in the first place. For a more thorough presentation of this application, the reader may also refer to [54] and the references therein.

7.3 Asset Allocation with Transaction Costs

The asset allocation problem consists in determining how to dynamically invest the available capital in a portfolio of different assets in order to maximize the expected total return or another relevant performance measure. Let us consider a financial market consisting of $I + 1$ different stocks that are traded only at discrete times $t \in \{0, 1, 2, \dots\}$ and denote by $Z_t = (Z_t^0, Z_t^1, \dots, Z_t^I)^T$ their prices at time t . Typically, Z_t^0 refers to a risk-free asset whose dynamic is given by $Z_t^0 = (1 + X)^t$ where X is the deterministic risk-free interest rate. The investment process works as follows: at time t , the investor observes the state of the market S_t , consisting for example of the past asset prices and other relevant economic variables, and subsequently chooses how to rebalance his portfolio, by specifying the units of each stock $n_t = (n_t^0, n_t^1, \dots, n_t^I)^T$ to be held between t and $t + 1$. In doing so, he needs to take into account the transaction costs that he has to pay to the broker to change his position. At time $t + 1$, the investor realizes a profit or a loss from his investment due to the stochastic variation of the stock values. The investor's goal is to maximize a given performance measure.

7.3.1 Wealth Dynamics

Let W_t denote the wealth of the investor at time t . The profit realized between t and $t + 1$ is simply given by the difference between the trading results and the transaction costs paid to the broker. More formally

$$\Delta W_{t+1} = W_{t+1} - W_t = \text{TP}_{t+1} - \text{TC}_t$$

where TP_{t+1} denotes the profit due to the variation of the asset prices between t and $t + 1$

$$\text{TP}_{t+1} = n_t \cdot \Delta Z_{t+1} = \sum_{i=0}^I n_t^i (Z_{t+1}^i - Z_t^i)$$

and TC_t denotes the fees paid to the broker to change the portfolio allocation and on the short positions

$$\text{TC}_t = \sum_{i=0}^I \delta_p^i |n_t^i - n_{t-1}^i| Z_t^i - \delta_f W_t \mathbb{1}_{n_t \neq n_{t-1}} - \sum_{i=0}^I \delta_s^i (n_t^i)^- Z_t^i$$

The transaction costs consist of three different components. The first term represents a transaction cost that is proportional to the change in value of the position in each asset. The second term is a fixed fraction of the total value of the portfolio which is paid only if the allocation is changed. The last term represents the fees paid to the broker for the shares borrowed to build a short position. The portfolio return between t and $t + 1$ is thus given by

$$X_{t+1} = \frac{\Delta W_{t+1}}{W_t} = \sum_{i=0}^I \left[a_t^i X_{t+1}^i - \delta_i |a_t^i - \tilde{a}_t^i| - \delta_s (a_t^i)^- \right] - \delta_f \mathbb{1}_{a_t \neq \tilde{a}_{t-1}} \quad (7.1)$$

where

$$X_{t+1}^i = \frac{\Delta Z_{t+1}^i}{Z_t^i}$$

is the return of the i -th stock between t and $t + 1$,

$$a_t^i = \frac{n_t^i Z_t^i}{W_t}$$

is the fraction of wealth invested in the i -th stock between time t and $t + 1$ and finally

$$\tilde{a}_t^i = \frac{n_{t-1}^i Z_t^i}{W_t} = \frac{a_{t-1}^i (1 + X_t^i)}{1 + X_t}$$

is the fraction of wealth invested in the i -th stock just before the reallocation. We assume that the agents invests all the available wealth at each step, so that W_t can be also interpreted as the value of the portfolio. This assumption leads to the following constraint on the portfolio weights

$$\sum_{i=0}^I a_t^i = 1 \quad \forall t \in \{0, 1, 2, \dots\} \quad (7.2)$$

We notice that we are neglecting the typical margin requirements on the short positions, which would reduce the available capital at time t . Considering margin requirements would lead to a more complex constraint on the portfolio weights which would be difficult to treat in the reinforcement learning framework. Plugging this constraint into Eq. (7.1), we obtain

$$X_{t+1} = X + \sum_{i=1}^I a_t^i (X_{t+1}^i - X) - \sum_{i=0}^I \left[\delta_i |a_t^i - \tilde{a}_t^i| - \delta_s^i (a_t^i)^- \right] - \delta_f \mathbb{1}_{a_t \neq \tilde{a}_{t-1}} \quad (7.3)$$

which highlights the role of the risk-free asset as a benchmark for the portfolio returns. The total profit realized by the investor between $t = 0$ and T is

$$\Pi_T = W_T - W_0 = \sum_{t=1}^T \Delta W_t = \sum_{t=1}^T W_t X_t$$

The portfolio return between $t = 0$ and T is given by

$$X_{0,T} = \frac{W_T}{W_0} - 1 = \prod_{t=1}^T (1 + X_t) - 1$$

In order to cast the asset allocation problem in the reinforcement learning framework, we consider the log-return of the portfolio between $t = 0$ and T

$$R_{0,T} = \log \frac{W_T}{W_0} = \sum_{t=1}^T \log(1 + X_t) = \sum_{t=1}^T R_t \quad (7.4)$$

where R_{t+1} is the log-return of the portfolio between t and $t + 1$

$$R_{t+1} = \log \left\{ 1 + \sum_{i=0}^I \left[a_t^i X_{t+1}^i - \delta_i |a_t^i - \tilde{a}_t^i| - \delta_s(a_t^i) \right] - \delta_f \mathbb{1}_{a_t \neq \tilde{a}_{t-1}} \right\} \quad (7.5)$$

The portfolio log-return can be directly used as the reward function of the asset allocation problem. However, more sophisticated approaches can be followed.

7.3.2 Rewards and Objective Functions

The results of the previous sections need to be adapted to the typical reinforcement learning formulation. Two distinct routes can be followed: first the asset allocation problem is seen as an infinite horizon task, either in the discounted or in the average reward formulation. Second, an episodic formulation is proposed. As discussed in the previous chapters, these scenarios require different learning algorithms.

7.3.2.1 Infinite Horizon Task

The portfolio log-return between $t = 0$ and a maturity T is given by

$$R_{0,T} = \sum_{t=0}^{T-1} R_{t+1}$$

where R_{t+1} is given in Eq. (7.5). The maturity T is arbitrary and one may be tempted to let $T \rightarrow \infty$. In addition to posing some convergence issues, this would be rather unrealistic as traders are typically evaluated on the short-run. As discussed in Chapter 2, the typical solution in this situation is to consider the discounted return

$$G_\gamma = \sum_{t=0}^{T-1} \gamma^t R_{t+1}$$

where $0 < \gamma < 1$ is the discount factor. This choice leads to the classical formulation for infinite horizon optimal control problems, which can be addressed in a fairly standard way using value-based or policy-based reinforcement learning algorithms. However, as discussed in the previous chapters, for the discounted rewards formulation it is difficult to extend the policy gradient theorem to the risk-sensitive framework. On the other hand, this is easier for the average reward formulation. In this case, the gain is measured as

$$G_{\text{avg}} = \lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=0}^{T-1} R_{t+1}$$

The relevant risk-neutral and risk-sensitive performance measures can be defined as in the previous chapters.

7.3.2.2 Episodic Task

Both formulations discussed above treat the asset allocation problem on an infinite horizon. On the other hand, it would be interesting to formulate this problem as an episodic task, for which many policy gradient algorithms are available as discussed in Chapter 4. To achieve this, we can follow the idea of [24].

Given an initial capital W_0 and two constants $L < 1 < U$, the goal of the trader is to increase his wealth to UW_0 , before going down to LW_0 . Let us reformulate this goal in terms of log-returns. First, let $l = \log L$ and $u = \log U$. We introduce the first time when the portfolio log-return breaches these thresholds

$$\tau_u = \inf\{t > 0 : R_{0,t} \geq u\} \quad (7.6)$$

$$\tau_l = \inf\{t > 0 : R_{0,t} \leq l\} \quad (7.7)$$

Hence, the goal of the agent can be formalized as follows

$$\max_{\pi} \mathbb{P}_{\pi}(\tau_u < \tau_l) = \mathbb{E}_{\pi}[\mathbb{1}_{\tau_u < \tau_l}] \quad (7.8)$$

This control problem can be interpreted as an episodic task in which the agent receives a reward either 1 if $R_{0,t} > u$ or 0 if $R_{0,t} < l$ and then terminates. This mechanism is illustrated in Figure 7.1. The upper cumulative log-return hits the upper barrier before hitting the lower barrier and consequently the agent receives a reward of 1. On the other hand, the lower path hits the lower barrier before and the agent thus receives zero reward.

This formulation may appear simple, but it offers a nice practical interpretation. The lower barrier can be seen as a stop-loss order, which forces the trader to exit from his position when the losses exceed a given threshold. We remark that the reward does not depend on the time step at which one of the two barrier is achieved. This could lead to extremely conservative policies which maximize the probability of success taking however a long time to achieve it. To attenuate this issue we might introduce a discount factor to represent the trader's preference to succeed as soon as possible. Formally, the optimization problem becomes

$$\max_{\pi} \mathbb{E}_{\pi}[\gamma^{\tau_u} \mathbb{1}_{\tau_u < \tau_l}] \quad (7.9)$$

where $0 < \gamma < 1$ denotes as usual the discount factor.

7.3.3 States

At each time step, the agent observes the state of the system and subsequently selects an action. First, we consider the $P + 1$ past returns of all risky assets $\{X_t, X_{t-1}, \dots, X_{t-P}\}$. These input variables may be used to construct more complex features for example using some deep learning techniques, such as a deep auto-encoder.

In order to properly incorporate the effects of transaction costs into his decision

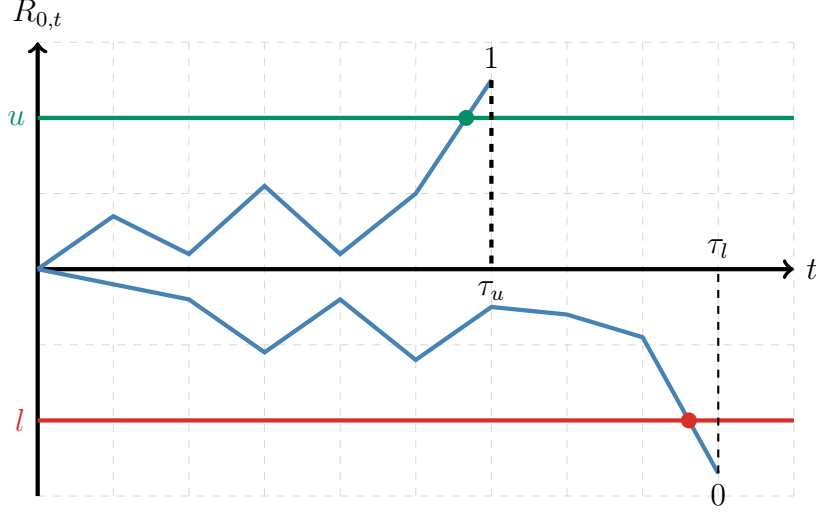


Figure 7.1: Episodic formulation for the asset allocation problem.

process, the agent must keep track of its current position \tilde{a}_t . For simplicity, we consider the past action as frozen and we will neglect the dependence from the policy, which will lead to the recurrent formulation proposed in [74].

Finally, we might consider some external variables Y_t that may be relevant to the trader, such as the common technical indicator used in practice. Summing up, the state of the system is given by

$$S_t = \{X_t, X_{t-1}, \dots, X_{t-P}, \tilde{a}_t, Y_t, Y_{t-1}, \dots, Y_{t-P}\} \quad (7.10)$$

In the episodic formulation, we might also include some feature measuring the distance of the cumulative log-return from the barriers

$$Y_t^u = u - R_{0,t} \quad Y_t^d = R_{0,t} - d \quad (7.11)$$

7.3.4 Actions

The agent, or trading system, only specifies the portfolio weights $a_t = (a_t^0, \dots, a_t^I)^T$ and therefore determines the allocation for the time interval $[t, t+1)$. If we assume that the agent invests all the available capital at each time step and that short-selling is not allowed, then the portfolio weights should satisfy, for every $t \in \{0, 1, \dots\}$, the following constraints

$$\begin{cases} a_t^i \geq 0, \forall i \in \{0, \dots, I\} \\ \sum_{i=0}^I a_t^i = 1 \end{cases} \quad (7.12)$$

This constraint can be easily enforced by considering a parametric softmax policy. If short-selling is allowed, weights might also be negative. A simple approach is to assume that, for every $t \in \{0, 1, \dots\}$, the weights satisfy

$$\begin{cases} a_t^i \in \mathbb{R}, \forall i \in \{1, \dots, I\} \\ a_t^0 = 1 - \sum_{i=1}^I a_t^i \end{cases} \quad (7.13)$$

Since a_t^0 is uniquely determined by the other weights, it is enough to define a policy that specifies the allocation in the risky assets, for example a Gaussian policy in \mathbb{R}^I . The obvious shortcoming of this approach is that the agent might enter in huge short positions, which is not realistic. A first observation is that, if the stochastic policy is concentrated around the origin, huge long or short positions would have very small probabilities of being selected. Moreover, we notice that the agent would pay large fees for entering into large short positions, independently of the trading profits. Therefore, we expect the agent to learn that short positions are very expensive and would therefore try to avoid them. Working in a continuous action space is computationally difficult and only few reinforcement learning algorithms are well-suited to this setting, for instance policy gradient methods. A simpler approach is to reduce the action space to a discrete space. For instance, in the two assets scenario we might assume that $a_t \in \{-1, 0, +1\}$. Thus the agent may be long (+1), neutral (0) or short (-1) on the risky-asset. Working in a discrete action space is more simple, and standard value-based approaches might also be employed.

Chapter 8

Numerical Results for the Asset Allocation Problem

In this chapter we present the numerical results of some of the policy gradient algorithms discussed in Chapter 4 for the asset allocation problem. Two different type of markets are analyzed: a market with only one risky asset and a market where multiple risky assets are available, for which finding a trading strategy is more difficult since the state and action spaces are much larger. The learning algorithms are first applied both in their risk-neutral version and risk-sensitive formulation to synthetically generated data, which present profitably tradable features. Once the behavior of these algorithms is validated in this controlled environment, the application on historical price series is considered.

8.1 Synthetic Risky Asset

To test the different reinforcement learning methods in a controlled environment, we generated log-price series for the risky asset as random walks with autoregressive trend processes. The two-parameter model is thus given by

$$\begin{aligned}z_t &= z_{t-1} + \beta_{t-1} + \kappa \epsilon_t \\ \beta_t &= \alpha \beta_{t-1} + \nu_t\end{aligned}$$

We then define the synthetic price series as

$$Z_t = \exp \left(\frac{z_t}{\max_t z_t - \min_t z_t} \right)$$

This model is often taken as a benchmark test in the automated trading literature, see for instance [75], because the price series generated in this way present some patterns that can be profitably exploited. Moreover the model is stationary and therefore the policy learned on the training set should generalize well on the test set, also known as backtest in the financial jargon. Thus we would expect our learning algorithms to perform well on this test case. If this wasn't the case, we should go back and improve the learning algorithms.

In this setting, we compare three the results of three long-short strategies obtained with ARAC, PGPE and NPGPE in both the risk-neutral and risk-sensitive framework. This means that the agent can either go long on the risky asset (i.e. $a_t^1 = 1$) or short the security (i.e. $a_t^1 = -1$) and invest the proceedings in the risk-less asset. Given the current conditions of the financial markets, we always assume a risk-free rate $X = 0$. Let us describe in more detail the the choice we made for each of the algorithms.

8.1.1 Specifications of the Learning Algorithms

Let us detail the choice of the parametric policies selected for each learning algorithm.

ARAC We considered a Boltzmann exploration policy on the two actions $a_t^1 \in \{-1, 1\}$ and a linear critic in which the features coincide with the agent's observation of the system state. This critic is extremely simple and there is surely some work to be done to improve it.

PGPE We considered a binary deterministic controller

$$F_\theta(s) = \text{sign}(\theta \cdot s)$$

where the parameters and the state also include a bias term. The controller parameters are sampled from a multi-variate Gaussian distribution

$$\theta \sim \mathcal{N}(\mu, \text{diag}(\sigma))$$

NPGPE We used the same controller as for PGPE but we assumed that the controller parameters are sampled from a Gaussian distribution parameterized by its mean and Cholesky factor

$$\theta \sim \mathcal{N}(\mu, C^T C)$$

8.1.2 Experimental Setup

All the algorithms were tested on the same price series of size 9000, generated from the process described above using $\alpha = 0.9$ and $\kappa = 3$. The learning process consisted of 1000 training epochs on the first 7000 days of the series with a learning rate that decreased at each epoch according to a polynomial schedule. The trained agents were subsequently backtested on the final 2000 days, during which the agents kept learning online in order to try to adapt to the changing environment. Since the price series is generated using a stationary model, it is not necessary to backtest the algorithm using the rolling-window approach typically employed in practice. The results that we present are the average of 10 independent experiments that used slightly different random initialization of the policy parameters.




Figure 8.1: Risk-neutral learning process for the asset allocation problem with one synthetic risky asset.

8.1.3 Risk-Neutral Framework

8.1.3.1 Convergence

Let us first discuss the case with no transaction costs. Figure 8.1 shows the learning curves for the three risk-neutral algorithms in terms of average daily reward, which is the quantity being maximized by the algorithms, the daily reward standard deviation and the annualized Sharpe ratio. The first thing we observe is the ARAC algorithm seems not to be improving the trading strategy as the training epochs go by. The average reward obtained is close to zero and will be surely be negative once transaction costs are introduced. On the other hand, NPGPE slowly converges to a profitable strategy which is however suboptimal compared to the one found by PGPE, that is better in all three measures considered. It is interesting to notice that PGPE and NPGPE yield a learning curve for the Sharpe ratio very similar to the one for the average reward. Even if the algorithm is risk-neutral, it manages to improve a risk-sensitive measure at the same time of the average reward. This might be simply a peculiarity of the very simple model assumed for the synthetic risky asset. Moreover, since the price process is stationary, the trading strategy learned on the training set perfectly generalizes to the test set.

8.1.3.2 Performances

Figure 8.2 compares the backtest performances of the three learned policies and a Buy and Hold strategy, which simply consists in investing all the available capital in the risky asset. Let us repeat that the solid lines are the averages of 10 independent experiments, which allows us to determine the 95% confidence intervals represented with the dashed lines. We clearly see that NPGPE and PGPE easily beat the market, realizing a total profit of 231.63% and 314.34% respectively against the 7.81% profit of the Buy and Hold strategy over the same period. Table 8.1 reports more performance statistics for the trading strategies




Figure 8.2: Backtest performance of trained trading systems for the asset allocation problem with one synthetic risky asset.

averaged over the independent experiments. We remark that PGPE and NPGPE beat the simple Buy and Hold strategy with respect to all measures, impressively achieving almost 100% of profitable years and consecutive 12 months periods. These statistics confirm that ARAC is not able to detect the profitable patterns in the synthetic price series and the learned strategy is close to randomness, with a 50% probability of reallocation (i.e. a coin flip). On the other hand, PGPE and NPGPE presents much lower reallocation frequencies. This seems promising for dealing with transaction costs, which penalize reallocations and short positions. In the next section we analyze in detail how the behavior of the learned strategies change with the introduction of transaction costs.

8.1.3.3 Impact of Transaction Costs

In the algorithmic trading literature there are many examples of strategies based on the prediction of future rewards starting from more or less complex indicators. However, as discussed in the previous chapter, the performances of these methods quickly degrade when transaction costs for changing the portfolio composition or for shorting a security are considered. Indeed, these methods simply invest based on the prediction of the future returns, without explicitly taking into account transaction costs. On the other hand, reinforcement learning algorithms should learn to avoid frequent reallocations or shorts thanks to the feedback mechanism between the learning agent and the system, thus generating better trading performances. In this section we analyze how the strategies learned by PGPE and by NPGPE change when gradually increasing the proportional transaction costs and the short-selling fees. Intuitively, we expect a progressive reduction of the frequency of reallocation and of shorting the risky asset.

	Buy and Hold	ARAC	NPGPE	PGPE
Total Return	7.81%	-0.86%	231.63%	314.34%
Daily Sharpe	0.27	-0.02	4.13	4.95
Monthly Sharpe	0.19	-0.07	2.90	3.26
Yearly Sharpe	0.23	-0.10	1.55	1.76
Max Drawdown	-22.35%	-12.60%	-3.72%	-3.27%
Avg Drawdown	-1.75%	-1.81%	-0.49%	-0.43%
Avg Up Month	2.87%	1.14%	2.47%	2.74%
Avg Down Month	-2.58%	-1.10%	-0.73%	-0.67%
Win Year %	40.00%	44.00%	98.00%	100.00%
Win 12m %	56.36%	48.00%	100.00%	100.00%
Reallocation Freq	0.00%	50.01%	19.99%	15.43%
Short Freq	0.00%	50.13%	41.59%	44.25%

Table 8.1: Backtest statistics of the risk-neutral trading strategies for the asset allocation problem with one synthetic risky asset. *Total Return* is the cumulative return obtained following the strategy. *Daily Sharpe* is the daily Sharpe ratio, annualized. *Monthly Sharpe* is the monthly Sharpe ratio, annualized. *Yearly Sharpe* is the yearly Sharpe ratio. *Max Drawdown* is the maximum drawdown observed, i.e. the maximum loss from a peak to a trough of a portfolio, before a new peak is attained. *Avg Drawdown* is the average drawdown observed, i.e. the average loss from a peak to a trough of a portfolio. *Avg Up Month* is the average profit on a positive month. *Avg Down Month* is the average loss on a negative month. *Win Year %* is the percentage of positive years. *Win 12m %* is the percentage of profitable consecutive 12 months. *Reallocation Freq* is the frequency with which the agent changes its position. *Short Freq* is the frequency with which the agent shorts the risky asset.

Figure 8.3 shows the impact of proportional transaction costs on the trading strategies learned by PGPE and by NPGPE. As expected, the frequency of reallocation for both strategies quickly drops to zero as the transaction costs increase, converging to the profitable buy and hold strategy. It is peculiar that the reallocation frequency for the PGPE strategy initially drops more quickly than for the NPGPE strategy, but then slows down and even increases when $\delta_P = 20$ bps. In summary, both algorithms are able to identify reallocation as the cause for lower rewards and to subsequently reduce the rate of reallocation. However, it seems that both algorithms prescribe a zero return buy and hold strategy on the risk-free asset, instead of the more profitable buy and hold strategy on the risky asset.

Figure 8.4 shows the impact of short-selling fees on the trading strategies learned by PGPE and NPGPE. Both algorithms behave as expected, displaying a progressive reduction of the frequency of short positions as the fees increase. For large values of short-selling fees, both strategies converge to the profitable buy and hold strategy, which completely avoids paying the fees. In particular, PGPE




Figure 8.3: Impact of proportional transaction costs on the trading strategies learned by PGPE and NPGPE.




Figure 8.4: Impact of short-selling fees on the trading strategies learned by PGPE and NPGPE.

quickly replicates the buy and hold strategy. On the other hand, NPGPE is not able to exactly reproduce the buy and hold strategy but it seems to converge to it for very large values of the short-selling fee.

8.1.4 Risk-Sensitive Framework

In this section we present the results in the risk-sensitive framework, in which the learning algorithms optimize the Sharpe ratio of the policy. Figure 8.5 shows the learning curves for the three risk-sensitive algorithms RSPGPE and RSNPGPE. We observe that the natural-gradient algorithm converges to a larger Sharpe ratio compared to the non-natural algorithm, which converges very quickly to a suboptimal policy. It is interesting that the risk-sensitive algorithms converge faster than their risk-neutral counterparts. Surprisingly, the strategies learned by the risk-sensitive algorithms have smaller Sharpe ratio than those learned in

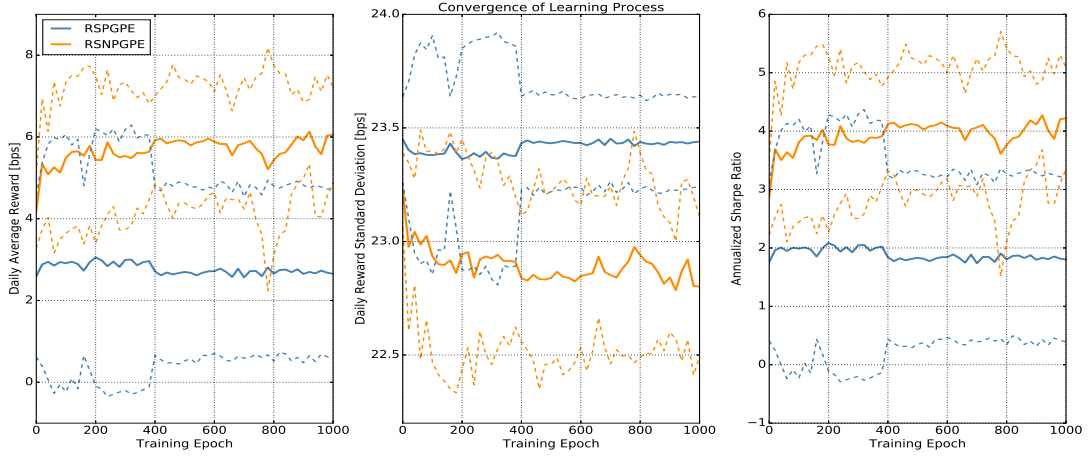


Figure 8.5: Risk-sensitive learning process for the asset allocation problem with one synthetic risky asset.

the risk-neutral version, in which we don't explicitly control risk. We still don't have a clear explanation for this counter-intuitive behavior. A first hypothesis is that, for the stationary dynamics considered to generate the price series, the optimal policy in the risk-neutral sense also minimize risk, but directly optimizing the Sharpe ratio is more difficult as more parameters and noise are involved in the learning process. Further research should be carried out to investigate this issue in more depth. Figure 8.6 shows the backtest performances for the risk-sensitive trading strategies, which both outperform the simple buy-and-hold strategy. More statistics are reported in Table 8.2. We notice that, contrarily to the risk-neutral setting, NPGPE performs better than PGPE with respect to all the performance measures. In particular, the policy learned by PGPE takes much less short positions, which negatively impacts the performance of the strategy.




Figure 8.6: Backtest performance of the risk-sensitive trading strategies for the asset allocation problem with one synthetic risky asset.

	Buy and Hold	RSPGPE	RSNPGPE
Total Return	7.81%	82.29%	219.81%
Daily Sharpe	0.27	1.95	3.97
Monthly Sharpe	0.19	1.19	2.84
Yearly Sharpe	0.23	0.67	1.71
Max Drawdown	-22.35%	-11.56%	-3.60%
Avg Drawdown	-1.75%	-1.00%	-0.47%
Avg Up Month	2.87%	2.65%	2.41%
Avg Down Month	-2.58%	-1.48%	-0.71%
Win Year %	40.00%	76.00%	100.00%
Win 12m %	56.36%	87.45%	100.00%
Reallocation Freq	0.00%	15.50%	28.00%
Short Freq	0.00%	16.17%	45.77%

Table 8.2: Backtest statistics of the risk-sensitive trading strategies for the asset allocation problem with one synthetic risky asset.

	NPGPE	RSNPGPE
Total Return	231.63%	219.81%
Daily Sharpe	4.13	3.97
Monthly Sharpe	2.90	2.84
Yearly Sharpe	1.55	1.71
Max Drawdown	-3.72%	-3.60%
Avg Drawdown	-0.49%	-0.47%
Avg Up Month	2.47%	2.41%
Avg Down Month	-0.73%	-0.71%
Win Year %	98.00%	100.00%
Win 12m %	100.00%	100.00%
Reallocation Freq	19.99%	28.00%
Short Freq	41.59%	45.77%

Table 8.3: Comparison of NPGPE and RSNPGPE for the asset allocation problem with one synthetic risky asset.

8.1.4.1 Risk-Neutral vs. Risk-Sensitive

The risk-neutral NPGPE and the risk-sensitive NPGPE strategies achieve very similar backtest performances. A natural question that arises is if the two algorithms are actually learning the same policy. The two versions of the NPGPE algorithm are compared in Table ?? . The statistics for the two algorithms are similar. Again, we point out the counter-intuitive fact that risk-neutral version produces higher Sharpe ratios, even if it is not the quantity directly being optimized. On the other hand, the risk-sensitive version seems to perform slightly better in terms of drawdowns. The most interesting feature is the different behavior of the two policies in terms of reallocation and short-selling frequencies. The risk-sensitive policy reshuffle the portfolio more frequently than the risk-neutral policy, also resorting more often to short positions. This shows that the policies learned by the two algorithms are actually different, leading however to very similar performances.

8.1.4.2 Impact of Transaction Costs

Once again we study the sensitivity of the trading strategies learned by RSPGPE and RSNPGPE with respect to transaction costs and short-selling fees. Intuitively, as the transaction costs increase we expect a progressive reduction of the frequency of reallocation and of shorting the risky asset.

Figure 8.7 shows the impact of proportional transaction costs on the trading strategies learned while Figure 8.8 shows that of short-selling fees. We observe that RSNPGPE quickly converge to the profitable Buy and Hold strategy in both situations. This was not the case in the risk-neutral framework, where NPGPE underperformed the benchmark when the transaction costs were large enough. Disappointingly, RSPGPE doesn't adapt to the increasing transaction




Figure 8.7: Impact of proportional transaction costs on the trading strategies learned by RSPGPE and RSNPGPE.




Figure 8.8: Impact of short-selling fees on the trading strategies learned by RSPGPE and RSNPGPE.

costs as expected, leading to large losses. While, the reallocation frequency decreases (but not enough to avoid a loss), the frequency of short position appears to be independent of the short-selling fees. Further analysis is necessary to explain this unexpected behavior of RSPGPE which is not observed in the risk-neutral framework.

8.2 Historical Risky Asset

In this section we discuss the attempt to apply the algorithms above to an historical price series. As a test case, we considered a Banca Monte dei Paschi di Siena (BMPS IM) stock as the risky asset. The experimental setup is the following: the learning algorithms presented above have been trained during 5000 epochs over a period of 1000 days. Then, the performance of the learned trading strategy has been assessed on the following 100 days. As an alternative, we could have employed a more standard “rolling-window” backtesting procedure. In this case, the algorithms perform poorly and are not able to identify a profitable trading strategy. In some measure, this was expected given the simplicity of the parametric policy considered. When working on historical price series, there are multiple challenges to be addressed. First, it is not clear if the price series contains some patterns that could be traded profitably. This is the central issue of the [EMH](#) which was discussed in [Section 7.1](#). Now, even if the prices contained a certain degree of predictability, the parametric policy should be sufficiently expressive to capture these patterns. This is clearly an issue of the current state of work, as the policy considered here is so simple that it would be surprising if it managed to capture a profitably tradable structure in the prices of a liquid stock. A possible development this work would be to employ a much more complex parametric policy, such as a neural network, in order to automatically extract more powerful features from the raw price series. Another related issue is the signal-to-noise ratio of financial time-series. As the policy is updated online, too much noise could mislead the policy gradient method and prevent its convergence.

8.2.1 Risk-Neutral Framework

These seems to be the issue in the application of the risk-neutral algorithms developed above to the historical data considered. As showed in [Figure 8.9](#), the policy gradient methods fail to converge and the learning curves appear to be dominated by noise. [Figure 8.10](#) shows the average profits generated by the learned strategies on the backtest set in 10 independent experiments. We see that on average the strategies learned with the ARAC and NPGPE algorithms slightly overperform the simple Buy and Hold strategy. However, the noise is so large that we cannot conclude that the learned strategies are actually better than the Buy and Hold. In addition, since the stock price remains almost flat during the test period a purely random trading strategy would not perform much differently than the Buy and Hold strategy. We suspect that the strategies learned would underperform in a more trending situation. [Table 8.4](#) reports a breakdown of the performance measures of the learned strategies on the backtest set. Again, we observe that the ARAC algorithm has a reallocation frequency of around 50%, which reinforce our suspicion that this strategy actually consists in a coin flip.




Figure 8.9: Risk-neutral learning process for the asset allocation problem with an historical risky asset.




Figure 8.10: Backtest performance of the trading systems learned for the asset allocation problem with one historical risky asset.

	Buy and Hold	ARAC	NPGPE	PGPE
Total Return	3.29%	7.68%	4.11%	-11.72%
Daily Sharpe	0.48	0.79	0.55	-1.37
Monthly Sharpe	-0.25	-0.48	-0.33	-2.38
Max Drawdown	-14.68%	-12.33%	-12.36%	-20.46%
Avg Drawdown	-6.90%	-7.15%	-5.12%	-9.61%
Avg Up Month	4.43%	8.27%	3.66%	3.02%
Avg Down Month	-2.65%	-2.49%	-4.00%	-5.98%
Reallocation Freq	0.00%	50.20%	19.39%	23.47%
Short Freq	0.00%	49.80%	19.80%	38.40%

Table 8.4: Backtest statistics of the risk-neutral trading strategies for the asset allocation problem with one historical risky asset.

8.2.2 Risk-Sensitive Framework

When applied to the historical data, the risk-sensitive algorithms shows a quite different behavior compared to the risk-neutral methods. Figure 8.11 shows the learning curves of the algorithms during the different training epochs. While the RSPGPE algorithms seems not to improve the policy over time, the policy learned with the RSNPGPE algorithm clearly improves with the training epochs and achieves a positive Sharpe ratio. This means that the simple autoregressive controller considered is able to capture some patterns in the training set which leads to positive performance. However it is not guaranteed that what is learned on the training set will hold during the backtest. This is exactly the situation as showed in Figure 8.12. Indeed, we remark that the policy learned by the RSNPGPE algorithm, although producing a profit on the training set, clearly underperforms the naive Buy and Hold strategy in the backtest. We conclude that the pattern learned in the training set does not generalize to the backtest set. This phenomenon could be seen as a form of overfitting which, however, is not due to the choice of an over-complicated parametric policy (it is clearly not the case here), but to the inconsistency between the training set and the backtest set. This is typical of financial time series, which are often characterized by a strong non-stationarity. For this reason, even if the controller is able to identify a signal in the training set, it is not guaranteed that this signal will persist in the backtest set. Even if the risk-sensitive algorithms that were developed in this thesis fail to learn a trading strategy which is profitable on the backtest set, it is still remarkable that the RSNPGPE algorithm manages to learn a strategy that is profitable on the training set while the risk-neutral algorithms could not. This indicates that there is some value in the algorithms we proposed, even if they could not succeed in the extremely difficult task of beating the market

8.2.3 The Challenge of Historical Data

In the previous sections, we presented the application of the reinforcement learning algorithms we developed in this thesis to the historical price series. In particular, we discussed how these algorithms fail to learn a trading strategy which consistently outperforms the naive Buy and Hold strategy in the backtest phase. This outcome was partly expected given the simplicity of the parametric policy considered. However, even if we considered a more complex controller, success would not be a-priori guaranteed. This is due to the inherent complexity of the financial markets. Indeed, while identifying patterns in historical data is not particularly hard, finding signals that work reliably in the real world is. Man AHL, a quant unit of Man Group Plc, needed three years of work to gain enough confidence in a machine learning strategy to devote client money to it. Therefore, it is clear that there is still much work to be done to be able to beat the markets with a strategy produced by a reinforcement learning algorithm.




Figure 8.11: Risk-sensitive learning process for the asset allocation problem with an historical risky asset.




Figure 8.12: Backtest performance of the trading systems learned for the asset allocation problem with one historical risky asset.

	Buy and Hold	RSPGPE	RSNPGPE
Total Return	3.29%	0.20%	-8.12%
Daily Sharpe	0.48	0.12	-0.93
Monthly Sharpe	-0.25	0.42	0.12
Max Drawdown	-14.68%	-17.02%	-19.40%
Avg Drawdown	-6.90%	-9.25%	-9.22%
Avg Up Month	4.43%	4.67%	2.53%
Avg Down Month	-2.65%	-6.45%	-5.15%
Reallocation Freq	0.00%	19.80%	39.59%
Short Freq	0.00%	24.00%	37.80%

Table 8.5: Backtest statistics of the risk-sensitive trading strategies for the asset allocation problem with one historical risky asset.

Chapter 9

Conclusions

In this chapter, we take the time to revisit what has been achieved in this thesis and to sum up our original contributions to the reinforcement learning literature. Moreover, we draw some conclusions about the impact that these modern optimization techniques can have on the financial industry and the challenges that must be overcome to make it possible. Finally, we suggest some possible axis for future research.

9.1 What Has Been Achieved

In this thesis, we presented an innovative application of some state-of-the-art policy gradient algorithms to the classical asset allocation problem with transaction costs.

In Chapter 2 we introduced the basic concepts of stochastic optimal control in discrete time, which is the standard theoretical framework used to model sequential decision problems. In particular, we presented both the traditional risk-neutral formulation, in which the goal of the agent is to maximize its total returns, and the less common risk-sensitive formulation, in which the agent wishes instead to optimize the trade-off between the rewards and the risk required to achieve them. This chapter laid the theoretical foundations and introduced the notation used in the rest of the thesis.

In Chapter 3 we presented the main ideas of [Reinforcement Learning \(RL\)](#), a general class of algorithms in the field of [Machine Learning \(ML\)](#) that allows an agent to learn how to behave in a stochastic and possibly unknown environment only by trial-and-error. These algorithms are typically based on the feedback mechanism between the stochastic environment and the agent, which receives a numerical reward for his actions. After having discussed the characteristic features of the [RL](#) problem, we provided a high-level overview of the different typologies of algorithms.

In Chapter 4 we gave an in-depth presentation of Policy Gradient algorithms for the risk-neutral control problem. After having introduced the key ideas of these methods, we provided a thorough review of the state-of-the-art algorithms that can be found in the literature. For episodic environments, most of these

methods are based on the well-known likelihood-ratio technique from stochastic optimization that allows to rewrite the policy gradient as an expected value. For continuing environments, the main result from which most of the learning algorithms are derived is the policy gradient theorem.

In Chapter 5 we discussed policy gradient methods for the risk-sensitive control problem, which is still an active field of research. In particular, we presented an extension of the policy gradient theorem to the risk-sensitive framework, both in the average reward and in the discounted reward formulations. To the best of our knowledge, this is the first time that a risk-sensitive policy gradient theorem is proved for a general discounted Markov decision process, in which rewards might as well depend on the next state of the system.

In Chapter 6 we proposed an original parameter-based version of the policy gradient theorem, both for the risk-neutral and the risk-sensitive framework. This extension allowed to derive efficient online learning algorithms similar in spirit to the well-known PGPE algorithm, which was originally conceived only for episodic environments. Moreover, these new algorithms can be easily enhanced using a critic or the natural policy gradient idea. This chapter undoubtedly represents the most innovative contribution of this thesis to the RL literature.

In Chapter 7, after a brief discussion about why finance represents an extremely challenging field of research, we provided a bibliographical survey of successful applications of RL techniques to financial problems. In particular, to numerically test the algorithms proposed in the previous chapters, we focused on the classical asset allocation problem with transaction cost.

In Chapter 8 we presented the numerical results for the asset allocation problem. We showed that the learning algorithms proposed in the previous chapters perform extremely well on synthetic data, consistently outperforming a simple buy-and-hold strategy and automatically adapting to transaction costs. On the other hand, the algorithms encountered more difficulties on historical data and we tried to provide an explanation for this behavior.

9.2 Conclusion

The main contribution of this thesis to the RL literature is the development of some original, online and parameter-based policy gradient algorithms both for the traditional risk-neutral framework and the less common risk-sensitive framework. These algorithms are highly versatile and can be easily combined with some powerful techniques such as the use of a critic or the use of a natural policy gradient.

The numerical application of this innovative algorithms to the classical financial problem of determining a profitable long-short trading strategy confirmed their potential. For a synthetic asset, the strategies produced outperform the simple Buy & Hold strategy, even when investment decisions are only based on extremely basic autoregressive features. Contrarily to standard prediction-based trading systems, the learning algorithms employed are able to adapt as expected

to the introduction of transaction costs by reducing the frequency of reallocation and of short positions.

The algorithms encountered more difficulties in learning a profitable strategy from historical data. In this case, the algorithms failed to converge to a profitable trading strategy probably because of the low signal-to-noise ratio of the historical price series considered and the low representative capabilities of the controller employed. However, even if we considered a more complex controller success would not be a-priori guaranteed. This is due to the inherent complexity of financial markets. Indeed, while identifying patterns in historical data is not particularly hard, finding signals that work reliably in the real world is. Man AHL, a quant unit of Man Group Plc, needed three years of work to gain enough confidence in a machine learning strategy to devote client money to it. Therefore, it is clear that there is still much work to be done to be able to beat the markets with a strategy produced by a reinforcement learning algorithm. Despite the difficulties we encountered in applying the algorithms developed in this thesis to historical data, we feel confident that these techniques will have a large impact on financial markets in the future. The first reason is that these algorithms have already proven successful in solving very difficult decision problems, such as mastering the game of Go or controlling robots in complex environments. The second reason is the versatility of these techniques, which can be applied with only minor modifications to very different domains. Finally, their modularity allows them to be combined with other modern techniques, such as [Deep Learning \(DL\)](#). These features have attracted the interest of the financial community as well, with investment banks and hedge funds investing large amount of resources to develop new innovative solutions based on [ML](#).

9.3 Further Developments

To conclude, we suggest some research directions that could be explored to improve the work presented in this thesis.

The first idea is to develop more complex features so as to allow the agent to take more informed decisions. A first approach would be to construct the classical indicators used in the technical analysis of stock prices. This measures embed the expert knowledge of many financial analysts and could help in guiding the agent towards better decisions. Another approach would be to employ some [DL](#) techniques to learn more powerful features directly from data. A first possibility would be to construct new features using a deep auto-encoder in a completely unsupervised way that would then be fed to a simple parametric controller. A second possibility would be to directly replace the simple autoregressive controller considered in this thesis with a deep neural network. Using a deep neural network was one of the key ingredient of many of the recent [RL](#) successes. For a modern presentation of these techniques, the interested reader may refer to [\[44\]](#). The second idea would be to consider high-frequency data, such as intra-day prices and traded volumes. It is well documented that high-frequency data show some characteristic patterns that could be captured by a deep reinforcement

learning trading system. An example of application of [RL](#) to high-frequency data can be found in [\[76\]](#), where the authors applied some basic algorithms such as Q-learning to the optimal trade execution problem in a limit-order book. It would thus be interesting to apply more advanced techniques to this problem. As the [RL](#) and [DL](#) literatures keep growing at a frenetic rhythm, new research directions keep opening up and new applications in different fields provide plenty of food for thought. Given the versatility of these techniques, we are convinced that a breakthrough application to finance is not far away.

Appendix A

Implementation

In this chapter we give a brief overview of the implementation of the algorithms presented above. In the prototyping phase, the asset allocation learning task has been implemented in Python as an extension of the PyBrain library. In a second phase, the architecture has been translated into C++ to achieve better performances. Here we only give an overview of the program, addressing the reader to the Git repository www.github.com/pnecchi/Thesis for the scripts and the detailed documentation.

A.1 Python Prototype

The first step of this project has been to implement a prototype in Python, a high-level, general-purpose, interpreted, dynamic programming language which is gaining a widespread popularity both in the academic world and in the industry. Python natively supports the object-oriented paradigm which makes it perfect to quickly develop a prototype of the class architecture, which can then be translated in C++. Moreover, thanks to external libraries such as Numpy, Scipy and Pandas, Python offers an open-source alternative to Matlab for scientific computing applications.

For the basic RL algorithms we exploited PyBrain¹, a modular ML library for Python whose goal is to offer flexible, easy-to-use yet still powerful algorithms for ML tasks and a variety of predefined environments to test and compare different algorithms [90]. An RL task in PyBrain always consists of an **Environment**, an **Agent**, a **Task** and an **Experiment** interacting with each other as illustrated in Figure A.4.

The **Environment** is the world in which the **Agent** acts and is characterized by a state which can be accessed through the `getSensors()` method. The **Agent** receives an observation of this state through the `integrateObservation()` method and selects an action through the `getAction()` method. This action is applied to the **Environment** with the `performAction()` method. However, the interactions between the **Environment** and the **Agent** are not direct but are

¹<http://pybrain.org/>




Figure A.1: PyBrain standard architecture for an RL problem.

mediated by the **Task**. The **Task** specifies what the goal is in an **Environment** and how the agent is rewarded for its actions. Hence, the composition of an **Environment** and a **Task** fully defines the MDP. An **Agent** always contains an **Actor**, which represents the policy used to select actions. Based on the rewards that the **Agent** receives via the `getReward()` method, the **Learner** improves the policy via a `learn()` procedure. In this step, an **Actor** may be used to evaluate a state with the goal of reducing the variance of the learning process. This entire learning process is controlled by an **Experiment** object.

This structure is quite standard for a RL problem and can be easily adapted to the problem at hand and extended to the learning algorithms developed in this thesis. Based on this architecture, we thus developed a fully-working Python prototype of the asset allocation problem. This prototype yielded some interesting results both on simulated data and on historical data, in particular for the PGPE algorithm. However, the learning process resulted too slow to be run systematically for a large number of time-steps and training epochs. By consequent, we quickly decided to pass to C++.

A.2 C++ Implementation

Passing from Python to C++ presents some challenges in the design of a suitable architecture for the RL algorithms discussed above. Following the approach of [50], our main goal has been code reusability, which is based on the important attributes of clarity and elegance of design. In addition, we always kept in mind the possibility that our original design might need to be extended in the future. In some cases we therefore favored extendability compared to efficiency. First we describe the C++ adaptation of the PyBrain's **Environment**, **Task**, **Agent** and **Experiment** objects with a particular attention to their concrete implementations for the asset allocation problem. Secondly, we discuss the design for an Average-Reward Actor-Critic agent (ARAC), which provides a concrete

implementation of the **Agent** interface.

A.2.1 Environment, Task, Agent and Experiment

Figure A.2 schematically represents the design of the base components of an RL application, which closely replicates Pybrain’s architecture. The pure abstract classes **Environment**, **Task**, **Agent** and **Experiment** define the generic interfaces to which all the concrete implementations of these objects must adhere. To achieve code modularity, we make most of the objects in our design clonable in order to allow for the polymorphic composition of classes. Exploiting this design pattern, we couple a **Task** with an **Environment** by storing a `std::unique_ptr<Environment>` as a private member of the class. Similarly, an **Experiment** is coupled via composition with a **Task** and an **Agent**. The methods of these classes are similar to those in Pybrain. For all the linear algebra operations we decided to use Armadillo², a high quality linear algebra library which provides high-level syntax (API) deliberately similar to Numpy and Matlab aiming towards a good balance between speed and ease of use. Therefore, the state of the system and the actions performed by the agent are represented as `arma::vector` objects. Let us now present the concrete implementation of these objects for the asset allocation task. **MarketEnvironment** implements a financial market by storing the historical daily returns of the risky assets in an `arma::matrix`. These values are read from a given input `.csv` file, which is generated automatically running a Python script and which either contains real historical returns downloaded from Yahoo Finance³ or synthetic returns generated according to a certain stochastic process. Therefore, we always work on samples of the system without making any assumption on its Markov transition kernel.

²<http://arma.sourceforge.net/>

³<https://uk.finance.yahoo.com/>

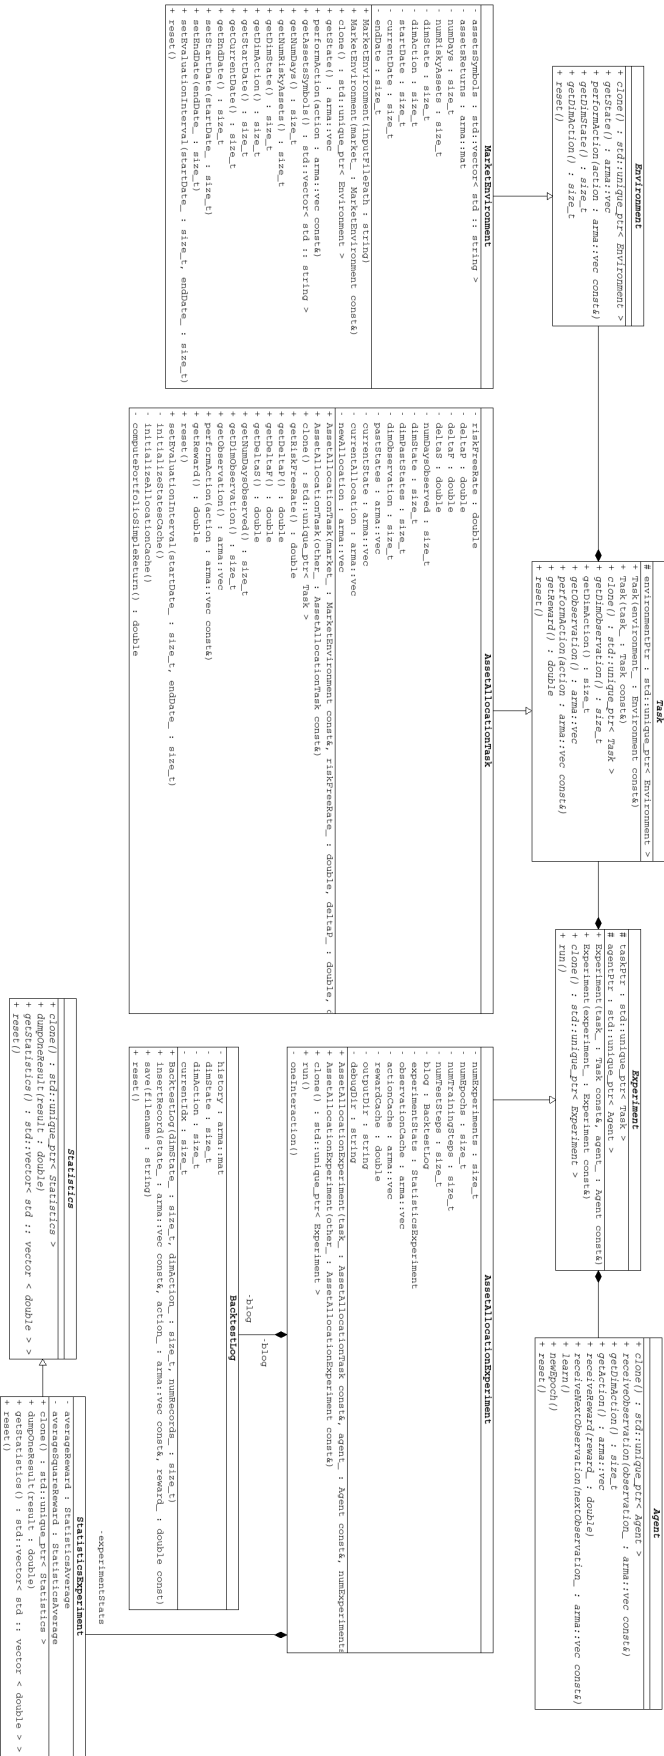


Figure A.2: Class architecture for the learning process in the asset allocation problem.

The `AssetAllocationTask` implements the MDP associated to the asset allocation problem by providing a method to compute the reward that the `Agent` obtains from investing on the risky assets. The `AssetAllocationTask` enlarges the system state so that the `Agent` also observes the past P states and the current portfolio weights. The `AssetAllocationExperiment` handles the interactions between the `Agent` and the `AssetAllocationTask`. The learning process is divided in two phases: the training phase consists of a certain number of learning epochs over the same time period during which the `Agent` improves the parameters of its policy via the `learn` method. Some estimates of the objective function are dumped in the `StatisticsExperiment` object and are used in the post-processing phase to plot the learning curves of the algorithm. In the backtest phase, the `Agent` applies the learned policy on the time period which follows the one used for training and the relevant performance measures are stored in the `BacktestLog` for successive analysis and comparison between different learning algorithms.

A.2.2 ARACAgent

A concrete implementation of the `Agent` interface completes the standard structure of an RL task. In this section we discuss the design for an Average Reward Actor-Critic agent (ARAC), which includes most of the features of the other algorithms tested in this thesis. The full architecture of this agent is illustrated in Figure A.3, but for brevity we will only focus on the more interesting aspects. The `ARACAgent` builds upon a `StochasticActor` and a `Critic` via composition. A `StochasticActor` is simply a wrapper around a `StochasticPolicy`, a pure abstract class defining the generic interface for a stochastic policy used by an agent to select actions. In addition to `getAction` and `get/set` methods for the policy parameters, a `StochasticPolicy` must implement the `likelihoodScore` method which computes the likelihood score $\nabla_{\theta} \log \pi_{\theta}(s, a)$ for a given state and action and which plays a crucial role in any policy gradient algorithm.

We provide two examples of concrete implementations of the `StochasticPolicy`. The first example is the `BoltzmannPolicy` typically used in discrete action spaces. The implementation of this policy is quite straightforward and we won't discuss the details. The second and more interesting stochastic policy implemented is the `PGPEPolicy`. This class is based on the decorator design pattern which is typically used to extend the interface of a certain class. Indeed, the `PGPEPolicy` is a `StochasticPolicy` which contains by polymorphic composition a `Policy`, a pure abstract class which provides the generic interface for a policy, potentially deterministic. This `Policy` object represents the deterministic controller F_{θ} used in the PGPE algorithm. Moreover, the `PGPEPolicy` contains by polymorphic composition a `ProbabilityDistribution`, a pure abstract class defining the generic interface for a probability distribution. This probability distribution represents the hyper-distribution p_{ξ} on the controller parameters. In order to be used in a PGPE algorithm, a `ProbabilityDistribution` must implement a `likelihoodScore` method to compute the likelihood score

of the hyper-distribution $\nabla_{\xi} \log p_{\xi}(\theta)$. Hence, the `likelihoodScore` method of `PGPEPolicy` simply broadcasts the call to the `likelihoodScore` method of its underlying `ProbabilityDistribution`. A concrete implementation of a `ProbabilityDistribution` is provided by the `GaussianDistribution`, which implements a multi-variate and axis-aligned normal distribution $\mathcal{N}(\mu, \text{diag}(\sigma))$. The objects discussed so far are sufficient to implement an actor-only learning algorithm, potentially using a baseline to evaluate the rewards. A more advanced variance reduction technique consists in using a `Critic`, which approximates the value function and provides an evaluation of a given state. The `Critic` class is simply a wrapper around a `FunctionApproximator` which provides a generic interface for a parametric function $B_{\omega}(s)$. The key methods of this class are `evaluate`, which evaluates the approximator at a given point, and `gradient`, which computes its gradient at a given point.

Finally, the `ARACAgent` employs some `LearningRate` to control the speed of the gradient descent optimization algorithm. A naive approach is to use a `ConstantLearningRate`, but this leads to a large-variance in the objective function value attained by the stochastic optimization algorithm. A more sensible choice is to use a `DecayingLearningRate` which decreases with the number of learning epochs performed by the agent according to $\alpha_n = \frac{a}{n^b}$. In this way, the learning process progressively “cools down” (using a simulated annealing terminology) and stabilizes to a given policy. This concludes our quick overview of the class architecture used for this project. In the thesis, other applications and algorithms were considered and we refer the reader to the full document for a more complete discussion.



Figure A.3: Class architecture for an Average Reward Actor-Critic agent (ARAC).

A.3 Execution Pipeline

In this section we describe the full pipeline of the program, which is schematically represented in Figure A.3. This pipeline allows to run the learning algorithm for the asset allocation problem and automatically determine a trading strategy. The execution consists of the following steps.

A.3.1 Compilation

To build the `thesis` library it is sufficient to run the `Makefile` generated with `cmake`. This produces two executables: `main` which is used to debug the program in the `Code::Blocks` IDE and `main_thesis` which is used to run the experiment in the full execution pipeline.

A.3.2 `generate_synthetic_series.py`

This Python script simulates the returns of a synthetic asset and prints them on a `.csv` file which is then read by `main_thesis` and used to initialize the `MarketEnvironment` object. Alternatively, `market_data_collector.py` collects the historical returns for a list of given assets from Yahoo finance.

A.3.3 `experiment_launcher.py`

This Python script manages the execution pipeline. First, the experiment parameters are specified and dumped in a `.pot` file which is then read by `main_thesis`. Given the parameter values, the script determines the folders where the output should be written so that the results can be easily associated to a specific set of parameters. Subsequently, it launches the `main_thesis` executable passing the correct parameters via the command line. Finally it runs the `postprocessing.py` scripts which processes the output files.

A.3.4 `main_thesis`

This executable takes some inputs from the command line, such as the algorithm to use, the paths to the input files and the paths where the output files should be generated. The experiment parameters are then read from the `.pot` file generated by the `experiment_launcher.py` using `GetPot`. The type of learning algorithm is specified by a string passed to the executable via command line and then used by the factory `FactoryOfAgents` to instantiate the corresponding `Agent`. When the `AssetAllocationExperiment` is run, it outputs various statistics to the given destination folders. More in detail, it prints two files for every independent run of the experiment: a `debug.csv` file which contains the learning curves of the algorithm and an `output.csv` file which contains the backtest performance measures for the trading strategy learned by the `Agent` during training.

A.3.5 postprocessing.py

This Python script processes the various files produced by `main_thesis` and generates an aggregate analysis of the various learning algorithms, so that they can be easily compared and assessed. In particular, it computes the average and confidence intervals for the learning curves of the algorithms and the backtest cumulative profits of the learned strategies. Moreover, it computes some performance measures typically used in Finance to evaluate a trading strategy, such as the Sharpe ratio and the maximum drawdown. The results of this analysis are stored in some `.csv` files and some images are generated using the Python library `matplotlib`.

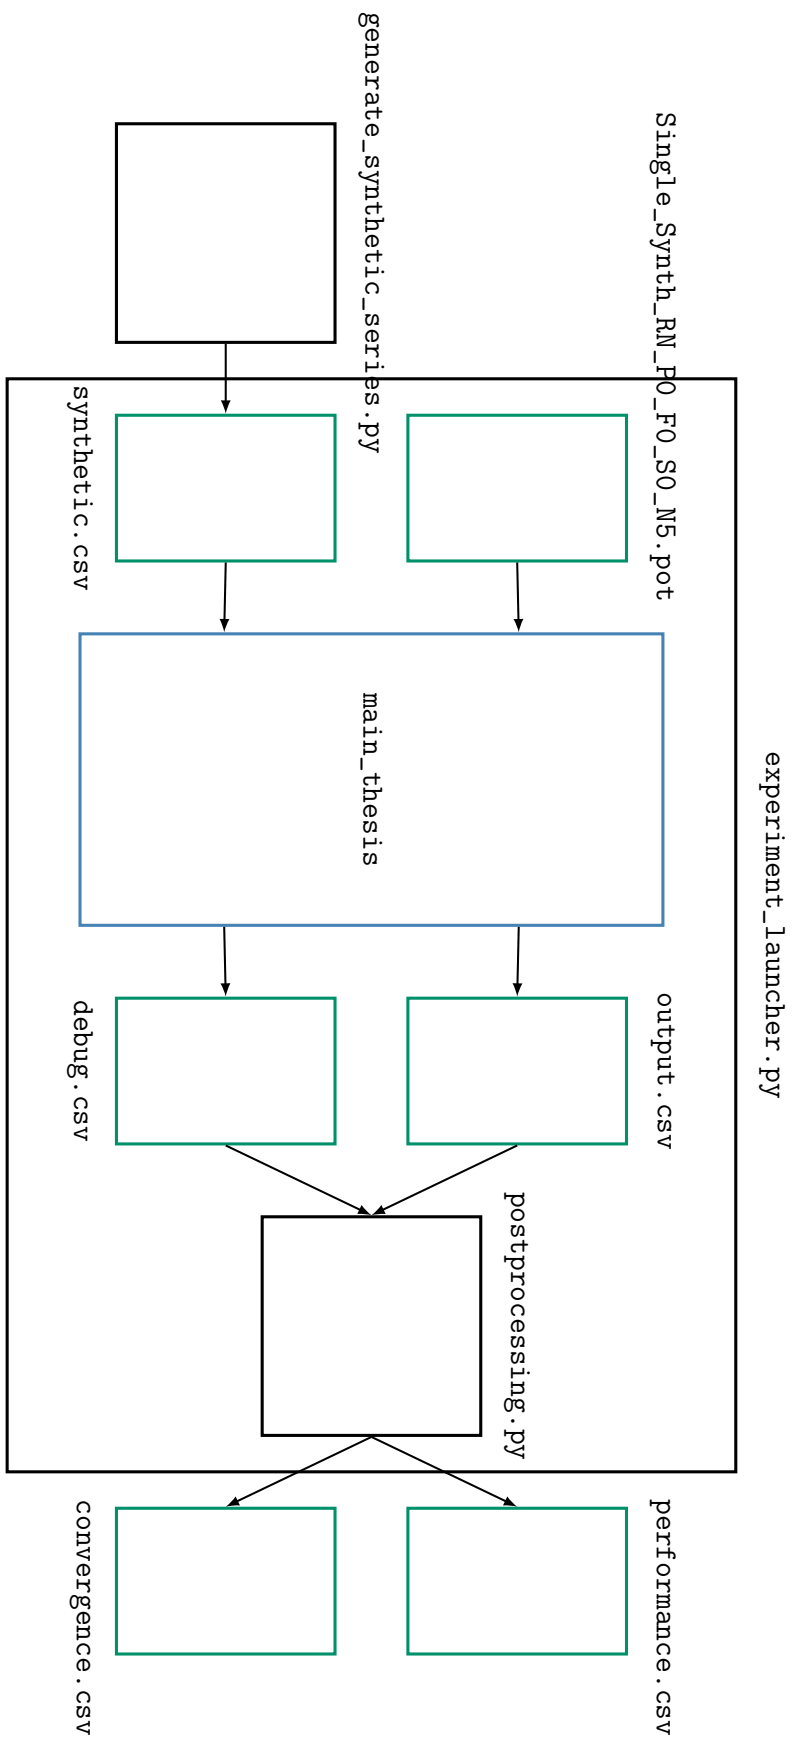


Figure A.4: Execution flow of an asset allocation experiment. Black boxes denote Python scripts, blue boxes executables while red boxes input/output files.

Bibliography

- [1] March of the machines. The Economist, June 2016.
- [2] The return of the machinery question. The Economist, June 2016.
- [3] AGARWAL, A., BARTLETT, P., AND DAMA, M. Optimal allocation strategies for the dark pool problem. *arXiv preprint arXiv:1003.2245* (2010).
- [4] AKIMOTO, Y., NAGATA, Y., ONO, I., AND KOBAYASHI, S. Bidirectional relation between cma evolution strategies and natural evolution strategies. In *International Conference on Parallel Problem Solving from Nature* (2010), Springer, pp. 154–163.
- [5] ALMGREN, R., AND CHRISS, N. Optimal execution of portfolio transactions. *Journal of Risk* 3 (2001), 5–40.
- [6] ALVAREZ TELEÑA, S. *Systematic Trading: Calibration Advances through Machine Learning*. PhD thesis, UCL (University College London), 2015.
- [7] ARAPOSTATHIS, A., BORKAR, V. S., FERNÁNDEZ-GAUCHERAND, E., GHOSH, M. K., AND MARCUS, S. I. Discrete-time controlled markov processes with average cost criterion: a survey. *SIAM Journal on Control and Optimization* 31, 2 (1993), 282–344.
- [8] BASU, A., BHATTACHARYYA, T., AND BORKAR, V. S. A learning algorithm for risk-sensitive cost. *Mathematics of Operations Research* 33, 4 (2008), 880–898.
- [9] BÄUERLE, N., AND RIEDER, U. *Markov decision processes with applications to finance*. Springer Science & Business Media, 2011.
- [10] BAXTER, J., AND BARTLETT, P. L. Infinite-horizon policy-gradient estimation. *Journal of Artificial Intelligence Research* 15 (2001), 319–350.
- [11] BEKIROU, S. D. Heterogeneous trading strategies with adaptive fuzzy actor–critic reinforcement learning: A behavioral approach. *Journal of Economic Dynamics and Control* 34, 6 (2010), 1153–1170.

- [12] BERTOLUZZO, F., AND CORAZZA, M. Testing different reinforcement learning configurations for financial trading: Introduction and applications. *Procedia Economics and Finance* 3 (2012), 68–77.
- [13] BERTOLUZZO, F., AND CORAZZA, M. Q-learning-based financial trading systems with applications. *University Ca'Foscari of Venice, Dept. of Economics Working Paper Series No 15* (2014).
- [14] BERTOLUZZO, F., AND CORAZZA, M. Reinforcement learning for automated financial trading: Basics and applications. In *Recent Advances of Neural Network Models and Applications*. Springer, 2014, pp. 197–213.
- [15] BERTSEKAS, D. P. *Dynamic programming and optimal control*, vol. 1. Athena Scientific, Belmont, 1995.
- [16] BERTSEKAS, D. P., AND SHREVE, S. E. *Stochastic optimal control: the discrete-time case*, vol. 23. Academic Press New York, 1978.
- [17] BERTSEKAS, D. P., AND TSITSIKLIS, J. N. *Neuro-Dynamic Programming*, 1 ed. Optimization and neural computation series. Athena Scientific, Belmont, 1996.
- [18] BHATNAGAR, S., SUTTON, R. S., GHAVAMZADEH, M., AND LEE, M. Natural actor–critic algorithms. *Automatica* 45, 11 (2009), 2471–2482.
- [19] BISHOP, C. M. *Pattern Recognition and Machine Learning*. Springer, 2006.
- [20] BORKAR, V. S. A sensitivity formula for risk-sensitive cost and the actor–critic algorithm. *Systems & Control Letters* 44, 5 (2001), 339–346.
- [21] BORKAR, V. S. Q-learning for risk-sensitive control. *Mathematics of operations research* 27, 2 (2002), 294–311.
- [22] BORKAR, V. S., AND MEYN, S. P. Risk-sensitive optimal control for markov decision processes with monotone cost. *Mathematics of Operations Research* 27, 1 (2002), 192–209.
- [23] BROGAARD, J., HENDERSHOTT, T., AND RIORDAN, R. High-frequency trading and price discovery. *Review of Financial Studies* 27, 8 (2014), 2267–2306.
- [24] BROWNE, S. Optimal investment policies for a firm with a random risk process: exponential utility and minimizing the probability of ruin. *Mathematics of operations research* 20, 4 (1995), 937–958.
- [25] BROWNE, S., ET AL. Reaching goals by a deadline: Digital options and continuous-time active portfolio management. *Advances in Applied Probability* 31, 2 (1999), 551–577.

- [26] BUSONI, L., BABUSKA, R., DE SCHUTTER, B., AND ERNST, D. *Reinforcement learning and dynamic programming using function approximators*, vol. 39. CRC press, 2010.
- [27] CARTEA, Á., JAIMUNGAL, S., AND PENALVA, J. *Algorithmic and high-frequency trading*. Cambridge University Press, 2015.
- [28] CASQUEIRO, P. X., AND RODRIGUES, A. J. Neuro-dynamic trading methods. *European Journal of Operational Research* 175, 3 (2006), 1400–1412.
- [29] CHAPADOS, N., AND BENGIO, Y. Cost functions and model combination for var-based asset allocation using neural networks. *IEEE Transactions on Neural Networks* 12 (2001), 890–906.
- [30] CHOEY, M., AND WEIGEND, A. S. Nonlinear trading models through sharpe ratio maximization. *International Journal of Neural Systems* 8 (1997), 417–431.
- [31] CHOW, Y., TAMAR, A., MANNOR, S., AND PAVONE, M. Risk-sensitive and robust decision-making: a cvar optimization approach. In *Advances in Neural Information Processing Systems* (2015), pp. 1522–1530.
- [32] CORAZZA, M., AND SANGALLI, A. Q-learning vs. sarsa: comparing two intelligent stochastic control approaches for financial trading.
- [33] CUMMING, J., ALRAJEH, D., AND DICKENS, L. An investigation into the use of reinforcement learning techniques within the algorithmic trading domain.
- [34] DEMPSTER, M. A., AND LEEMANS, V. An automated fx trading system using adaptive reinforcement learning. *Expert Systems with Applications* 30, 3 (2006), 543–552.
- [35] DEMPSTER, M. A. H., AND ROMAHI, Y. S. *Intraday FX trading: An evolutionary reinforcement learning approach*. Springer, 2002.
- [36] DENG, Y., BAO, F., KONG, Y., REN, Z., AND DAI, Q. Deep direct reinforcement learning for financial signal representation and trading. *IEEE Transactions on Neural Networks and Learning Systems* (2016).
- [37] DENG, Y., KONG, Y., BAO, F., AND DAI, Q. Sparse coding inspired optimal trading system for hft industry. *IEEE Transactions on Industrial Informatics* 11, 2 (2015), 467–475.
- [38] ELDER, T. Creating algorithmic traders with hierarchical reinforcement learning.
- [39] FAMA, E. F. Efficient capital markets: Ii. *The journal of finance* 46, 5 (1991), 1575–1617.

- [40] FELDKAMP, L. A., PROKHOROV, D. V., EAGEN, C. F., AND YUAN, F. Enhanced multi-stream kalman filter training for recurrent networks. In *Nonlinear Modeling*. Springer, 1998, pp. 29–53.
- [41] GANCHEV, K., NEVMYVAKA, Y., KEARNS, M., AND VAUGHAN, J. W. Censored exploration and the dark pool problem. *Communications of the ACM* 53, 5 (2010), 99–107.
- [42] GITTINS, J., GLAZEBROOK, K., AND WEBER, R. *Multi-armed bandit allocation indices*. John Wiley & Sons, 2011.
- [43] GOLD, C. Fx trading via recurrent reinforcement learning. In *IEEE 2003 IEEE International Conference on Computational Intelligence for Financial Engineering. Proceedings* (2003).
- [44] GOODFELLOW, I., BENGIO, Y., AND COURVILLE, A. Deep learning. Book in preparation for MIT Press, 2016.
- [45] HASTIE, T., TIBSHIRANI, R., AND FRIEDMAN, J. *The elements of statistical learning: data mining, inference, and prediction*. Springer, 2009.
- [46] HENDRICKS, D., AND WILCOX, D. A reinforcement learning extension to the almgren-chriss framework for optimal trade execution. In *IEEE Conference on Computational Intelligence for Financial Engineering & Economics (CIFEr)* (2014), pp. 457–464.
- [47] JAEGER, H. *Tutorial on training recurrent neural networks, covering BPPT, RTRL, EKF and the echo state network approach*. GMD-Forschungszentrum Informationstechnik, 2002.
- [48] JENSEN, M. C. Some anomalous evidence regarding market efficiency. *Journal of financial economics* 6, 2/3 (1978), 95–101.
- [49] JOHNSON, B. *Algorithmic Trading & DMA: An introduction to direct access trading strategies*, vol. 200. 4Myeloma Press London, 2010.
- [50] JOSHI, M. S. *C++ design patterns and derivatives pricing*, vol. 2. Cambridge University Press, 2008.
- [51] KAKADE, S. A natural policy gradient. In *NIPS* (2001), vol. 14, pp. 1531–1538.
- [52] KAKADE, S. Optimizing average reward using discounted rewards. In *International Conference on Computational Learning Theory* (2001), Springer, pp. 605–615.
- [53] KAMIJO, K., AND TANIGAWA, T. Stock price pattern recognition-a recurrent neural network approach. In *IEEE 1990 IJCNN International Joint Conference on Neural Networks* (1990).

- [54] KEARNS, M., AND NEVMYVAKA, Y. Machine learning for market microstructure and high frequency trading. *High-Frequency Trading—New Realities for Traders, Markets and Regulators* (2013), 91–124.
- [55] KONDA, V. R., AND TSITSIKLIS, J. N. Actor-critic algorithms. In *NIPS* (1999), vol. 13, pp. 1008–1014.
- [56] KUMAR, N., AND HALL, T. Why machine still can’t learn so good. Bloomberg, November 2016.
- [57] KUSHNER, H., AND YIN, G. G. *Stochastic approximation and recursive algorithms and applications*, vol. 35. Springer Science & Business Media, 2003.
- [58] L.A., P., AND GHAVAMZADEH, M. Actor-critic algorithms for risk-sensitive mdps. In *Advances in Neural Information Processing Systems 26*, C. J. C. Burges, L. Bottou, M. Welling, Z. Ghahramani, and K. Q. Weinberger, Eds. Curran Associates, Inc., 2013, pp. 252–260.
- [59] LARUELLE, S., LEHALLE, C.-A., AND PAGES, G. Optimal split of orders across liquidity pools: a stochastic algorithm approach. *SIAM Journal on Financial Mathematics* 2, 1 (2011), 1042–1076.
- [60] LARUELLE, S., LEHALLE, C.-A., AND PAGES, G. Optimal posting price of limit orders: learning by trading. *Mathematics and Financial Economics* 7, 3 (2013), 359–403.
- [61] LEWIS, M., AND BAKER, D. *Flash boys*. Allen Lane, 2014.
- [62] LI, H., DAGLI, C. H., AND ENKE, D. Short-term stock market timing prediction under reinforcement learning schemes. In *2007 IEEE International Symposium on Approximate Dynamic Programming and Reinforcement Learning* (2007), pp. 233–240.
- [63] LIANG, J., SONG, W., AND WANG, M. Stock price prediction based on procedural neural networks. *Adv. Artif. Neu. Sys. 2011* (2011), 1–11.
- [64] LITTMAN, M. L. *Algorithms for sequential decision making*. PhD thesis, Brown University, 1996.
- [65] LO, A. W., MAMAYSKY, H., AND WANG, J. Foundations of technical analysis: Computational algorithms, statistical inference, and empirical implementation. *The journal of finance* 55, 4 (2000), 1705–1770.
- [66] MAHADEVAN, S. Average reward reinforcement learning: Foundations, algorithms, and empirical results. *Machine learning* 22, 1-3 (1996), 159–195.
- [67] MALKIEL, B. G. The efficient market hypothesis and its critics. *The Journal of Economic Perspectives* 17, 1 (2003), 59–82.

- [68] MALKIEL, B. G., AND FAMA, E. F. Efficient capital markets: A review of theory and empirical work. *The journal of Finance* 25, 2 (1970), 383–417.
- [69] MALLABY, S. *More money than god: Hedge funds and the making of the new elite*. A&C Black, 2010.
- [70] MARKOWITZ, H. Portfolio selection. *The journal of finance* 7, 1 (1952), 77–91.
- [71] MIYAMAE, A., NAGATA, Y., ONO, I., AND KOBAYASHI, S. Natural policy gradient methods with parameter-based exploration for control tasks. In *Advances in neural information processing systems* (2010), pp. 1660–1668.
- [72] MOODY, J., AND SAFFELL, M. Learning to trade via direct reinforcement. *Neural Networks, IEEE Transactions on* 12, 4 (2001), 875–889.
- [73] MOODY, J., SAFFELL, M., LIAO, Y., AND WU, L. Reinforcement learning for trading systems. In *Decision Technologies for Computational Finance: Proceedings of the fifth International Conference Computational Finance* (2013), vol. 2, Springer Science & Business Media, p. 129.
- [74] MOODY, J., AND WU, L. Optimization of trading systems and portfolios. In *Proceedings of the IEEE/IAFE 1997 Computational Intelligence for Financial Engineering (CIFEr)* (1997), pp. 300–307.
- [75] MOODY, J., WU, L., LIAO, Y., AND SAFFELL, M. Performance functions and reinforcement learning for trading systems and portfolios. *Journal of Forecasting* 17 (1998), 441–470.
- [76] NEVMYVAKA, Y., FENG, Y., AND KEARNS, M. Reinforcement learning for optimized trade execution. In *Proceedings of the 23rd international conference on Machine learning* (2006), ACM, pp. 673–680.
- [77] NOCEDAL, J., AND WRIGHT, S. *Numerical optimization*. Springer Science & Business Media, 2006.
- [78] O, J., LEE, J., LEE, J. W., AND ZHANG, B.-T. Adaptive stock trading with dynamic asset allocation using reinforcement learning. *Information Sciences* 176, 15 (2006), 2121–2147.
- [79] PAGES, G. *Introduction to Numerical Probability for Finance*. 2016.
- [80] PETERS, J., MÜLLING, K., AND ALTUN, Y. Relative entropy policy search. In *AAAI* (2010), Atlanta.
- [81] PETERS, J., AND SCHAAL, S. Policy gradient methods for robotics. In *Intelligent Robots and Systems, 2006 IEEE/RSJ International Conference on* (2006), IEEE, pp. 2219–2225.

- [82] PETERS, J., AND SCHAAL, S. Reinforcement learning of motor skills with policy gradients. *Neural networks* 21, 4 (2008), 682–697.
- [83] PRASHANTH, L., AND GHAVAMZADEH, M. Actor-critic algorithms for risk-sensitive reinforcement learning. *arXiv preprint arXiv:1403.6530* (2014).
- [84] PUTERMAN, M. L. *Markov decision processes: discrete stochastic dynamic programming*. John Wiley & Sons, 1994.
- [85] RUSSO, C., AND DETRIXHE, J. Wall street’s speed demons are heroes. Bloomberg, October 2016.
- [86] SAAD, E. W., PROKHOROV, D. V., AND WUNSCH, D. C. Comparative study of stock trend prediction using time delay, recurrent and probabilistic neural networks. *IEEE Transactions on Neural Networks* 9, 6 (1998), 1456–1470.
- [87] SANDERSON, C. Armadillo: An open source c++ linear algebra library for fast prototyping and computationally intensive experiments.
- [88] SATO, M., AND KOBAYASHI, S. Variance-penalized reinforcement learning for risk-averse asset allocation. In *Intelligent Data Engineering and Automated Learning—IDEAL 2000. Data Mining, Financial Engineering, and Intelligent Agents*. Springer, 2000, pp. 244–249.
- [89] SATO, M., AND KOBAYASHI, S. Average-reward reinforcement learning for variance penalized markov decision problems. In *Proceedings of the Eighteenth International Conference on Machine Learning* (San Francisco, CA, USA, 2001), ICML ’01, Morgan Kaufmann Publishers Inc., pp. 473–480.
- [90] SCHAUL, T., BAYER, J., WIERSTRA, D., SUN, Y., FELDER, M., SEHNKE, F., RÜCKSTIESS, T., AND SCHMIDHUBER, J. PyBrain. *Journal of Machine Learning Research* 11 (2010), 743–746.
- [91] SEHNKE, F., ET AL. *Parameter exploring policy gradients and their implications*. PhD thesis, Technische Universität München, 2012.
- [92] SEHNKE, F., OSENDORFER, C., RÜCKSTIESS, T., GRAVES, A., PETERS, J., AND SCHMIDHUBER, J. Policy gradients with parameter-based exploration for control. In *Artificial Neural Networks-ICANN 2008*. Springer, 2008, pp. 387–396.
- [93] SEHNKE, F., OSENDORFER, C., RÜCKSTIESS, T., GRAVES, A., PETERS, J., AND SCHMIDHUBER, J. Parameter-exploring policy gradients. *Neural Networks* 23, 4 (2010), 551–559.
- [94] SHARPE, W. F. The sharpe ratio. *The journal of portfolio management* 21, 1 (1994), 49–58.

- [95] SILVER, D., HUANG, A., MADDISON, C. J., GUEZ, A., SIFRE, L., VAN DEN DRIESSCHE, G., SCHRITTWIESER, J., ANTONOGLU, I., PANNEERSHELVAM, V., LANCTOT, M., ET AL. Mastering the game of go with deep neural networks and tree search. *Nature* 529, 7587 (2016), 484–489.
- [96] SILVER, D., LEVER, G., HEES, N., DEGRIS, T., WIERSTRA, D., AND RIEDMILLER, M. Deterministic policy gradient algorithms. In *ICML* (2014).
- [97] SOBEL, M. J. The variance of discounted markov decision processes. *Journal of Applied Probability* (1982), 794–802.
- [98] SUN, Y., WIERSTRA, D., SCHAUL, T., AND SCHMIDHUBER, J. Efficient natural evolution strategies. In *Proceedings of the 11th Annual conference on Genetic and evolutionary computation* (2009), ACM, pp. 539–546.
- [99] SUTTON, R. S., AND BARTO, A. G. *Introduction to reinforcement learning*, vol. 135. MIT Press Cambridge, 1998.
- [100] SUTTON, R. S., MCALLESTER, D. A., SINGH, S. P., MANSOUR, Y., ET AL. Policy gradient methods for reinforcement learning with function approximation. In *NIPS* (1999), vol. 99, pp. 1057–1063.
- [101] SZEPESVÁRI, C. Algorithms for reinforcement learning. *Synthesis lectures on artificial intelligence and machine learning* 4, 1 (2010), 1–103.
- [102] TAMAR, A., CASTRO, D. D., AND MANNOR, S. Temporal difference methods for the variance of the reward to go. In *Proceedings of the 30th International Conference on Machine Learning (ICML-13)* (2013), pp. 495–503.
- [103] TAMAR, A., CHOW, Y., GHAVAMZADEH, M., AND MANNOR, S. Policy gradient for coherent risk measures. In *Advances in Neural Information Processing Systems* (2015), pp. 1468–1476.
- [104] TAMAR, A., DI CASTRO, D., AND MANNOR, S. Policy gradients with variance related risk criteria. In *Proceedings of the 29th International Conference on Machine Learning* (2012), pp. 387–396.
- [105] TAMAR, A., AND MANNOR, S. Variance adjusted actor critic algorithms. *arXiv preprint arXiv:1310.3697* (2013).
- [106] TAN, Z., QUEK, C., AND CHENG, P. Y. Stock trading with cycles: A financial application of anfis and reinforcement learning. *Expert Systems with Applications* 38, 5 (2011), 4741–4755.
- [107] TIMMERMAN, A., AND GRANGER, C. W. Efficient market hypothesis and forecasting. *International Journal of forecasting* 20, 1 (2004), 15–27.

- [108] TSAY, R. S. *Analysis of financial time series*, vol. 543. John Wiley & Sons, 2005.
- [109] WERBOS, P. J. Backpropagation through time: what it does and how to do it. *Proceedings of the IEEE* 78, 10 (1990), 1550–1560.
- [110] WIERING, M., AND VAN OTTERLO, M. *Reinforcement Learning: State-of-the-Art*, 1 ed., vol. 12 of *Adaptation, Learning, and Optimization*. Springer, 2012.
- [111] WILLIAMS, R. J., AND ZIPSER, D. A learning algorithm for continually running fully recurrent neural networks. *Neural computation* 1, 2 (1989), 270–280.
- [112] YANG, S., PADDRIK, M., HAYES, R., TODD, A., KIRILENKO, A., BELING, P., AND SCHERER, W. Behavior based learning in identifying high frequency trading strategies. In *IEEE Conference on Computational Intelligence for Financial Engineering & Economics (CIFEr)* (2012).
- [113] ZHAO, T., HACHIYA, H., NIU, G., AND SUGIYAMA, M. Analysis and improvement of policy gradient estimation. In *Advances in Neural Information Processing Systems* (2011), pp. 262–270.
- [114] ZHAO, T., NIU, G., XIE, N., YANG, J., AND SUGIYAMA, M. Regularized policy gradients: Direct variance reduction in policy gradient estimation. In *Proceedings of The 7th Asian Conference on Machine Learning* (2015), pp. 333–348.